



UMD DATA605 - Big Data Systems

NoSQL Document Stores MongoDB CouchDB

Instructor: Dr. GP Saggese - gsaggese@umd.edu**

TAs: Krishna Pratardan Taduri, kptaduri@umd.edu Prahar

Kaushikbhai Modi, pmodi08@umd.edu

with thanks to Prof. Alan Sussman Prof. Amol

Deshpande Prof. Nguyen Vo (Utah State U.)

Prof. Kathleen Durant (Northeastern U.)

v1.1

UMD DATA605 - Big Data Systems

Key-Value Store vs Document DBs

- **Key-value stores**

- Basically a map or a dictionary
 - E.g., HBase, Redis
- Typically only look up values by key
 - Sometimes can do search in value field with a pattern
- Uninterpreted value (e.g., binary blob) associated with a key
- Typically one namespace for all key-values

- **Document DBs**

- Collect sets of key-value pairs into *documents*
 - E.g., MongoDB, CouchDB
- Documents represented in JSON, XML, or BSON (binary JSON)
- Documents organized into *collections*
 - Similar to *tables* in relational DBs
 - Large collections can be partitioned and indexed

Key-Value



Resources

- All concepts in slides
- MongoDB tutorial
- Web
 - <https://www.mongodb.com/>
 - Official docs
 - pymongo
- Book
 - Seven Databases in Seven Weeks, 2e

The
Pragmatic
Programmers

Seven Databases in Seven Weeks

Second Edition

A Guide to Modern
Databases and the
NoSQL Movement



MongoDB

- Developed by MongoDB Inc.
 - Founded in 2007
 - Based on DoubleClick experience with large-scale data
 - Mongo comes from “hu-mongo-us”
- One of the most used NoSQL DBs (if not the most used)
- **Document-oriented NoSQL database**
 - Schema-less
 - No Data Definition Language (DDL), like for SQL
 - You can store maps with any keys and values
 - Application tracks the schema, mapping between documents and their meaning
 - Keys are hashes stored as strings
 - Document Identifiers **_id** created for each document (field name reserved by Mongo)
 - Values use BSON format
 - Based on JSON (B stands for Binary)
- Written in C++
- Supports APIs (drivers) in many languages
 - E.g., JavaScript, Python, Ruby, Java, Scala, C++, ...

MongoDB: Example of Document

- **A document is a JSON data structure**
- It corresponds to a row in a relational DB
 - Without schema
 - Primary key is _id
 - Values nested to an arbitrary depth

```
{
  "_id" : ObjectId("4d0b6da3bb30773266f39fea"),
  "country" : {
    "$ref" : "countries",
    "$id" : ObjectId("4d0e6074deb8995216a8309e")
  },
  "famous_for" : [
    "beer",
    "food"
  ],
  "last_census" : "Sun Jan 07 2018 00:00:00 GMT -0700 (PDT)",
  "mayor" : {
    "name" : "Ted Wheeler",
    "party" : "D"
  },
  "name" : "Portland",
  "population" : 582000,
  "state" : "OR"
}
```

MongoDB: Functionalities

- **Design goals**
 - Performance
 - Availability / scalability
 - Rich data storage (not rich querying!)
- **Dynamic schema**
 - No DDL (Data Definition Language)
 - Secondary indexes
 - Query language via an API
- **Several levels of data consistency**
 - E.g., atomic writes and fully-consistent reads (at document level)
- **No joins nor transactions across multiple documents**
 - Makes distributed queries easy and fast
- **High availability through replica sets**
 - E.g., primary replication with automated failover
- **Built-in sharding**
 - Horizontal scaling via automated range-based partitioning of data
 - Reads and writes distributed over shards

MongoDB: Hierarchical Objects

- A Mongo **instance** has:
 - Zero or more “databases”
 - Mongo instance ~ Postgres instance
- A Mongo **database** has:
 - Zero or more “collections”
 - Mongo collection ~ Postgres tables
 - Mongo database ~ Postgres database
- A Mongo **collection** has:
 - Zero or more “documents”
 - Mongo document ~ Postgres rows
- A Mongo **document** has:
 - One or more “fields”
 - It has always primary key `_id`
 - Mongo field ~ Postgres columns

Collection

Document

Database



Relational DBs vs MongoDB: Terms and Concepts

RDBMS Concept	MongoDB Concept	Meaning in MongoDB
database	database	Container for collections
relation / table / view	collection	Group of documents
row / instance	document	Group of fields
column / attribute	field	A name-value pair
index	index	Automatic
primary keys	_id field	Always the primary key
foreign key	reference	Pointers
table joins	embedded documents	Nested name-value pairs

```
{
  "_id" : ObjectId("4d0b6da3bb30773266f39fea"),
  "country" : {
    "$ref" : "countries",
    "$id" : ObjectId("4d0e6074deb8995216a8309e")
  },
  "famous_for" : [
    "beer",
    "food"
  ],
}
```


Relational vs Document DB: Workflows

- **Relational DBs**


- E.g., PostgreSQL
- Know what you want to store
 - Tabular data
- Do not know how to use it
 - Static schema allows query flexibility (e.g., joins)
- Complexity is at insertion time
 - Decide how to represent the data (i.e., schema)

- **Document DBs**

- E.g., MongoDB
- No assumptions on what to store
 - E.g., irregular JSON data
- Know a bit how to access data
 - You want to access the data by key
 - E.g., it's a nested key-value map
- Complexity is at access time
 - Get the data from the server
 - Process data on the client side

Why Use MongoDB?

- Simple to query
 - Do the work on client side
- It's fast
 - 2-10x faster than Postgres
- Data model / functionalities suitable for most web applications
 - Semi-structured data
 - Quickly evolving systems
- Easy and fast integration of data
- Not well suited for heavy and complex transactions systems
 - E.g., banking system



MongoDB: Data Model

- **Documents** are composed of field and value pairs
 - **Field names** are strings
 - **Values** are any BSON type
 - Arrays of documents
 - Native data types
 - Other documents
- E.g.,
 - `_id` holds an ObjectId
 - `name` holds a document that contains the fields `first` and `last`
 - `birth` and `death` are of Date type
 - `contribs` holds an array of strings
 - `views` holds a value of the NumberLong type

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```

← field: value
← field: value
← field: value
← field: value

MongoDB: Data Model

```
{
  _id: <ObjectId1>,
  username: "123xyz",
  contact: {
    phone: "123-456-7890",
    email: "xyz@example.com"
  },
  access: {
    level: 5,
    group: "dev"
  }
}
```

Embedded sub-document

Embedded sub-document

- Documents can be nested
 - Embedded sub-document

SCIENCE Denormalized data models

- Store multiple related pieces of information in the same record

Schema Free

- MongoDB does not need any pre-defined data schema
- Every **document** in a **collection** can have different fields and values
 - No need for NULL values / union of fields like in relational DBs
- E.g., dishomogeneous data instances


```
{name: "will",  
  eyes: "blue",  
  birthplace: "NY",  
  aliases: ["bill", "la ciacco"],  
  loc: [32.7, 63.4],  
  boss: "ben"}
```

```
{name: "jeff",  
  eyes: "blue",  
  loc: [40.7, 73.4],  
  boss: "ben"}
```

```
{name: "brendan",  
  aliases: ["el diablo"]}
```

```
{name: "ben",  
  hat: "yes"}
```

```
{name: "matt",  
  pizza: "DiGiorno",  
  height: 72,  
  loc: [44.6, 71.3]}
```



JSON Format

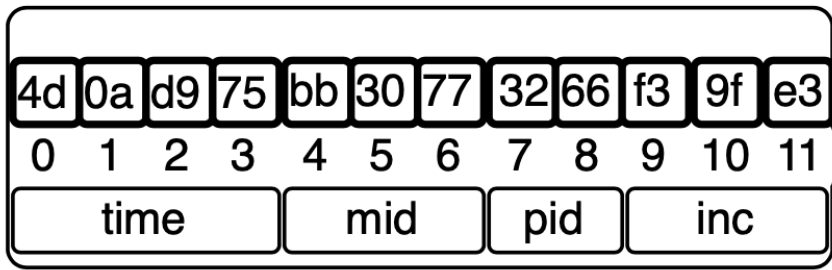
- JSON = JavaScript Object Notation
- Data is stored in field / value pairs
- A field / value pair consists of:
 - A field name (always a string)
 - Followed by a colon :
 - Followed by a typed value "name": "R2-D2"
- Data in documents is separated by commas , "name": "R2-D2", race : "Droid"
- Curly braces {} hold documents {"name": "R2-D2", race : "Droid", affiliation: "rebels"}
- An array is stored in brackets [] [{"name": "R2-D2", race: "Droid", affiliation: "rebels"}, {"name": "Yoda", affiliation: "rebels"}]
- Supports:
 - Embedding of nested objects within other objects
 - Just references

BSON Format

- Binary-encoded serialization of JSON-like documents
 - <https://bsonspec.org>
- Zero or more key/value pairs are stored as a single entity
 - Each entry consists of:
 - a field name (string)
 - a data type
 - a value
- Similar to Protocol Buffer, but more schema-less
- Large elements in a BSON document are prefixed with a length field to facilitate scanning
- MongoDB understands the internals of BSON objects, even nested ones
 - Can build indexes and match objects against query expressions for BSON keys

ObjectId

- Each JSON data contains an `_id` field of type `ObjectId`
 - Same as a SERIAL constraint incrementing a numeric primary key in PostgreSQL
- An `ObjectId` is 12 bytes, composed of:
 - a timestamp
 - client machine ID
 - client process ID
 - a 3-byte auto-incremented counter
- Each Mongo process can handle its own ID generation without colliding
 - Mongo has a distributed nature
- Details here



Indexes

- **Primary index**
 - Automatically created on the `_id` field
 - B+ tree indexes
- **Secondary index**
- Users can create secondary indexes to:
 - Improve query performance
 - Enforce unique values for a particular field
- Single field index and compound index (like SQL)
 - Order of the fields in a compound index matters
- Sparse property of an index
 - The index contains only entries for documents that have the indexed field
 - Ignore records that do not have the field defined
- Reject records with duplicate key value if an index is unique and sparse
- Details at <https://www.mongodb.com/docs/manual/indexes/>

CRUD Operations

- CRUD = Create, Read, Update, Delete
- **Create** `db.collection.insert(<document>)`
`db.collection.update(<query>, <update>, upsert: true)` Upsert = update (if exists) or insert (if it doesn't)
- **Read** `db.collection.find(<query>, <projection>)`
`db.collection.findOne(<query>, <projection>)`
- **Update** `db.collection.update(<query>, <update>, <options>)`
- **Delete** `db.collection.remove(<query>, <justOne>)` Details at <https://www.mongodb.com/docs/manual/crud/>

Create Operations

- **db.collection** specifies the collection (like an SQL table) to store the document **db.collection.insert(<document>)**
 - Without **_id** field, MongoDB generates a unique key **db.parts.insert(type: "screwdriver", quantity: 15)**
 - Use **_id** field if it has a special meaning **db.parts.insert(_id: 10, type: "hammer", quantity: 1)**
- Update 1 or more records in a collection satisfying **query** **db.collection.update(<query>, <update>, upsert: true)**
- Update an existing record or create a new record **db.collection.save(<document>)**
- A more modern OOP-like syntax than the COBOL / FORTRAN-inspired SQL

Read Operations

- **find** provides functionality similar to **SQL SELECT** command
db.collection.find(<query>, <projection>).cursor with:
 - = WHERE condition
 - = fields in result set
- **db.parts.find(parts: "hammer").limit(5)**
 - Return cursor to handle a result set
 - Can modify the query to impose limits, skips, and sort orders
 - Can specify to return the 'top' number of records from the result set
- **db.collection.findOne(<query>, <projection>)**

More Query Examples

```
SELECT * FROM users WHERE age>33
```

```
SELECT * FROM users WHERE age!=33
```

```
SELECT * FROM users WHERE name LIKE "%Joe%"
```

```
SELECT * FROM users WHERE a=1 and b='q'
```

```
SELECT * FROM users WHERE a=1 or b=2
```

```
**SELECT * FROM foo
```

```
WHERE name='bob' and (a=1 or b=2 )
```

```
SELECT * FROM users
```

```
WHERE age>33 AND age<=40
```

```
**Mongo**
```

```
**`db.users.find({age: {$gt: 33}})`
```

```
**
```

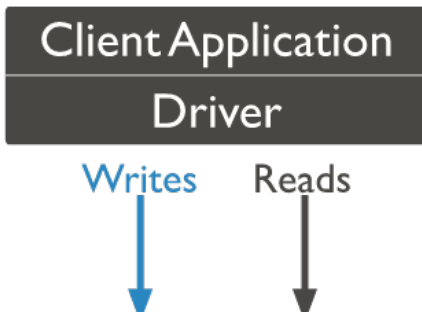
```
**`db.users.find({age: {$ne: 33}})`
```

```
**ACADEMY
```

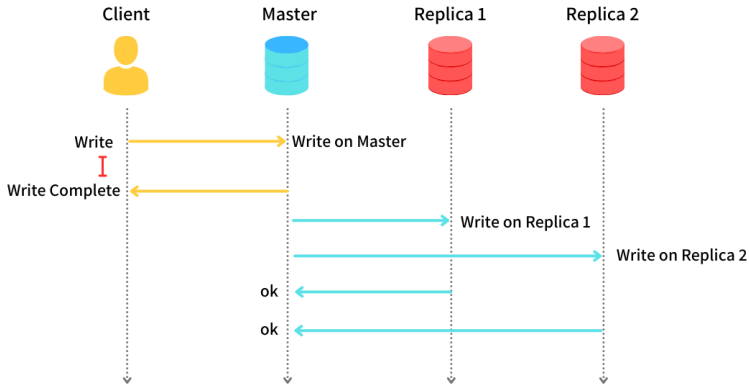


Data Replication

- **Data replication** ensure:
 - Redundancy
 - Backup
 - Automatic failover
- Replication occurs through groups of servers known as **replica sets**
 - **Primary set**: set of servers that client asks direct updates to
 - **Secondary set**: set of servers used for duplication of data
 - Different properties can be associated with a secondary set,
 - E.g., secondary-only, hidden delayed, arbiters, non-voting
- If the primary fails the secondary sets “vote” to elect the new primary set



Sync vs Async Replication



Asynchronous Replication

- **Synchronous replication:** updates are propagated to other replicas as part of a single transaction
- Implementations


SCIENCE 21
ACADEMY
Phase Commit (2PC)
Paxos

Data Consistency

- **Client decides how to enforce consistency for reads**
- Reads to a primary have **strict consistency**
 - Reads reflect the latest changes to the data
 - All writes and *consistent* reads go to the primary
- Reads to a secondary have **eventual consistency**
 - Updates propagate gradually
 - Client may read a previous state of the database
 - All *eventually consistent* reads are distributed among the secondaries

MongoDB: Sharding

- **Shard** = subset of data
 - A collection is split in pieces based on the shard key
 - Data distributed based on shard key or intervals $[a, b]$
- **Sharding** = method for distributing data across different machines
- **Horizontal scaling** can be achieved through sharding
 - Divide data and workload over multiple servers
 - Complexity in infrastructure and maintenance
- **mongos** acts as a query router interfacing clients and sharded cluster
 - Each shard can be deployed as a replica set
 - Config servers store metadata and configuration settings for cluster



RDMBS Internals

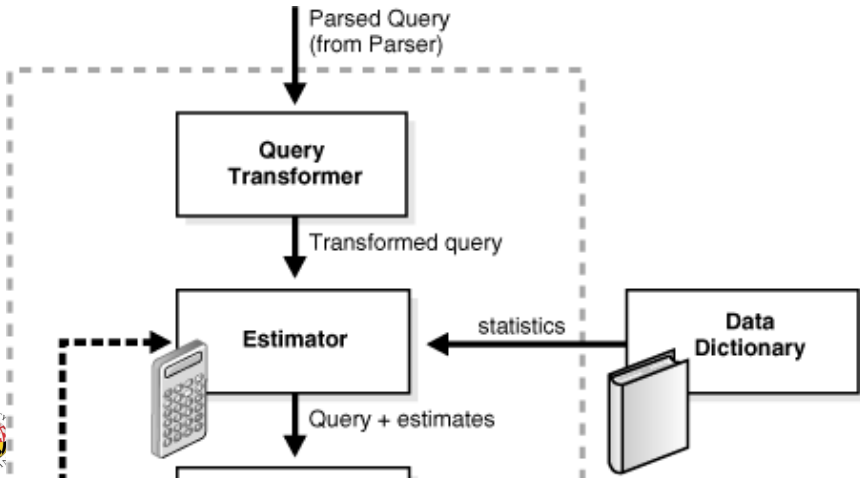
Storage hierarchy - How are tables mapped to files? - How are tuples mapped to disk blocks?

Buffer Manager - Bring pages from disk to memory - Manage the limited memory

Query Processing Engine - Given a user query, decide how to “execute” it - Specify sequence of pages to be brought in memory - Operate upon the tuples to produce results

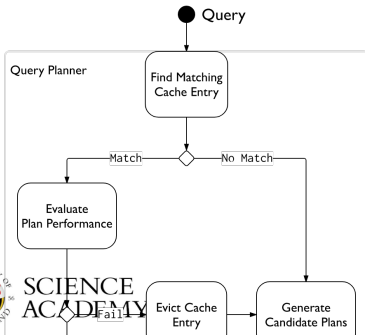
Query Optimizer

- **RDBMSs: query optimizer is static**
 - Assign a cost to each query plan
 - Estimate some cost params (e.g., time to access data)
 - Search for the best query
 - At least traditional RDBMs



Query Optimizer

- **MongoDB: query optimizer is dynamic**
 - Try different query plans and learn which ones perform well
 - The space of query plans is not so large, because there are no joins
 - When testing new plans
 - Execute multiple query plans in parallel
 - As soon as one plan finishes, terminate the other plans
 - Cache the result
 - If a plan that was working well starts performing poorly try again different plans
 - E.g, data in the DB has changed, parameter values to a query are different

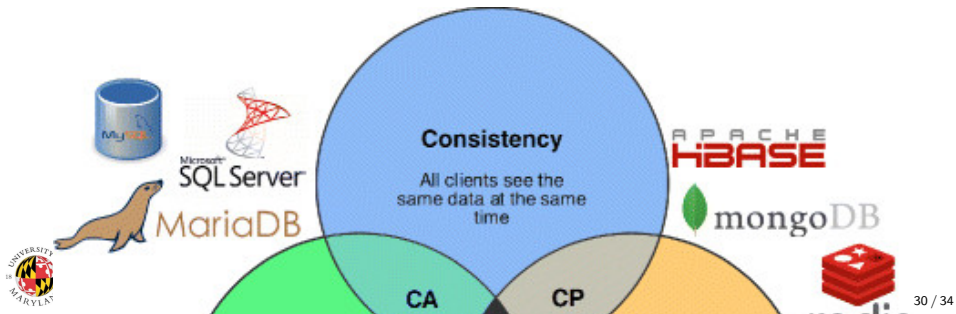


MongoDB: Strengths

- Provide a flexible and modern query language
- High-performance
 - Implemented in C++
- Very rapid development, open source
 - Support for many platforms
 - Many language drivers
- Built to address a distributed database system
 - Sharding
 - Replica sets of data
- Tunable consistency
- Useful for working with a huge quantity of data not requiring a relational model
 - The relationships between the elements does not matter
 - What matters is the ability to store and retrieve great quantities of data

MongoDB: Limitations

- No referential integrity
 - Aka foreign key constraint
- Lack of transactions and joins
- High degree of denormalization
 - Need to update data in many places instead of one
- Lack of predefined schema is a double- edged sword
 - You must have a data model in your application
 - Objects within a collection can be completely inconsistent in their fields
- CAP Theorem: targets consistency and partition tolerance, giving up on availability



UMD DATA605 - Big Data Systems

Couchbase


- NoSQL document-oriented DB (like MongoDB)
- Couchbase = merge of CouchDB and membase
 - *CouchDB*
 - Open source document store
 - HTTP RESTful API to add, update, delete documents
 - Support all 4 ACID properties
 - *membase*
 - Distributed key-value store (like Redis)
 - Designed to scale both up and down
 - Highly available and partition tolerant
 - Uses HTTP protocol to query and interact with objects in the DB
 - No query language
 - Objects stored in *buckets*
 - Collection of JSON docs, with no special relation to one another
- From CAP point of view:
 - Supports consistency and partition tolerance
 - High availability is achieved through use of multiple clusters



Couchbase

Architecture

- Every Couchbase node consists of **different services**:
 - Data service
 - Index service
 - Query service
 - Cluster manager component
- Services can run on separate nodes of the cluster, if needed
- **Data replication**
 - Across nodes of a cluster
 - Across data centers
- **Data service**
 - Writes data *asynchronously* to disk after acknowledging to the client
 - Optionally *synchronous*: ensure data is written to more than one server before acknowledging a write



Queries

- **Can create multiple views over documents**
 - Views are optimized / indexed by Couchbase for fast queries
 - Re-indexed when underlying documents changes
 - Can do full-text searches using the indexes
- **Perform well when:**
 - There are infrequent changes to the structure of documents
 - Know in advance what kinds of queries you want to execute
- **Query**
 - Uses a custom query language called N1QL (“nickel”)
 - Extends SQL to JSON documents
 - Queries over multiple documents using (server-side) joins
- **Map-reduce support**
 - (Map) First define a view with the columns of the document your are interested in
 - (Reduce) Optionally define aggregate functions over the data