



DATA605: Big Data Systems

Introduction

Instructor: Dr. GP Saggese - gsaggese@umd.edu

References:

- *Class Intro*
- Big Data
- Data Models

Invariants of a Class Lecture

- **Invariants**

- Focus on intuition over math
- Emphasize realistic assumptions, numerical methods
 - Analytical solutions are outdated
- Interactive Jupyter notebook tutorials
 - Tutorials done at home
 - Videos added over time

- **Class flow**

- Alternate between slides, whiteboard, tutorials

- **Labs**

- Review complete class project examples
- Collaborate on class project



Books of the Class

- Slides
 - Are extracted from books, technical articles, Internet
 - Should be self-sufficient

Learning Outcomes

- Model and reason about data
- Process and manipulate data
 - E.g., Python, Pandas
- Introduce a variety of data models
 - E.g., relational, NoSQL, graph DBs
 - Decide appropriate data model for different applications
- Use data management systems
 - E.g., PostgreSQL, MongoDB, HBase
 - Decide appropriate system for scenarios
- Build data processing pipelines
 - E.g., Docker, Airflow
- Build a big-data system end-to-end
 - Class project
 - Contribute to an open-source project



Tools We Will Learn To Use

- **Programming languages**
 - Python
- **Development tools**
 - Bash/Linux OS
 - Git: data model, branching
 - GitHub: Pull Requests (PR), issues
 - Jupyter notebooks
 - Docker
- **Big data tools**
 - Extract-Transform-Load (ETL) pipelines
 - Relational DBs (PostgreSQL)
 - NoSQL DBs (HBase, MongoDB, Couchbase, Redis)
 - Graph DBs (Neo4j, GraphX, Giraph)
 - Computing framework (Hadoop, Spark, Dask)
 - Workflow manager (Airflow)
 - Cloud services (AWS)
- **Tutorials** for tools used in the class projects

Todos

- Study slides and materials
- DATA605 - ELMS/Canvas site
 - Enable notifications
 - Contact info for me/TAs
- DATA605 - Schedule
- DATA605 - GitHub repo
- Setup computing environment
 - Install Linux/VMware
 - Install Docker on laptop
 - Instructions in class repo
- Bring laptop to class
 - Quizzes at class start
- Lessons recorded
 - Attend class!



Grading

- **Quizzes**
 - 40% of grade
 - Multi-choice on previous 2 lessons
 - 20 questions in 20 minutes
 - 4-5 quizzes to encourage study during semester
- **Final Project**
 - 60% of grade
 - Comprehensive application of course concepts
 - Big data project in Python from a list of topics
 - Individual or group

Class Projects

- The project is “*Build X with Y*”, where *X* is a “use case” and *Y* is a “technology”
 - Study and describe technology *Y*
 - Implement use case *X* using technology *Y*
 - Create Jupyter notebooks to demo your project
 - Commit code to GitHub, contribute to open-source repo
 - Write a blog entry
 - Present your project in a video
- Choose from list of *X* and *Y*, e.g.,
 - Big data
 - Large language models
 - ...
- Each project:
 - Individual or group ($n < 4$)
 - Varying difficulty levels



Soft Skills to Succeed in the Workplace

- **Goal:** model class project for workplace preparation
 - Work in a team
 - Design software architecture (OOP, Agile, Design Patterns)
 - Comment your code
 - Write external documentation (tutorials, manuals, how-tos)
 - Write understandable code (including for future-you)
 - Read others' code
 - Follow code conventions (PEP8, Google Code)
 - Communicate clearly (emails, Slack)
 - File a bug report
 - Reproduce a bug
 - Intuition of CS constants
 - Basic understanding of OS (virtual memory, processes)



Links

- [ELMS](#)
- [Syllabus](#)
 - Schedule
 - GitHub project
 - Class FAQs
- [Project specs](#)

Yours Truly

- **GP Saggese**
 - 2001-2006, PhD / Postdoc at the University of Illinois at Urbana-Champaign
 - [LinkedIn](#)
 - gsaggese@umd.edu
- **University of Maryland:**
 - 2023-, Lecturer for UMD DATA605: Big Data Systems
 - 2025-, Lecturer for UMD MSML610: Advanced Machine Learning
- **In the real-world**
 - Research scientist at NVIDIA, Synopsys, Teza, Engineers' Gate
 - 3x AI and fin-tech startup founder (ZeroSoft, June, Causify AI)
 - 20+ academic papers, 2 US patents



- Class Intro
- *Big Data*
- Data Models

Data Science

- **Promises of data science**

- Give a competitive advantages
- Make better strategic and tactical business decisions
- Optimize business processes

- **Data science is not new**, it was called:

- Operation research (~1970-80s)
- Decision support, business intelligence (~1990s)
- Predictive analytics (Early 2010s)
- ...

- **What has changed**

- Now learning and applying data science is **easy**
 - No need for hiring a consulting company
- Tools are open-source
 - E.g., Python + pydata stack (numpy, scipy, Pandas, sklearn)
- Large data sets available
- Cheap computing
 - E.g., cloud computing (AWS, Google Cloud), GPUs



Motivation: Data Overload

- “*Data science is the number one catalyst for economic growth*” (McKinsey, 2013)
- **Explosion of data in every domain**
 - Sensing devices/networks monitor processes 24/7
 - E.g., temperature of your room, your vital signs, pollution in the air
 - Sophisticated smart-phones
 - 80% of the world population has a smart-phone
 - Internet and social networks make it easy to publish data
 - Internet of Things (IoT): everything is connected to the internet
 - E.g., power supply, toasters
 - Datafication turns all aspects of life into data
 - E.g., what you like/enjoy turned into a stream of your “likes”
- **Challenges**
 - How to handle the increasing amount data?
 - How to extract actionable insights and scientific knowledge from data?

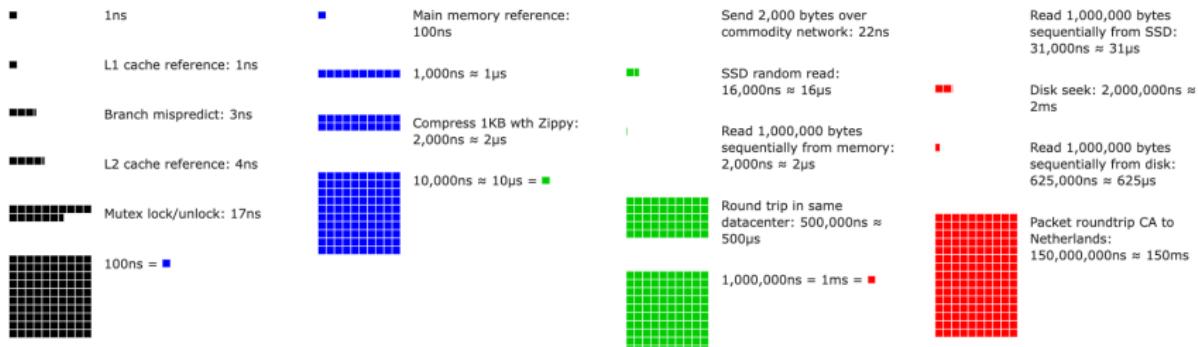
Scale of Data Size

- **Megabyte** = $2^{20} \approx 10^6$ bytes
 - Typical English book
- **Gigabyte** = 10^9 bytes = 1,000 MB
 - 1/2 hour of video
 - Wikipedia (compressed, no media) is 22GB
- **Terabyte** = 1 million MB
 - Human genome: ~1 TB
 - 100,000 photos
 - \$50 for 1TB HDD, \$23/mo on AWS S3
- **Petabyte** = 1000 TB
 - 13 years of HD video
 - \$250k/year on AWS S3
- **Exabyte** = 1M TB
 - Global yearly Internet traffic in 2004
- **Zetabyte** = 1B TB = 10^{21} bytes
 - Global yearly Internet traffic in 2016
 - Fill 20% of Manhattan, New York with data centers
- **Yottabytes** = 10^{24} bytes
 - Yottabyte costs \$100T
 - Fill Delaware and Rhode Island with a million data centers
- **Brontobytes** = 10^{27} bytes



Constants Everybody Should Know

- CPU at 3GHz: 0.3 ns per instruction
- L1 cache reference/register: 1 ns
- L2 cache reference: 4 ns
- Main memory reference: 100 ns
- Read 1MB from memory: 20-100 us
- SSD random read: 16 us
- Send 1KB over network: 1 ms
- Disk seek: 2 ms
- Packet round-trip CA to Netherlands: 150 ms



Big Data Applications

- **Personalized marketing**
- Target each consumer individually
 - E.g., Amazon personalizes suggestions using:
 - Shopping history
 - Search, click, browse activity
 - Other consumers and trends
 - Reviews (NLP and sentiment analysis)
- Brands understand customer-product relationships
 - Use sentiment analysis from:
 - Social media, online reviews, blogs, surveys
 - Positive, negative, neutral sentiment
- E.g.,
 - In 2022, \$600B spent on digital marketing

Big Data Applications

- **Mobile advertisement**
- Mobile phones are ubiquitous
 - 80% of world population has one
 - 6.5 billion smartphones
- Integrate online and offline databases,
e.g.,
 - GPS location
 - Search history
 - Credit card transactions
- E.g.,
 - You've bought a new house
 - You google questions about house renovations
 - You watch shows about renovations
 - Your phone tracks where you are
 - Google sends you coupons for the closest Home Depot
 - *"I feel like Google is following me"*



Big Data Applications

- **Biomedical data**
- Personalized medicine
 - Patients receive treatment tailored to them for efficacy
 - Genetics
 - Daily activities
 - Environment
 - Habits
- Genome sequencing
- Health tech
 - Personal health trackers (e.g., smart rings, phones)



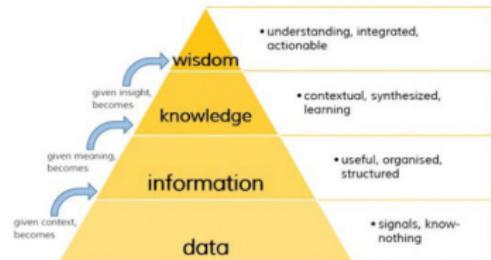
Big Data Applications

- Smart cities
- Interconnected mesh of sensors
 - E.g., traffic sensors, camera networks, satellites
- Goals:
 - Monitor air pollution
 - Minimize traffic congestion
 - Optimal urban services
 - Maximize energy savings



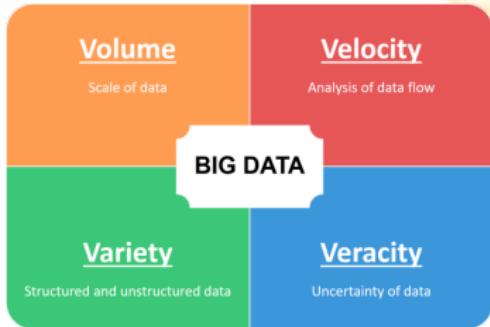
Goal of Data Science

- **Goal:** from data to wisdom
 - Data (raw bytes)
 - Information (organized, structured)
 - Knowledge (learning)
 - Wisdom (understanding)
- Insights enable decisions and actions
- Combine streams of big data to generate new data
 - New data can be “big data” itself



The Six V'S of Big Data

- **Volume**
 - Vast amount of data is generated
- **Variety**
 - Different forms
- **Velocity**
 - Speed of data generation
- **Veracity**
 - Biases, noise, abnormality in data
 - Uncertainty, trustworthiness
- **Valence**
 - Connectedness of data in the form of graphs
- **Value**
 - Data must be valuable
 - Benefit an organization



The Six V's of Big Data

- **Volume**

- Exponentially increasing data
- 2.5 exabytes (1m TB) generated daily
 - 90% of data generated in last 2 years
 - Data doubles every 1.2 years
- Twitter/X: 500M tweets/day (2022)
- Google: 8.5B queries/day (2022)
- Meta: 4PB data/day (2022)
- Walmart: 2.5PB unstructured data/hour (2022)

- **Variety**

- Different data forms
 - Structured (e.g., spreadsheets, relational data)
 - Semi-structured (e.g., text, sales receipts, class notes)
 - Unstructured (e.g., photos, videos)
- Different formats (e.g., binary, CSV, XML, JSON)



The Six V's of Big Data

- **Velocity**

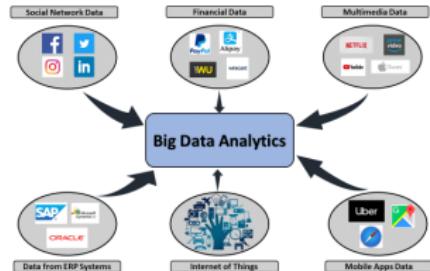
- Speed of data generation
 - E.g., sensors generate data streams
- Process data off-line or in real-time
- Real-time analytics: consume data as fast as generated

- **Veracity**

- Relates to data quality
- How to remove noise and bad data?
- How to fill in missing values?
- What is an outlier?
- How do you decide what data to trust?

Sources of Big Data

- Distinguish Big Data by source
 - **Machines**
 - **People**
 - **Organizations**



Sources of Big Data: Machines

- **Machines** generate data
 - Real-time sensors (e.g., sensors on Boeing 787)
 - Cars
 - Website tracking
 - Personal health trackers
 - Scientific experiments
- **Pros**
 - Highly structured
- **Cons**
 - Difficult to move, computed in-place or centralized
 - Streaming, not batch

Sources of Big Data: People

- **People** and their activities generate data
 - Social media (Instagram, Twitter, LinkedIn)
 - Video sharing (YouTube, TikTok)
 - Blogging, website comments
 - Internet searches
 - Text messages (SMS, Whatsapp, Signal, Telegram)
 - Personal documents (Google Docs, emails)
 - **Pros**
 - Enable personalization
 - Valuable for business intelligence
 - **Cons**
 - Semi-structured or unstructured data
 - Text, images, movies
 - Requires investment to extract value
 - Acquire → Store → Clean → Retrieve → Process → Insights
- SCIENCE ACADEMY Surveillance capitalism



Sources of Big Data: Orgs

- **Organizations** generate data

- Commercial transactions
- Credit cards
- E-commerce
- Banking
- Medical records
- Website clicks

- **Pros**

- Highly structured

- **Cons**

- Store every event to predict future
 - Miss opportunities
- Stored in “data silos” with different models
 - Each department has own system
 - Additional complexity
 - Data outdated/not visible
 - Cloud computing helps (e.g., data lakes, data warehouses)

Is Data Science Just Hype?

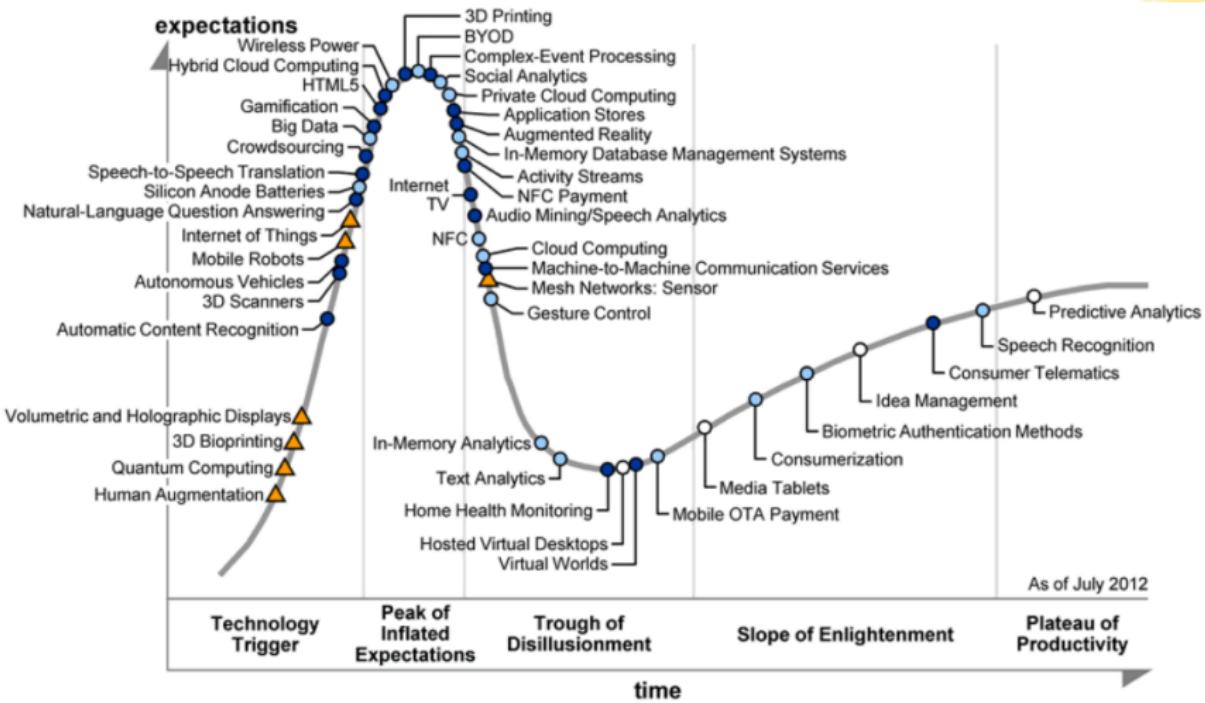
- Big data (or data science)
 - “Any process where interesting information is inferred from data”
- Data scientist called the “sexiest job” of the 21st century
 - The term has become very muddled at this point
- **Is it all hype?**



Is Data Science Just Hype?

- **No**
 - Extract insights and knowledge from data
 - Big data techniques revolutionize many domains
 - E.g., education, food supply, disease epidemics
- **But**
 - Similar to what statisticians have done for years
- **What is different?**
 - More data is digitally available
 - Easy-to-use programming frameworks (e.g., Hadoop) simplify analysis
 - Cloud computing (e.g., AWS) reduces costs
 - Large-scale data + simple algorithms often outperform small data + complex algorithms

What Was Cool in 2012?



Plateau will be reached in:

○ less than 2 years ● 2 to 5 years ● 5 to 10 years

▲ more than 10 years ✗ obsolete ✗ before plateau

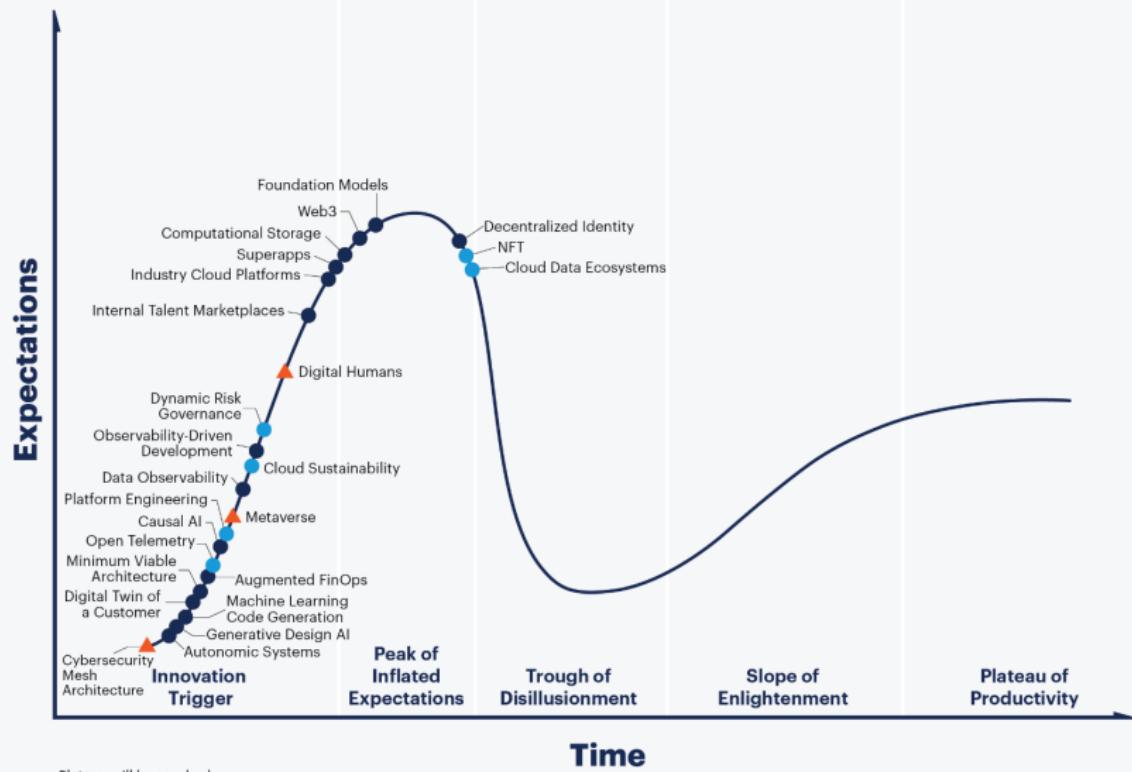


What Was Cool in 2017?



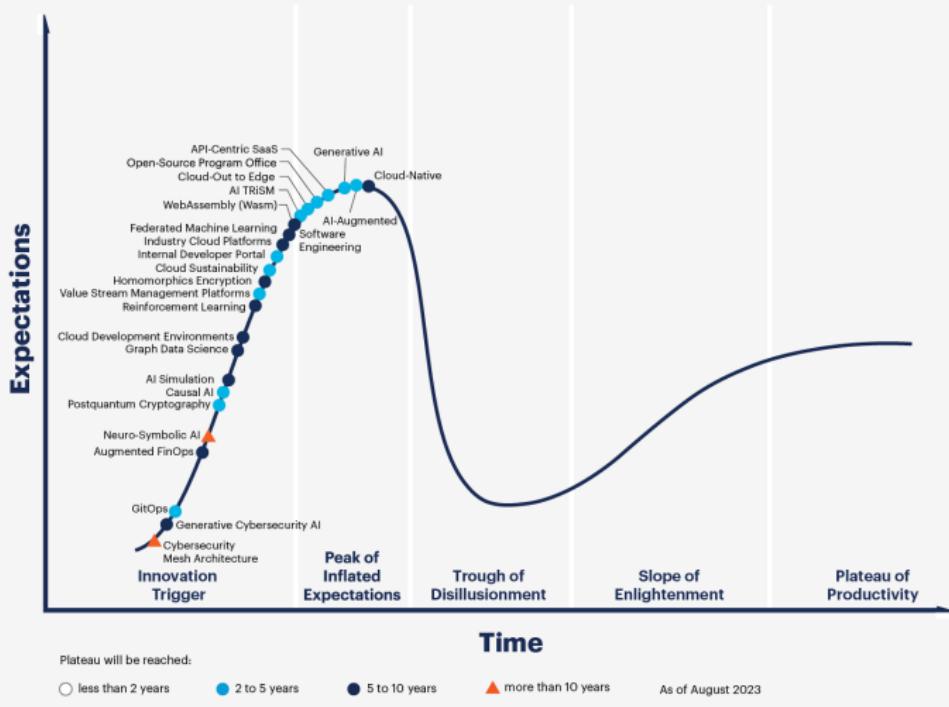
What Was Cool in 2022?

Hype Cycle for Emerging Tech, 2022



What Was Cool in 2023?

Hype Cycle for Emerging Technologies, 2023



Key Shifts Before/After Big-Data

- **Datasets: small, curated, clean → large, uncurated, messy**
 - Before:
 - Statistics based on small, carefully collected random samples
 - Costly and careful planning for experiments
 - Hard to do fine-grained analysis
 - Today:
 - Easily collect huge data volumes
 - Feed into algorithms
 - Strong signal overcomes noise
- **Causation → Correlation**
 - Goal: determine cause and effect
 - Causation hard to determine → focus on correlation
 - Correlation is often sufficient
 - E.g., diapers and beer bought together
- **"Data-fication"**
 - = converting abstract concepts into data
 - E.g., "sitting posture" data-fied by sensors in your seat
 - Preferences data-fied into likes
- From: Rise of Big Data, 2013



Examples: Election Prediction

- Nate Silver and the 2012 Elections
 - Predicted 49/50 states in 2008 US elections
 - Predicted 50/50 states in 2012 US elections
- Reasons for accuracy
 - Multiple data sources
 - Historical accuracy incorporation
 - Statistical models
 - Understanding correlations
 - Monte-Carlo simulations for electoral probabilities
 - Focus on probabilities
 - Effective communication

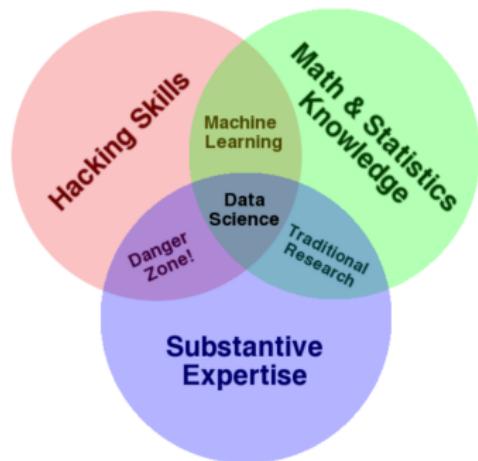


Examples: Google Flu Trends

- 5% to 20% of US population contracts flu yearly; 40k deaths
- Early warnings enable prevention and control
- Google Flu Trends
 - Early flu outbreak warnings via search query analysis
 - 45 search terms analyzed
 - IP used to determine location
 - Predict regional flu outbreaks 1-2 weeks before CDC
 - Active from 2008 to 2015
- Caveat: accuracy declined
 - Claimed 97% accuracy
 - Out of sample accuracy lower (overshot CDC data by 30%)
 - People search about flu without knowing diagnosis
 - E.g., searching for “fever” and “cough”
 - Google Flu Trends: The Limits of Big Data

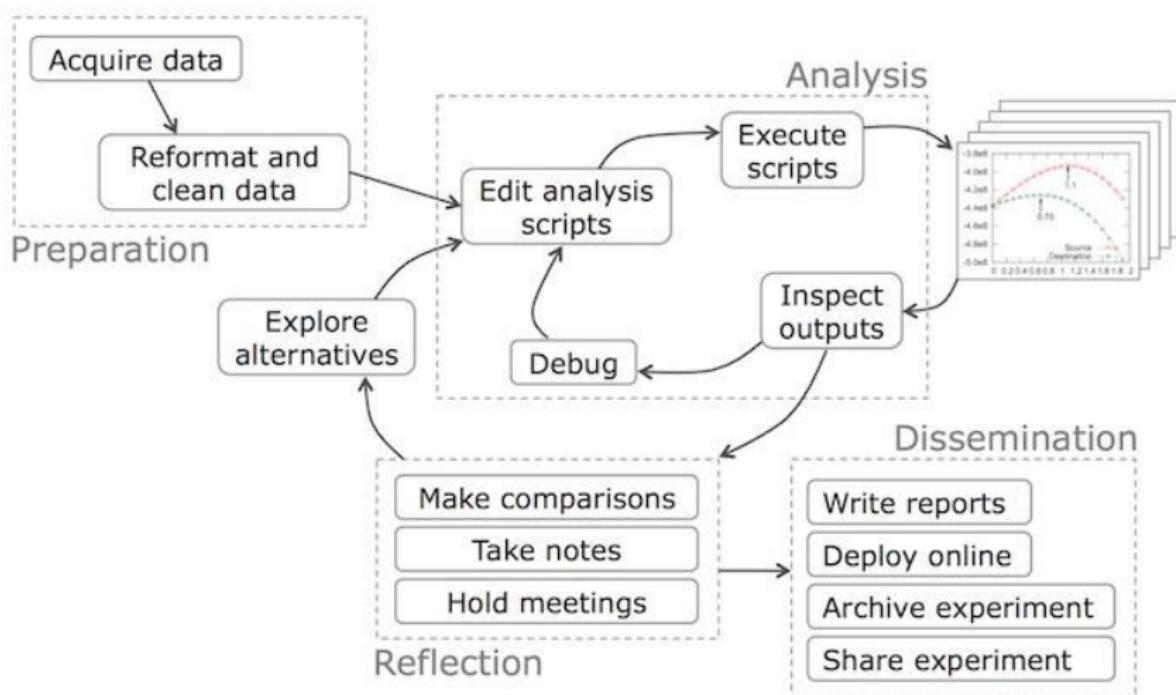
Data Scientist

- Ambiguous, ill-defined term
- From Drew Conway's Venn Diagram



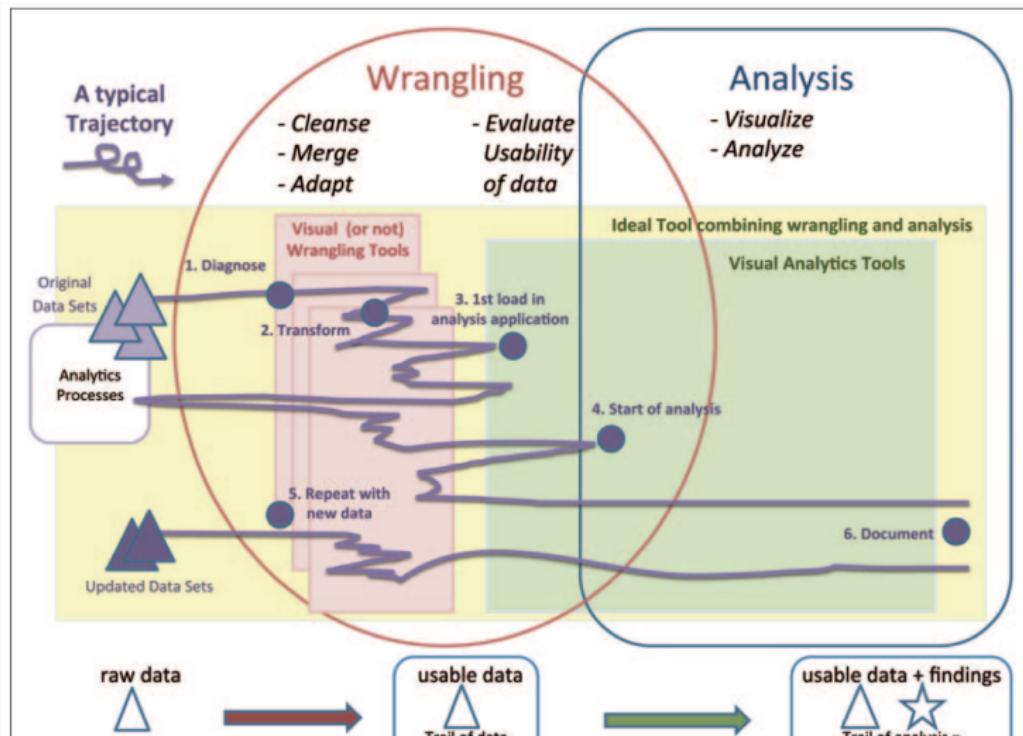
Typical Data Scientist Workflow

- From Data Science Workflow



Where Data Scientist Spends Most Time

- 80-90% of the work is data cleaning and wrangling
- “Janitor Work” in Data Science
- Research Directions in Data Wrangling



What a Data Scientist Should Know

- From: How to hire a data scientist
- **Data grappling skills**
 - Move and manipulate data with programming
 - Scripting languages (e.g., Python)
 - Data storage tools: relational databases, key-value stores
 - Programming frameworks: SQL, Hadoop, Spark
- **Data visualization experience**
 - Draw informative data visuals
 - Tools: D3.js, plotting libraries
 - Know what to draw
- **Knowledge of statistics**
 - Error-bars, confidence intervals
 - Python libraries, Matlab, R
- **Experience with forecasting and prediction**
 - Basic machine learning techniques
- **Communication skills**
 - Tell the story, communicate findings



- Class Intro
- Big Data
- *Data Models*

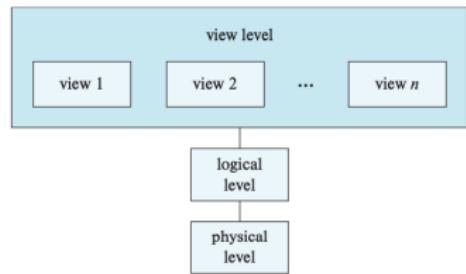
Data Models

- **Data modeling**
 - Represents and captures structure and properties of real-world entities
 - Abstraction: real-world → **representation**
- **Data model**
 - Describes how data is *represented* (e.g., relational, key-value) and *accessed* (e.g., insert operations, query)
 - Schema in a DB describes a specific data collection using a data model
- **Why need data model?**
 - Know data structure to write general-purpose code
 - Share data across programs, organizations, systems
 - Integrate information from multiple sources
 - Preprocess data for efficient access (e.g., building an index)



Multiple Layers of Data Modeling

- **Physical layer**
 - How is the data physically stored
 - How to represent complex data structures (e.g., B-trees for indexing)
- **Logical layer**
 - Entities
 - Attributes
 - Type of information stored
 - Relationships among the above
- **Views**
 - Restrict information flow
 - Security and/or ease-of-use



Data Models: Logical Layer

- **Modeling constructs**

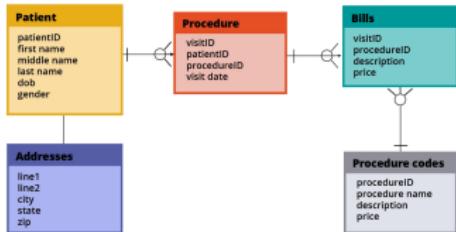
- Concepts to represent data structure
- E.g.,
 - Entity types
 - Entity attributes
 - Relationships between entities
 - Relationships between attributes

- **Integrity constraints**

- Ensure data integrity
 - Avoid errors and inconsistencies
 - E.g., field can't be empty, must be an integer

- **Manipulation constructs**

- E.g., insert, update, delete data



Examples of Data Models

- We will cover:
 - Relational model (SQL)
 - Entity-relationship (ER) model
 - XML
 - Object-oriented (OO)
 - Object-relational
 - RDF
 - Property graph
- Serialization formats as data models
 - CSV
 - Parquet
 - JSON
 - Protocol Buffer
 - Avro/Thrift
 - Python Pickle

Good Data Models

- Data model should be:
 - Expressive
 - Capture real-world data
 - Easy to use
 - Perform well
- Tension between characteristics
 - Powerful models
 - Represent more datasets
 - Harder to use/query
 - Less efficient (e.g., more memory, time)
- Evolution of data modeling tools captures data structure
 - Structured data → Relational DBs
 - Semi-structured web data → XML, JSON
 - Unstructured data → NoSQL DBs



Data Independence

- **Logical data independence**

- Change data representation without altering programs
- E.g., API abstracting backend

- **Physical data independence**

- Change data layout on disk without altering programs
 - Index data
 - Partition/distribute/replicate data
 - Compress data
 - Sort data



Databases: A Brief History

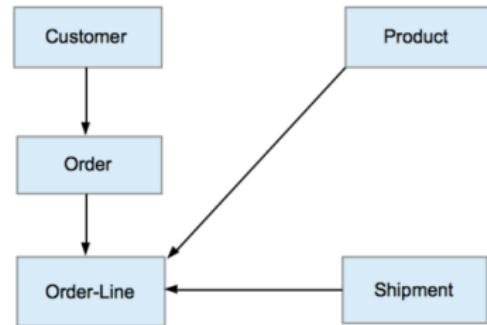
- **1960s: Early beginning**
 - Computers become attractive technology
 - Enterprises adopt computers
 - Applications use own data stores
 - Each application has its own format
 - Data unavailable to other programs
 - **Database:** term for “shared data banks” by multiple applications
 - Define data format
 - Store as “data dictionary” (schema)
 - Implement “database management” software to access data
 - Issues:
 - How to write data dictionaries?
 - How to access data?
 - Who controls the data?
 - E.g., integrity, security, privacy concerns

Databases: A Brief History

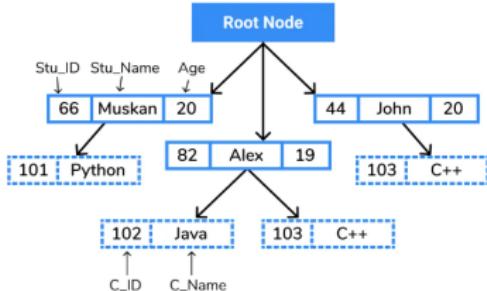
- **1960s, Hierarchical and Network Model**
 - Connect records of different types
 - Example: connect accounts with customers
 - Network model aimed for generality and flexibility
- IBM's IMS Hierarchical Database (1966)
 - Designed for Apollo space program
 - Predates hard disks
 - Used by over 95% of top Fortune 1000 companies
 - Processes 50 billion transactions daily, manages 15 million gigabytes of data
- Cons:
 - Exposed too much internal data (structures/pointers)
 - Leaky abstraction

Relational, Hierarchical, Network Model

- **Relational model**
 - Data as tuples in relations
 - SQL
- **Hierarchical model**
 - Tree-like structure
 - One parent, many children
 - Connected through links
 - XML DBs resurgence in 1990s
- **Network model**
 - Graph organization
 - Multiple parents and children
 - Graph DBs resurgence in 2010s



Customer ID	Tax ID	Name	Address	[More fields...]
1234567890	555-5512222	Ramesh	323 Southern Avenue	...
2223344556	555-5523232	Adam	1200 Main Street	...
3334445563	555-5533323	Shweta	871 Rani Jhansi Road	...
4232342432	555-5325523	Sarfaraz	123 Maulana Azad Sarani	...



Databases: A Brief History

- **1970s: Relational model**

- Set theory, first-order predicate logic
 - Ted Codd developed the Relational Model
- Elegant, formal model
 - Provided data independence
 - Users didn't worry about data storage, processing
- High-level query language
 - SQL based on relational algebra
- Notion of normal forms
 - Reason about data and relations
 - Remove redundancies

- Influential projects:

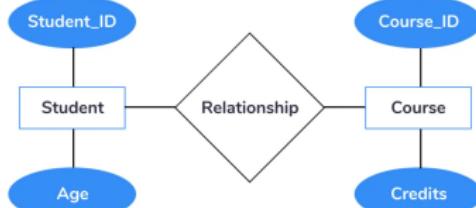
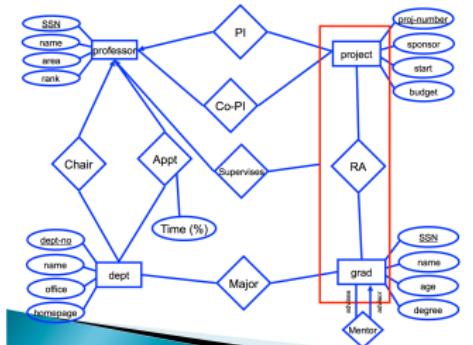
- INGRES (UC Berkeley), System R (IBM)
- Ignored IMS compatibility

- Debates: Relational Model vs Network Model proponents



Entity-Relationship Model

- 1976: Peter Chen proposed “Entity-Relationship Model”
- Describes knowledge as entities and relationships
- **Entities**: Physical or logical objects, “Nouns”
- **Relationships**: Connections between entities, “Verbs”
- Map ER model to relational DB: Entities, relationships → tables



Databases: A Brief History

- **1980s: Relational model acceptance**
 - SQL standard due to IBM's backing
 - Enhanced relational model
 - Set-valued attributes, aggregation
- Late 80's
 - Object-oriented DBs
 - Store objects, not tables
 - Overcome *impedance mismatch* between languages and databases
 - Object-relational DBs
 - User-defined types
 - Combine object-oriented benefits with relational model
 - No expressive difference from pure relational model

Object-Oriented

- OOP is a data model
 - Object behavior described through data (fields) and code (methods)
- **Composition**
 - Aka has-a relationships
 - E.g., Employee class has an Address class
- **Inheritance**
 - Aka is-a relationships
 - E.g., Employee class derives from Person class
- **Polymorphism**
 - Code executed depends on the class of the object
 - One interface, many implementations
 - E.g., draw() method on a Circle vs Square object, both descending from Shape class
- **Encapsulation**
 - E.g., private vs public fields/members
 - Prevents external code from accessing inner workings of an object



Databases: A Brief History

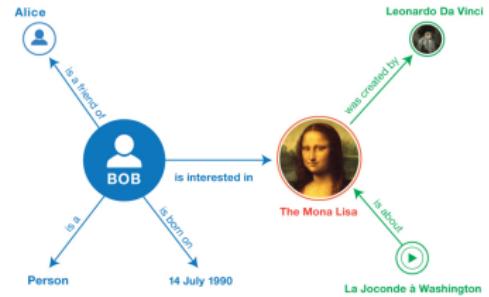
- Late 90's-today
- Web/Internet emerges
- XML: eXtensible Markup Language
 - For *semi-structured* data
 - Tree-like structure
 - Flexible schema

```
<?xml version="1.0" encoding="UTF-8"?>
<CATALOG>
    <CD>
        <TITLE>Empire Burlesque</TITLE>
        <ARTIST>Bob Dylan</ARTIST>
        <COUNTRY>USA</COUNTRY>
        <COMPANY>Columbia</COMPANY>
        <PRICE>10.90</PRICE>
        <YEAR>1985</YEAR>
    </CD>
    <CD>
        <TITLE>Hide your heart</TITLE>
        <ARTIST>Bonnie Tyler</ARTIST>
        <COUNTRY>UK</COUNTRY>
```



Resource Description Framework

- Aka RDF
- Key construct:
“subject-predicate-object” triple
 - Subject=sky
 - Predicate=has-the-color
 - Object=blue
- Maps to a labeled, directed multi-graph
 - More general than a tree
- Stored in:
 - Relational DBs
 - Dedicated “triple-stores” DBs



01 <<http://example.org/bol>>
02 <<http://example.org/bol>>
03 <<http://example.org/bol>>
04 <<http://example.org/bol>>
05 <<http://www.wikidata.org>>
06 <<http://www.wikidata.org>>
07 <<http://data.europeana.eu>>

Property Graph Model

- Graph:
 - Vertices and edges
 - Properties for each edge and vertex
- Stored in:
 - Relational DBs
 - Graph DBs

