



## UMD DATA605 - Big Data Systems


### Git, Data Pipelines

**Instructor:** Dr. GP Saggese - [gsaggese@umd.edu](mailto:gsaggese@umd.edu)

# Bash / Linux: Resources

---

- Command line
  - <https://ubuntu.com/tutorials/command-line-fundamentals>
- E.g., find, xargs, chmod, chown, symbolic, and hard links
- How Linux works
  - Processes
  - File ownership and permissions
  - Virtual memory
  - How to administer a Linux box as root
- Mastery
  - <https://linuxcommand.org/tlcl.php>
  -



# Git Resources

---

# Version Control Systems

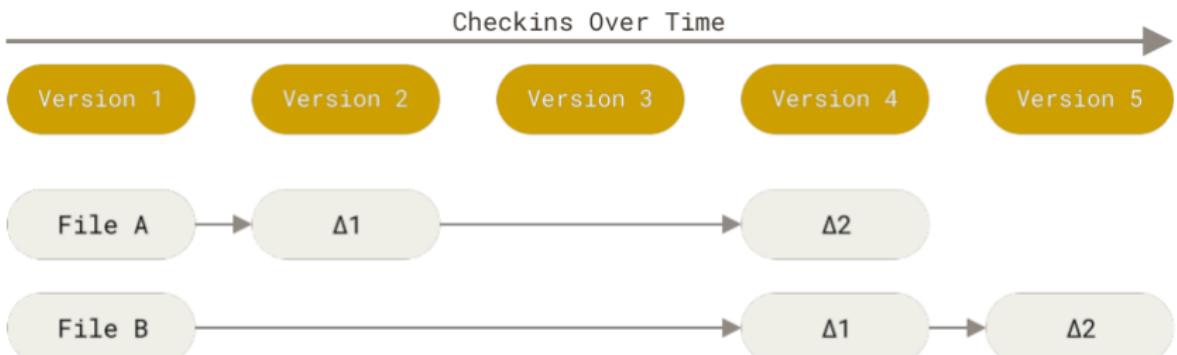
---

- A Version Control System (VCS) is a system that allows to:
    - Record changes to files
    - Recall specific versions later (like a “file time-machine”)
    - Compare changes to files over time
    - Track *who* changed *what* and *when* and *why*
  - **Simplest "VCS"**
    - Make a copy of a dir and add \_v1 (bad) or add a timestamp \_20220101 (better)
    - Make a copy of a dir and add \_v1 (bad) or add a timestamp \_20220101 (better)
    - **Cons:** It kind of works for one person, but doesn't scale
  - **Centralized VCS**
    - E.g., Perforce, Subversion
    - A server stores the code, clients connect to it
    - **Cons:** If the server is down, nobody can work
  - **Distributed VCS**
    - E.g., Git, Mercurial, Bazaar, Dart
    - Each client has the entire history of the repo locally
    - Each node is both a client and a server
- **Cons:** complex



# VCS: How to Track Data

- Consider a directory with project files inside
- **How do you track changes to the data?**
- Delta-based VCS
  - E.g., Subversion
  - Store the data in terms of patches (changes of files over time)
  - Can reconstruct the state of the repo by applying the patches
- **Stream of snapshots VCS**
  - E.g., Git
  - Store data in terms of snapshots of a filesystem
  - Take a “picture” of what files look like
  - Store reference (hash) to the snapshots
  - Save link to previous identical files



- **Almost everything is local for Git**
  - The history is stored locally in each node
  - Diff-ing is done locally
    - For centralized VCS you need to access the server
  - You can commit to your local copy
    - Upload changes when there is network connection
- **Almost everything is undoable in Git**
  - No data corruption
    - Everything is checksummed
  - Nothing can be lost
  - Disclaimer:
    - As long as you commit (at least locally) or stash
    - As long as you don't precipitate in "git hell"
    - You need to know how to do it
- **Git is a mini key-value store with a VCS built on top**
  - This is actually exactly true
  - Two layers:
    - "porcelain": key-value store for a file-system
    - "plumbing": VCS layer

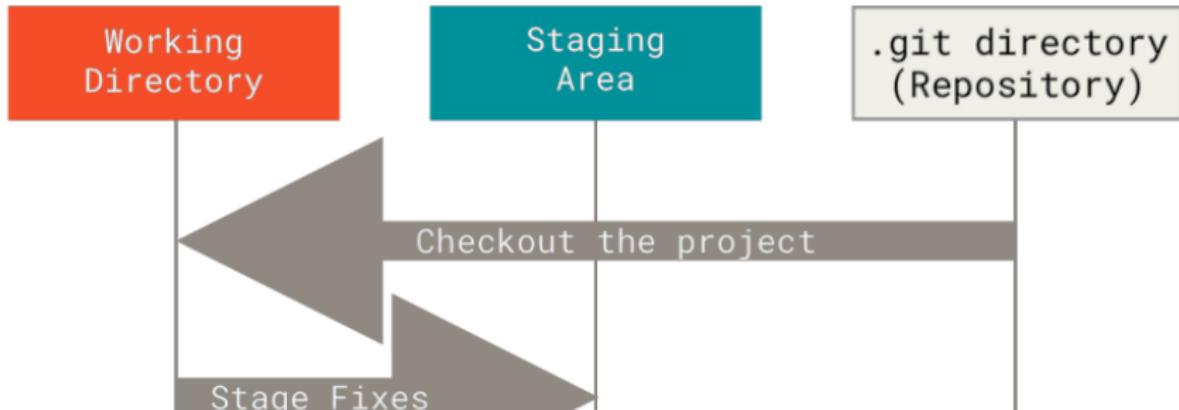
# States of a File in Git

There are 3 main sections of a Git project:

- **Working tree** (aka checkout) - Version of the code placed on the filesystem for the user to use and modify
- **Staging area** (aka cache, index) - A file in .git that stores information about the next commit
- **Git directory** (aka .git) - Store metadata and objects (like a DB)

It is the repo itself with all history - When you clone, you get the .git of the project Each file can be in 4 states from Git point-of-view -

**Untracked**: files that are not under Git version control - **Modified**: you have changed the file, but not committed to DB yet - **Staged**: you have marked a modified file to go into your next commit snapshot - **Committed**: data is safely stored in your local DB



# Git Tutorial

---

- Git tutorial on class repo
  - Read me
- How to use a tutorial
  - Type the commands from the tutorial one-by-one
    - Do not copy paste
  - Look at the results
    - Understand what each line means
  - Play with things
    - “What happens if I do this?”
    - “Does the result match my mental model?”
  - Learn command line before switching to the GUI
    - GUIs hide details and you become dependent on it
- Use one or more tutorials on-line
  - Ideally go through the Git book and the examples
- Build your own cheat sheet
  - Reusing other people cheat sheet works only if you already know
- Strive to achieve mastery of the basic tools of the trade
  - Bash, Git, editor
  - Python
  - Pandas

# Git: Daily Use

---

- Check out a project (git clone) or start from scratch (git init)
  - Only once per Git project client
- **Daily routine**
  - Modify files in working tree (vi ...)
  - Add files (git add ...)
  - Stage changes for the next commit (git add -u ...)
  - Commit changes to .git (git commit)
- **Use a branch to group commits together**
  - Isolate your code from changes in master
  - Merge master into your branch
  - Isolate master from your changes
  - Pull Request (aka PR) to get the code reviewed
  - Merge PR into upstream

# Git Remote

---

- Remote repos are versions of the project hosted on Internet/remote file system
  - To collaborate you need to manage remote repos
  - Push / pull changes
- You can have multiple forks of the same repo with different policies
  - E.g., read-only, read-write
- `git remote -v`: show what are the remotes
- `git fetch`: pull from the remote repo all the data (e.g., branches, commits) that you don't have locally
- `git pull`: short hand from `git fetch origin + git merge master --rebase`
- `git push <REMOTE> <BRANCH>`: push local data to a remote
  - E.g., `git push origin master`
- If somebody pushed to the remote, you can't push your changes right away, but you need to:
  - Fetch the changes
  - Merge changes to the branch in your client
  - Resolve conflicts, if needed
  - (Test project sanity, e.g., by running unit tests)
  - Push changes to the remote

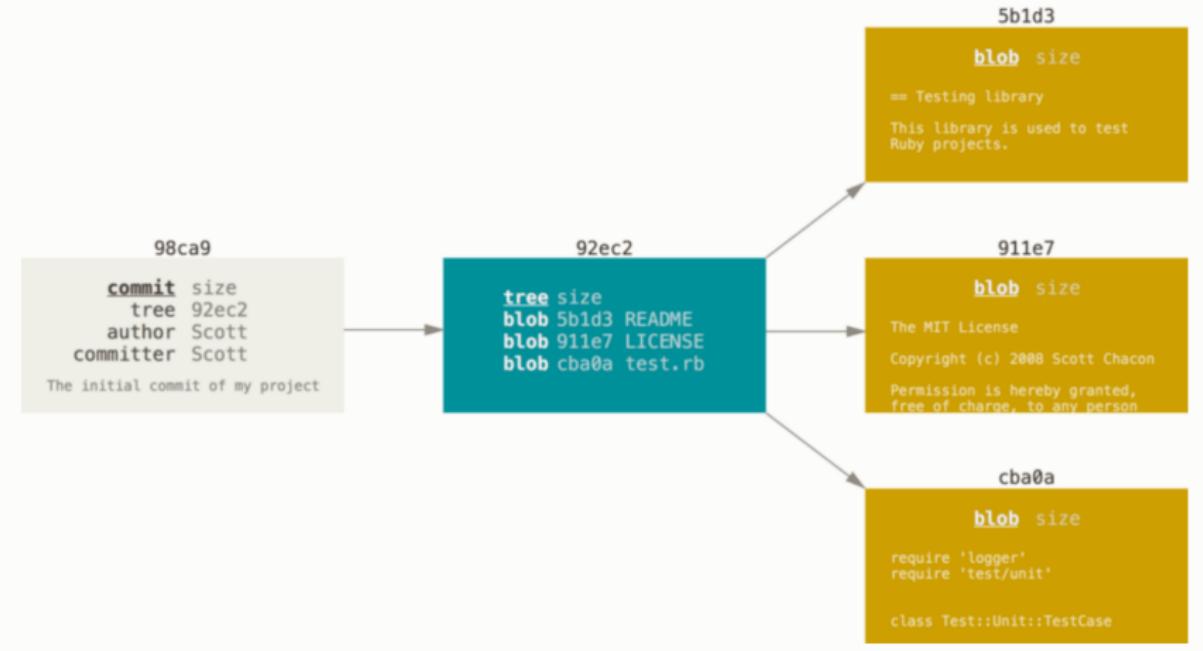
# Git Tagging

---

- Git allows to mark specific points in history with a tag
  - E.g., release points
- You can check out a tag
- You get in detached HEAD state
  - If you commit your change won't be added to the tag or to the branch
  - The commit will be "unreachable", i.e., reachable only by the commit hash



# Git Internals



- You can understand Git only if you understand the data model

- Git is a key-value store with a VCS user interface on top of it
- Key = hash of a file
- Value = content of a file

# Git Branching

---

- **Branching**
  - = diverging from the main line of development
- **Why branch?**
  - Work without messing the rest of the code
  - Work without being affected by changes in the main branch
  - Merge the code downstream to pick up the changes
  - Once you are done working in your branch, merge code upstream
- **Git branching is lightweight**
  - It's instantaneous
  - A branch is just a pointer to a commit
  - Git doesn't store data as difference of files, but as a series of snapshot
- **Git workflows branch and merge often**
  - Even multiple times a day
  - This is surprising for people used to distributed VCS
    - E.g., you would start branching before going for lunch
  - Branches are cheap in Git
  - Use them all the times to isolate your work and organize it

# Git Branching

- **master** (or **main**) is just a normal branch
  - master is a pointer to the last commit
  - As you commit the pointer moves forward
- **HEAD**
  - Pointer to the local branch you are on
  - E.g., master, testing
  - `git checkout <BRANCH>` moves you across branches
- `git branch testing`
  - Create a new pointer testing
  - Point to the commit you are on
  - Pointer can be moved around
- Divergent history
  - When work progresses in two branches that are “split”




# Git Checkout

- git checkout switches branch
  - Move HEAD pointer to the new branch
  - Change the files in the working dir to match the state corresponding to the branch pointer
- E.g., there are two branches master and testing
  - You are on master
  - git checkout testing
  - The pointer moves, the working dir is changed (not really in this case)
  - Then you can keep working by committing on **testing**
  - The pointer to testing moves forward (no divergent history)



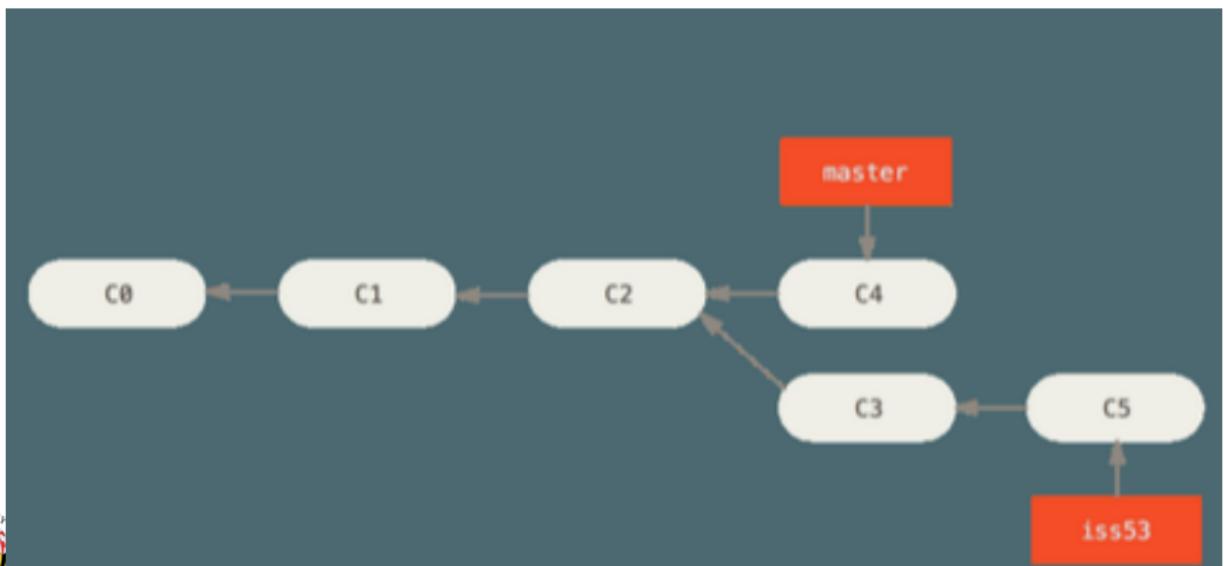
# Git Branching and Merging

-Tutorials: Work on main, Hot fix Start from a project with some commits  
Branch to work on a new feature “Issue 53” > git checkout -b iss53  
**work ... work ... work** > git commit Need a hotfix to master > git  
checkout master > git checkout -b hotfix fix ... fix ... fix > git  
commit -am "Hot fix" > git checkout master > git merge hotfix  
**Fast forward** Now there is a divergent history between master and iss53



# Git Branching and Merging

> git checkout iss53 **work ... work ... work** The branch keeps diverging  
At some point you are done with **iss53** You want to merge your work back to **master** Go to the target branch > git checkout master > git merge iss53 Git can't fast forward Git creates a new snapshot with the 3-way "merge commit" (commit with more than one parent) Delete the branch > git branch -d iss53



# Fast Forward Merge

- **Fast forward merge** = merge a commit X with a commit Y that can be reached by following the history of commit X\*\*
- There is not divergent history to merge
- Git simply moves the branch pointer forward from X to Y
- **Mental model:** a branch is just a pointer that says where the tip of the branch is\*\*
- E.g., C4' is reachable from C3 > git checkout master > git merge experiment
- Git moves the pointer of master to C4'

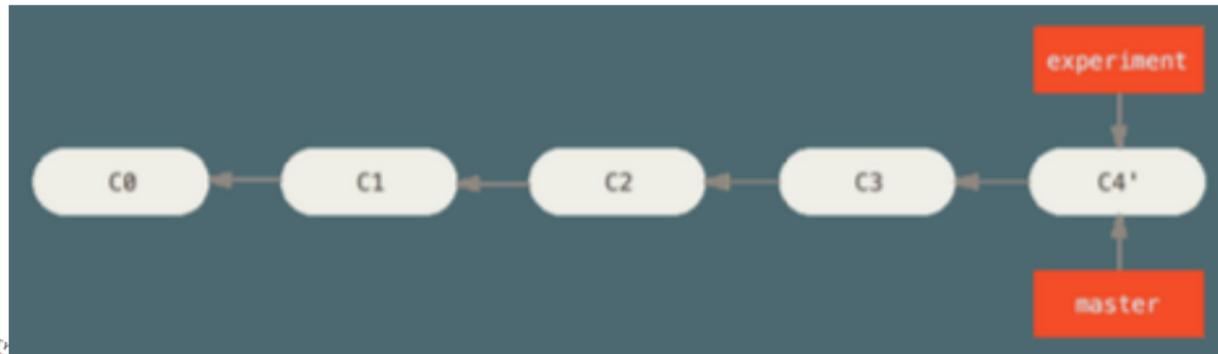


Figure 1: alt\_text

# Merging Conflicts

- Tutorial:
  - Merging conflicts
- Sometimes Git can't merge, e.g.,
  - The same file has been modified by both branches
  - One file was modified by one branch and deleted by another
- Git:
  - Does not create a merge commit
  - Pauses to let you resolve the conflict
  - Adds conflict resolution markers
- User merges manually
  - Edit the files git mergetool
  - git add to mark as resolved\*\*
  - git commit
  - Use PyCharm

```
<<<<< HEAD:index.html
<div id="footer">contact : email.support@github.com</div>
=====
<div id="footer">
    please contact us at support@github.com

```




# Git Rebasing

---

- In Git there are two ways of merging divergent history
  - E.g., **master** and **experiment** have a common ancestor C2
- **Merge**
  - Go to the target branch
  - > git checkout master
  - > git merge experiment
  - Create a new snapshot C5 and commit
- **Rebase**
  - Go to the branch to rebase
  - > git checkout experiment
  - > git rebase master
  - Rebase algo:
    - Get all the changes committed in the branch (C4) where we are on (experiment) since the common ancestor (C2)
    - Sync to the branch that we are rebasing onto (master at C3)
    - Apply the changes C4'
    - Only the branch where we are is affected
  - Finally fast forward master

# Uses of Rebase

- **Rebasing makes for a cleaner history**
  - The history looks like all the work happened in series
  - Although in reality it happened in parallel to the development in master
- **Rebasing to contribute to a project**
  - Developer
    - You are contributing to a project that you don't maintain
    - You work on your branch
    - When you are ready to integrate your work, rebase your work onto **origin/master**
  - The maintainer
    - Does not have to do any integration work
    - Does just a fast forward or a clean apply (no conflicts)




# Golden Rule of Rebasing

- **Remember:** rebasing means abandoning existing commits and creating new ones that are similar but different
- **Problem**
  - You push your commits to a remote somewhere
  - Others pull your commits down and base their work on them
  - You rewrite those commits with `git rebase`
  - You push them again with `git push --force`
  - Your collaborators have to re-merge their work
- **Solution Strict version:** “Do not **ever** rebase commits that exist outside your repository” **Loose version:** “It’s ok to rebase your branch even if you pushed to a server, as long as you are the only one to use it”



# Rebase vs Merge: Philosophical Considerations

- Rebase-vs-merge depend on the answer to the question:
- **What does the commit history of a repo mean?**
- a) History is the record of what actually happened
- \*“History should not be tampered with!”\*\*
- **Q: What if there is a series of messy merge commits?**
- **A: This is how it happened. The repo should preserve this**
- Use git merge
- **b) History represents how a project should have been made**
- “*You would not publish a book as a sequence of drafts and correction, but rather the final version*”
- *You should tell the history in the way that is best for future readers*
- Use git rebase and filter-branch\*\*



# Rebase vs Merge: Philosophical Considerations

---

- Many man-centuries have been wasted discussing rebase-vs-merge at the watercooler
  - Total waste of time! Tell people to get back to work!
- When you contribute to a project often people decide for you based on their preference
- **Best of the merge-vs-rebase approaches**
- Rebase changes you've made in your local repo
  - Even if you have pushed but you know the branch is yours
  - Use `git pull --rebase` to clean up the history of your work
  - If the branch is shared with others then you need to definitively `git merge`
- Only `git merge` to master to preserve the history of how something was built
- **Personally**
- I like squash-and-merge branches to master
  - My commits are just my checkpoints
  - Rarely they are “complete”



experiment

64

# Remote Branches

---

- Remote branches are pointers to branches in remote repos

```
git remote -v
origin  git@github.com:gpsaggese/umd_data605.git (fetch)
origin  git@github.com:gpsaggese/umd_data605.git (push)
```

## • Tracking branches

- Local references representing the state of the remote repo
- E.g., master tracks **origin/master**
- You can't change the remote branch (e.g., **origin/master**)
- You can change tracking branch (e.g., **master**)
- Git updates tracking branches when you do `git fetch origin` (or `git pull`)
- To share code in a local branch you need to push it to a remote
  - > `git push origin serverfix`
- To work on it
  - > `git checkout -b serverfix origin/serverfix`

# Git Workflows

**Git workflows** - = ways of working and collaborating using Git

**Long-running branches** - Branches at different level of stabilities, that are always open -

**master** is always ready to be released - **develop** branch to develop in - *topic / feature* branches - When branches are “stable enough” they are merged up

**Topic branches** - Short-lived branches for a single feature - E.g., **hotfix**, **wip-XYZ** - Easy to review - Silo-ed from the rest - This is typical of Git since other VCS support for branches is not good enough - E.g., - You start **iss91**, then you cancel some stuff, and go to **iss91v2** - Somebody starts **dumbidea** branch and merge to **master** (!) - **You squash-and-merge your iss91v2\*\***




# Centralized Workflow

**Centralized workflow in centralized VCS** - Developers:

- Check out the code from the central repo on their computer
- Modify the code locally
- Push it back to the central hub (assuming no conflicts with latest copy, otherwise they need to merge)

**Centralized workflow in Git** - Developers:

- Have push (i.e., write) access to the central repo
- Need to fetch and then merge
- Cannot push code that will overwrite each other code (only fast-forward changes)



# Forking Workflows

---

- Typically devs don't have permissions to update directly branches on a project
  - Read-write permissions for core contributors
  - Read-only for anybody else
- **Solution**
  - "Forking" a repo
  - External contributors
    - Clone the repo and create a branch with the work
    - Create a writable fork of the project
    - Push branches to fork
    - Prepare a PR with their work
  - Project maintainer
    - Reviews PRs
    - Accepts PRs
    - Integrates PRs
  - In practice it's the project maintainer that pulls the code when it's ready, instead of external contributors pushing the code
- **Aka "GitHub workflow"**
  - "Innovation" was forking (Fork me on GitHub!)
  - GitHub acquired by Microsoft for 7.5b USD

# Integration-Manager Workflow

---

- This is the classical model for open-source development
  - E.g., Linux, GitHub (forking) workflow
- **1. One repo is the "official" project**
  - Only the project maintainer pushes to the public repo
  - E.g., sorrentum/sorrentum
- **2. Each contributor**
  - Has read access to everyone else's public repo
  - Forks the project into a private copy
    - Write access to their own public repo
    - E.g., gpsaggese/sorrentum
  - Makes changes
  - Pushes changes to his own public copy
  - Sends email to maintainer asking to pull changes (pull request)
- **3. The maintainer**
  - Adds contributor repo as a remote
  - Merges the changes into a local branch
  - Tests changes locally
  - Pushes branch to the official repo



blessed  
repository

developer  
public

developer  
public

# Git log

- git log reports info about commits\*\*
- **refs** are references to:
  - HEAD (commit you are working on, next commit)
  - origin/master (remote branch)
  - experiment (local branch)
  - d921970 (commit)
- after a reference resolves to the parent of that commit
- HEAD = commit before HEAD, i.e., last commit
- 2 means
- A merge commit has multiple parents
- **Double-dot notation**
  - **1..2** = commits that are reachable from 2 but not from 1 (difference)\*\*
  - `git log master..experiment $\to$ D,C`
  - `git log experiment..master $\to$ F,E`
- **Triple-dot notation**
  - **1...2** = commits that are reachable from either branch but not from both (union excluding intersection)\*\*
  - `git log master...experiment $\to$ F,E,D,C`

# Advanced Git

---

- stashing
  - Copy state of your working dir (e.g., modified and staged files), save it in a stack, to apply later
- cherry-picking
  - Rebase for a single commit
- rerere
  - = “Reuse Recorded Resolution”
  - Git caches how to solve certain conflicts
- tagging
  - Give a name to a specific commit (e.g., v1.3)
- submodules / subtrees
  - Project including other Git projects
- bisect
  - Sometimes a bug shows up at top of tree
  - You don't know at which revision it started manifesting
  - You have a script that returns 0 if the project is good and non-0 if the project is bad
  - `git bisect` can find the revision at which the script goes from good to bad\*\*
- filter branch
  - Rewrite repo history in some scriptable way

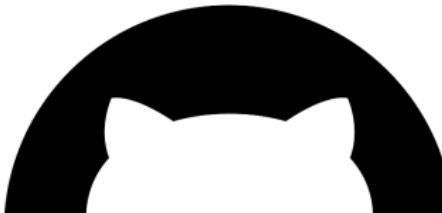


SCIENCE  
ACADEMY

# GitHub

---

- GitHub acquired by MSFT for 7.5b
- **GitHub: largest host for Git repos**
  - Git hosting (100m+ open source projects)
  - PRs, forks
  - Issue tracking
  - Code review
  - Collaboration
  - Wiki
  - Actions (CI / CD)
- **"Forking a project"**
  - In open-source communities
    - It had a negative connotation
    - Take a project, modify it, and make it a competing project
  - In GitHub parlance
    - Make a copy of a project so that you can contribute to it even if you don't have push / write access



# Data Pipelines - Resources

- Concepts in the slides
- Class project
- Mastery
  - Data Pipelines Pocket Reference: Moving and Processing Data for Analytics

O'REILLY®

# Data Pipelines Pocket Reference

Moving and  
Processing Data  
for Analytics




# Data as a Product

---

- Many services today "sell" data
  - Services are typically powered by data and machine learning
  - Data products, e.g.,
    - Personalized search engine (Google)
    - Sentiment analysis on user-generated data (Facebook)
    - A recommendation engine + e-commerce (Amazon)
    - Streaming data (Netflix, Spotify)
- Several steps are required to generate data products
  - Data ingestion
  - Data pre-processing
    - Cleaning, tokenization, feature computation
  - Model training
  - Model deployment
    - MLOps
  - Model monitoring
    - Is model working?
    - Is model getting slower?
    - Are model performance getting worse?
  - Collect feedback from deployment
    - E.g., recommendations vs what users bought
    - Ingest data from production for future versions of the model

# Data Pipelines

- “Data is the new oil”
  - ... but oil needs to be refined
- **Data pipelines**
  - Processes that move and transform data
  - Goal: derive new value from data through analytics, reporting, machine learning\*\*
- Data needs to be:
  - Collected
  - Pre-processed / cleaned
  - Validated
  - Processed
  - Combined
- **Data ingestion**
  - Simplest data pipeline
  - Extract data (e.g., from REST API)
  - Load data into DB (e.g., SQL table)



# Roles in Building Data Pipelines

- **Data engineers**

- Build and maintain data pipelines
- Tools:
  - Python / Java / Go / No-code
  - SQL / NoSQL stores
  - Hadoop / MapReduce / Spark
  - Cloud computing

- **Data scientists**

- Build predictive models
- Tools:
  - Python / R / Julia
  - Hadoop / MapReduce / Spark
  - Cloud computing

- **Data analysts**

- E.g., marketing, MBAs, sales, . . .
- Build metrics and dashboards
- Tools:
  - Excel spreadsheets
  - GUI tools (e.g., Tableaux)
  - Desktop

- **Recurring practical problems**

- Who is responsible for the data?

- Issues with scaling



# Data Ingestion


---

- **Data ingestion**
  - = extract data from one source and load it into another store
- **Data sources / sinks**
  - DBs
    - E.g., Postgres, MongoDB
  - REST API
    - Abstraction layer on top of DBs
  - Network file system / cloud
    - E.g., CSV files, Parquet files
  - Data warehouses
  - Data lakes
- **Source ownership**
  - An organization can use 10-1000s of data sources
  - Internal
    - E.g., DB storing shopping carts for a e-commerce site
  - 3rd-parties
    - E.g., Google analytics tracking website usage

## How Data Ingestion Works

# Data Pipeline Paradigms

- There are several styles of building data pipelines
- Multiple phases
  - Extract
  - Load
  - Transform
- Phases arranged in different ways depending on philosophy about data / roles
  - ETL
  - ELT
  - EtLT



# ETL Paradigm: Phases

---

- **Extract**

- Gather data from various data sources, e.g.,
  - Internal / external data warehouse
  - REST API
  - Data downloading from API
  - Web scraping

- **Transform**

- Raw data is combined and formatted to become useful for analysis step

- **Load**

- Move data into the final destination, e.g.,
  - Data warehouse
  - Data lake

- **Data ingestion pipeline = E + L**

- Move data from one point to another
- Format the data
- Make a copy
- Have different tools to operate on the data



# ETL Paradigm: Example

---

- **Extract**

- Buy-vs-build data ingestion tools
  - Vendor lock-in

- **Transform**


- Data conversion (e.g., parsing timestamp)
- Create new columns from multiple source columns
  - E.g., year, month, day -> yyyy/mm/dd
- Aggregate / filter through business logic
  - Try not to filter, better to mark
- Anonymize data

- **Load**

- Organize data in a format optimized for data analysis
  - E.g., load data in relational DB
- Finally data modeling

# Workflow Orchestration

- Companies have many (10-1000s) data pipelines
- Orchestration tools, e.g.,
  - Apache Airflow (from AirBnB)
  - Luigi (from Spotify)
  - AWS Glue
  - Kubeflow
- Schedule and manage flow of tasks according to their dependencies
  - Pipeline and jobs are represented through DAGs
- Monitor, retry, and send alarms



# ELT paradigm

- ETL has been the standard approach for long time

- Extract → Transform → Load

- **Cons**

- Need to understand the data at ingestion time
- Need to know how the data will be used

- Today ELT is becoming the pattern of choice\*\*


- Extract → Load → Transform

- **Pro:**

- No need to know how the data will be used
- Separate data engineers and data scientists / analysts
- Data engineers focus on data ingestion (E + L)
- Data scientists focus on transform (T)

- **ETL → ELT enabled by new technologies**

- Large storage to save all the raw data (cloud computing)
- Distributed data storage and querying (e.g., HDFS)
- Columnar DBs



# Row-based vs Columnar DBs

---


- **Row-based DBs**
  - E.g., MySQL, Postgres
  - Optimized for reading / writing rows
  - Read / write small amounts of data frequently
- **Columnar DBs**
  - E.g., Amazon Redshift, Snowflake
  - Read / write large amounts of data infrequently
  - Analytics requires a few columns
  - Better data compression

OrderId	CustomerId	ShippingCountry	OrderTotal
1	1258	US	55.25
2	5698	AUS	125.36
3	2265	US	776.95
4	8954	CA	32.16
Block 1	1, 1258, US, 55.25		
Block 2	2, 5698, AUS, 125.36		
Block 3	3, 2265, US, 776.95		
Block 4	4, 8954, CA, 32.16		

- **ETL**
  - Extract → Transform → Load
- **ELT**
  - Extract → Load → Transform
  - Transformation / data modeling (“T”) according to business logic
- **EtLT**
  - Sometimes transformations with limited scope (“t”) are needed
    - De-duplicate records
    - Parse URLs into individual components
    - Obfuscate sensitive data (for legal or security reasons)
  - Then implement rest of “LT” pipeline

# Structure in Data (or Lack Thereof)

- **Structured data:** there is a schema
  - Relational DB
  - CSV
  - DataFrame
  - Parquet
- **Semi-structured:** subsets of data have different schema
  - Logs
  - HTML pages
  - XML
  - Nested JSON
  - NoSQL data
- **Unstructured:** no schema\*\*
  - Text
  - Pictures
  - Movies
  - Blobs of data



# OLAP vs OLTP Workloads

---

- There are two classes of data workloads
- **OLTP**
  - On-Line Transactional Processing
  - Execute large numbers of transactions by a large number of processes in real-time
  - **Lots of concurrent small read / write transactions**
  - E.g., online banking, e-commerce, travel reservations
- **OLAP**
  - On-Line Analytical Processing
  - Perform multi-dimensional analysis on large volumes of data
  - **Few large read or write transactions**
  - E.g., data mining, business intelligence

# OLAP

# Challenges with Data Pipelines

---

- High-volume vs low-volume
  - Lots of small reads / writes
  - A few large reads / writes
- Batch vs streaming
  - Real-time constraints
- API rate limits / throttling
- Connection time-outs
- Slow downloads
- Incremental mode vs catch-up

# Data Warehouse vs Data Lake

- **Data warehouse**

- = DB storing data from different systems in a structured way
- Corresponds to ETL data pipeline style
- E.g., a large Postgres instance with many DBs and tables
- E.g.,
  - AWS Athena, RDS
  - Google BigQuery

- **Data lake**

- = data stored in a semi-structured or without structure
- Corresponds to ELT data pipeline style
- E.g., an AWS S3 bucket storing blog posts, flat files, JSON objects, data605/lectures\_source/images



# Data Lake: Pros and Cons

---

- Data lake = stores data in a semi-structured or without structure
- **Pros**
  - Storing data in cloud storage is cheaper
  - Making changes to types or properties is easier since it's unstructured or semi-structured (with no predefined schema)
    - E.g., JSON documents
  - Data scientists
    - Don't know initially how to access and use the data
    - Want to explore the raw data
- **Cons**
  - It is not optimized for querying like a structured data warehouse
    - There are tools that allow to query data in a data lake similar to SQL
    - E.g., AWS Athena, Redshift Spectrum

# Advantages of Cloud Computing

---

- Ease of building and deploying:
  - Data pipelines
  - Data warehouses
  - Data lakes
- Managed services
  - No need for admin and deploy
  - Highly scalable DBs
    - E.g., Amazon Redshift, Google BigQuery, Snowflake
- Rent-vs-buy
  - Easy to scale up and out
  - Easy to upgrade
  - Better cash-flow
- Cost of storage and compute is continuously dropping
  - Economies of scale
- Cons
  - The flexibility has a cost (2x-3x more expensive than owning)
  - Vendor lock-in