

UMD DATA605 - Big Data Systems

DevOps with Docker

Instructor: Dr. GP Saggese - gsaggese@umd.edu**

TAs: Krishna Pratardan Taduri, kptaduri@umd.edu Prahar
Kaushikbhai Modi, pmodi08@umd.edu

v1.1

Class Announcements

- **Install Docker**
- Assignments on ELMS
- **Do the Git and GitHub tutorial**
- **Quiz about Git next week**
- **Class project**
- Study and experiment with different big data technology
- We will give you different examples
 - Individual project
 - Different complexity

Docker - Resources

- We will use Docker during the class project and most tutorials
- Concepts in the slides
- Class tutorials:
 - `tutorial_docker`
 - `tutorial_docker_compose`
- Web resources:
 - Docker Tutorial for beginners
 - <https://labs.play-with-docker.com/>
 - <https://training.play-with-docker.com>
 - A Beginner-Friendly Introduction to Containers, VMs and Docker
 - Official Docker Getting Started Tutorial
- Mastery:
 - Poulton, Docker Deep Dive: Zero to Docker in a single book, 2020

Application Deployment

- **For (almost all) Internet companies, the application is the business**
 - If the application breaks, the business stops working
 - E.g., Amazon, Google, Facebook, on-line banks, travel sites (e.g., Expedia), . . . , OpenAI
- **Problem**
 - How to release / deploy / manage / monitor applications?
- **Solutions**
 - Before 2000s: “bare-metal era”
 - 2000s-2010s: “virtual machine era”
- ~2013: “container era”

DevOps

- **DevOps** = set of practices that combines:
 - Software development (*dev*)
 - IT operations (*ops*)
- **Containers revolutionized DevOps**
 - Enable true independence between application development and IT ops
 - One team creates an application
 - Another team deploys and manages the applications
 - Create a model for better collaboration (fewer conflicts) and innovation
 - IT: "It doesn't work!"
 - Devs: "What? It works for me"
- Plan
- Code
- Build
- Test
- Release
- Deploy
- Operate
- Monitor

Run on bare metal

- **< 2000s**
 - Running one or few applications on each server (without virtualization)
- **Pros**
 - No virtualization overhead
- **Cons**
 - Not safe / not secure since no separation between applications
 - Expensive
- **Expensive / low efficiency**
 - IT would buy a new server for each application
 - Difficult to spec out the machine → buy “big and fast servers”
 - Overpowered servers operating at 5-10% of capacity
 - Tons of money in the 2000 DotCom boom was spent on machines and networks
- It kind of came back in 2020 but with different use cases in Cloud Computing
- **Winners:** Cisco, Sun, Microsoft

Virtual Machine Era

- **Circa 2000-2010: Virtual Machine**

- Virtual machine technology = run multiple copies of OSES on the same hardware

- **Pros**

- VM runs safely and securely multiple applications on a single server
- IT could run apps on existing servers with spare capacity

- **Cons**

- Every VM requires an OS (waste of CPU, RAM, and disk)
- Buy an OS license
- Monitor and patch each OS
- VMs are slow to boot

- **Winners:** VMWare, RedHat, Citrix

Containers Era

- **Circa 2013: Docker becomes ubiquitous**
- **Docker**
 - Didn't invent containers
 - Made containers simple and mainstream
- Linux supported containers for some time
 - Kernel namespaces
 - Control groups
 - Union filesystems
- **Pros**
 - Containers are fast and portable
 - Containers don't require full-blown OS
 - All containers run on a single host
 - Reduce OS licencing cost
 - Reduce overhead of OS patching and maintenance
- **Cons**
 - CPU overhead
 - Toolchain to learn / use
- **Winners:** AWS, Microsoft Azure, Google (not Docker Inc.)

Serverless Computing

- **Containers run in an OS, OS runs on a host**
 - **Where is the host running?**
 - Local (your laptop)
 - On premise (your own computer in a rack)
 - Cloud instance (e.g., AWS EC2 instance)
 - **How is the host running?**
 - On bare-metal server
 - On a virtual machine
 - On a virtual machine running a virtual machine
- **Serverless computing**
 - As long your application runs somewhere, you don't care "how" or "where"
 - E.g., AWS Lambda

HW vs OS Virtualization

- **Hypervisor performs HW virtualization**
 - Carves out physical hardware resources into VMs
 - Resources (CPUs, RAM, storage) are allocated to a VM
 - It's like having multiple computers
- “Virtual machine tax”
 - To run 3 apps, you need 3 VMs
 - Each VM requires time to start
 - Consumes CPU, RAM, storage
 - VM needs a OS license
 - VM / OS need admins, patching
 - You just want to run 3 apps!
- **Containers perform OS virtualization**
 - It's like having multiple OSes

Docker: Client-Server

- Client-Server architecture
- **Docker client**
 - Command line interface
 - Communicate with the server through IPC socket
 - E.g., `/var/run/docker.sock` or IP port
- **Docker engine**
 - Run and manage containers
 - Modular and built from several OCI-compliant sub-systems
 - E.g., Docker daemon, `containerd`, `runc`, plug-ins for networking and storage

Docker Architecture

- **Docker run-time**
 - **runc**: start and stop containers
 - **containerd**
 - Pull images
 - Create containers, volumes, network interfaces
- **Docker engine**
 - **dockerd**
 - Expose remote API
 - Manage images, volumes, networks
- **Docker orchestration**
 - **docker swarm**
 - Manage clusters of nodes
 - Replaced by Kubernetes
- **Open Container Initiative (OCI)**
 - Standardize low-level components of container infrastructure
 - E.g., image format, run-time API
 - “Death” of Docker

Docker Container

- **Docker Container**
 - Unit of computation
 - Lightweight, stand-alone, executable software package
 - Include everything needed to run
 - E.g., code, runtime / system libraries, settings
 - It is a run-time object
 - vs Docker images are built-time objects
 - Like program running (container) vs program code (image)

Docker Image

- **Docker Image**
 - Unit of deployment
 - Contain everything needed by an app to run
 - Application code
 - Application dependencies
 - Minimal OS support
 - File system
 - Users can
 - Build images from Dockerfiles
 - Pull pre-built images from a registry
 - Multiple layers stacked on top of each other
 - Typically few 100s MBs

Docker Image Layers

- **Docker image** is a configuration file that lists the layers and some metadata
 - It is composed of read-only layers
 - Each layer is independent from each other
 - Each layer comprises of many files
- **Docker driver**
 - Stacks these layers as a unified filesystem
 - Implements a copy-on-write behavior
 - Files from the top layers can obscure the files from the bottom layers
- **Layer hash**
 - Each layer has an hash based on its content
 - Layers are pulled and pushed compressed
- **Image hash**
 - Each image has an hash
 - The hash is function of the config file and of the layers
 - When an image changes, a new hash is generated

Docker: Container Data

- A container has access to different data
- **Container storage**
 - It is a copy-on-write layer in the image
 - It is ephemeral (only temporary data)
 - Data inside of containers is persisted as long as the container is not killed
 - If you stop or pause a container data is not lost
 - Containers are designed to be immutable
 - It's not good practice to write "persistent" data into containers
- **Bind-mount a local dir**
 - = a local dir is mounted to a dir inside a container
- **Docker volumes**
 - Docker provides volumes that exist separately from the container
 - E.g., to store the content of a Postgres DB
 - State is permanent across container invocations
 - Can be shared across containers

Docker Repos

- **Docker Repo (Registry)**
 - Store Docker images ### Devops = Devs + Ops
- **Devs**
- Implement the app
 - E.g., Python, virtual env
- Containerize the app
 - Create Dockerfile
 - Contain the instruction on how to build an image
- Build image
- Run the app as a container
- Test “locally”
- **Ops**
- Download container images
 - Contain filesystem, application, app dependencies
- Start / destroy containers
- In case of issues, it's easy to repro the problem
 - “Here is the log”
 - Run command line
- Deploy on a test system and debug

Containerizing an App

- = create a container with your app inside
- Develop your application code using the needed dependencies
 - E.g., install dependencies
 - Directly inside a container
 - Inside a virtual env
- Create a Dockerfile describing:
 - your app
 - its dependencies
 - how to run it
- Build image with **docker image build**
- (Optional) Push image to a Docker image registry
- Run / test container from image
- Distribute your app as a container (no installation)

Building a Container

- **Dockerfile**
 - Describe how to create a container
- **Build context**
 - > `docker build -t web:latest .`
 - `.` is the build context
 - Directory containing the application and what's needed to build it
 - Sent to Docker engine to build the application
 - Typically the Dockerfile is in the root directory of the build context

Dockerfile Example

```
FROM python:3.8-slim-buster LABEL  
maintainer="gsaggese@umd.edu" WORKDIR /app COPY  
requirements.txt requirements.txt RUN pip3 install -r requirements.txt  
COPY . . CMD ["python3", "-m", "flask", "run", "--host=0.0.0.0"]
```

Docker: Commands

Show all the available images

```
**\> docker images**
```

```
REPOSITORY TAG IMAGE ID CREATED SIZE
```

```
counter_app-web-fe latest 4bf6439418a1 17 minutes ago 54.7MB
```

```
**...**
```

Show a particular image

```
**\> docker images counter_app-web-fe**
```

```
counter_app-web-fe latest 4bf6439418a1 17 minutes ago 54.7MB
```

```
...
```

Note that `docker images ls` is incorrect since it shows the image

Delete an image

```
**\> docker rmi ...**
```

Show the running containers

```
**\> docker container ls**
```

```
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS NAMES
```

```
505541bcfe8b counter_app-web-fe "python app.py" 7 minutes ago Up 7
```

```
c1889540cfd2 redis:alpine "docker-entrypoint.s..." 7 minutes ago U
```

Docker: Commands

Show running containers

```
docker container ls CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS NAMES
281d654f6b8d counter_app_web-fe "python app.py" 5 minutes ago Up 5 minutes 0.0.0.0:5000->5001/tcp counter_app-web-fe-1
de55ae4104da redis:alpine "docker-entrypoint.s..." 5 minutes ago Up 5 minutes 6379/tcp counter_app-redis-1
Show volumes and networks
docker volume ls DRIVER VOLUME NAME
local counter_app_counter-vol
docker network ls NETWORK ID NAME DRIVER SCOPE
b4c1976d7c27 bridge bridge local
33ff702253b3 counter_app_counter-net bridge local
```

Docker: Delete state

```
Commands: -> docker container ls -> docker container rm  
$(docker container ls -q) -> docker images -> docker rmi  
$(docker images -q) -> docker volume ls -> docker volume rm  
$(docker volume ls -q) -> docker network ls -> docker network  
rm $(docker network ls -q)
```

Docker Tutorial

- tutorial_docker.md

Docker Compose

- **Manage multi-container apps running on a single node**
 - Describe app in a single *declarative* configuration YAML file
 - Instead of scripts with long Docker commands
 - Compose talks to Docker API to achieve what you requested
 - E.g., you need a client app and Postgres DB
 - E.g., microservices
 - Web front-end
 - Ordering
 - Back-end DB
- In 2020 Docker Compose has become an open standard for “code-to-cloud” process
- **Manage multi-container apps running on multiple hosts**
 - Docker Stacks / Swarm
 - Kubernetes

Docker Compose: Tutorial Example

- The default name for a Compose file is `docker-compose.yml`
 - You can specify `-f` for custom filenames
- **Top-level keys** are:
 - **version:**
 - Mandatory first line to specify API version
 - Ideally always use the latest version
 - Typically 3 or higher
 - **services:**
 - Define the different microservices
 - **networks:**
 - Creates new networks
 - By default it creates a bridge network to connect multiple containers on the same Docker host
 - **volumes:**
 - Creates new volumes
- **Key in services** describe a different “service” in terms of container
 - **Inner keys** specify the params of Docker run command

```
version: "3.8"
```

```
services:
```

```
  web-fe:
```

```
    build: .
```

```
    command: python app.py
```

Docker Compose: Commands

docker compose --help

Usage: docker compose [OPTIONS] COMMAND

Options:

--env-file string	Specify an alternate environment file
-f, --file stringArray	Compose configuration files
-p, --project-name string	Project name

Commands:

build	Build or rebuild services
convert	Converts the compose file to platform's canonical format
cp	Copy files/folders between a service container and the host
create	Creates containers for a service.
down	Stop and remove containers, networks
events	Receive real time events from containers.
exec	Execute a command in a running container.
images	List images used by the created containers
kill	Force stop service containers.
logs	View output from containers
ls	List running compose projects
pause	Pause services

Docker Compose: Commands

Build the containers for the services > `docker compose build` Pull the needed images for the services > `docker compose pull` Show the running services > `docker compose ps` Show the status of the service > `docker compose ls` Bring up the entire service > `docker compose up` Rebuild after trying out some changes in dockerfile/compose file > `docker-compose up --build --force-recreate` Show the processes inside each container > `docker compose top`

counter_app-redis-1

UID	PID	PPID	C	STIME	TTY	TIME	CMD
999	49590	49549	0	10:40	?	00:00:02	redis-server *

counter_app-web-fe-1

UID	PID	PPID	C	STIME	TTY	TIME	CMD
root	49614	49574	0	10:40	?	00:00:00	python app.py
root	49734	49614	1	10:40	?	00:00:08	/usr/local/bi

Docker Compose: Commands

Build the containers for the services

docker compose down [+] Running 3/2 Container counter_app-redis-1 Removed

Container counter_app-web-fe-1 Removed Network
counter_app_counter-net Removed Shutdown service removing the volume (i.e., resetting state) *docker-compose down -v*
Shutdown service removing images and volume *docker-compose down \$\\downarrow\$ --rmi all*

Docker Compose: Tutorial

- Example taken from <https://github.com/nigelpoulton/counter-app>
- tutorial_docker_compose
- > cd tutorials/tutorial_docker_compose
- > vi tutorial_docker_compose.md

Class Announcements

- **1 Project teams** - Posted UMD DATA605 - Class Project Teams - Spring 2023 - If your name is not on any of the teams, please send me an email - No midterm or final exam, complete class project to get a grade - **2 Team composition** - Teams based on your self-assessed skills - In each group there should be someone with experience with Git, Docker, Python, and so on - Teams are not perfect, but none of your team at your future jobs will be - Working in a team is a skill that takes a long time to hone: let's start practicing it - **3 Class project complexity** - Projects have about the same complexity - If not, we will try to account for this when grading - Projects were assigned randomly to the teams

Class Announcements

-4 Add personal info - Fill the spreadsheet UMD DATA605 - Class Project Teams - Spring 2023 with your information - Email - GitHub username (free account at <https://github.com>) - Telegram Handle for IM (free account at <https://telegram.org>) **-5 Next steps** - Read carefully UMD DATA605 - Class Project - If things are not clear send us an email or add a comment to the Google Doc - Read the Google doc corresponding to your assigned project from UMD DATA605 - Class Project Teams - Spring 2023 - Read the first Deliverable 1 from UMD DATA605 - Class Project **-6 Golden rule** - *Treat others how you want to be treated* - Everybody comes from a different place and different skill level, somebody has a job, somebody has a full-time work - *If you want to go fast, go alone; if you want to go far, go together*