



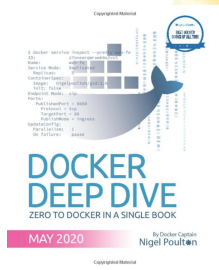
UMD DATA605 - Big Data Systems

Lesson 3.1: DevOps

Instructor: Dr. GP Saggese, gsaggese@umd.edu

Docker - Resources

- We will use Docker during the class project and most tutorials
- Concepts in the slides
- Class tutorials:
 - [tutorial_docker](#)
 - [tutorial_docker_compose](#)



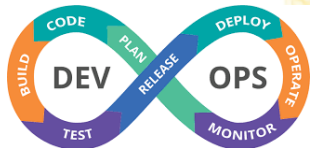
- Web resources:
 - [Docker Tutorial for beginners](#)
 - <https://labs.play-with-docker.com/>
 - <https://training.play-with-docker.com>
 - [A Beginner-Friendly Introduction to Containers, VMs and Docker](#)
 - [Official Docker Getting Started Tutorial](#)
- Mastery:
 - Poulton, [Docker Deep Dive: Zero to Docker in a single book](#), 2020

Application Deployment

- For (almost all) Internet companies, **the application is the business**
 - If the application breaks, the business stops
 - E.g., Amazon, Google, Facebook, online banks, travel sites, ...
- **Problem**
 - Release, deploy, manage, monitor applications
- **Solutions**
 - Before 2000s: “bare-metal era”
 - 2000s-2010s: “virtual machine era”
 - After ~2013: “container era”

DevOps

- **DevOps** = practices combining:
 - Software development (*dev*)
 - IT operations (*ops*)
- **Containers revolutionized DevOps**
 - Enable independence between app development and IT ops
 - One team creates an app
 - Another team deploys and manages apps
 - Foster collaboration and innovation
 - IT: "It doesn't work!"
 - Devs: "What? It works for me"



- Plan
- Code
- Build
- Test
- Release
- Deploy
- Operate
- Monitor

Run on bare metal

- **< 2000s**
 - Running few applications per server (no virtualization)
- **Pros**
 - No virtualization overhead
- **Cons**
 - Insecure due to no application separation
 - Expensive
- **Expensive / low efficiency**
 - New server for each application
 - Hard to spec machines → buy “big and fast servers”
 - Overpowered servers at 5-10% capacity
 - Significant spending during 2000 DotCom boom on machines and networks
- Resurfaced in 2020 with different Cloud Computing use cases
- **Winners:** Cisco, Sun, Microsoft



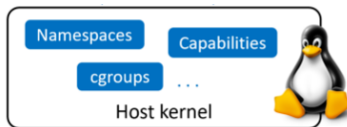
Virtual Machine Era

- **Circa 2000-2010: Virtual Machine**
 - Run multiple OS copies on the same hardware
- **Pros**
 - Run multiple apps safely on one server
 - Use existing servers with spare capacity
- **Cons**
 - Each VM needs an OS (wastes CPU, RAM, disk)
 - Buy OS license
 - Monitor and patch each OS
 - Slow boot times
- **Winners:** VMWare, RedHat, Citrix



Containers Era

- **Circa 2013: Docker becomes ubiquitous**
- **Docker**
 - Didn't invent containers
 - Made containers simple and mainstream
- Linux supported containers
 - Kernel namespaces
 - Control groups
 - Union filesystems
- **Pros**
 - Containers are fast and portable
 - Don't require full OS
 - All run on a single host
 - Reduce OS licensing cost
 - Reduce OS patching and maintenance
- **Cons**
 - CPU overhead
 - Toolchain to learn / use
- **Winners**: AWS, Microsoft Azure, Google (not Docker Inc.)



Serverless Computing

- **Containers run in an OS, OS runs on a host**
 - **Where is the host running?**
 - Local (your laptop)
 - On premise (your computer in a rack)
 - Cloud instance (e.g., AWS EC2)
 - **How is the host running?**
 - On bare-metal server
 - On a virtual machine
 - On a virtual machine running a virtual machine
- **Serverless computing**
 - Application runs without concern for “*how*” or “*where*”
 - E.g., AWS Lambda

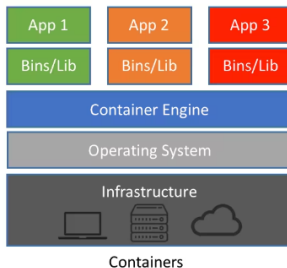
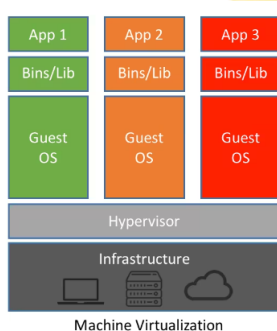
HW vs OS Virtualization

- **Hypervisor performs HW virtualization**

- Carves physical hardware into VMs
- Allocates CPUs, RAM, storage to a VM
- Like having multiple computers
- “Virtual machine tax”
 - Running 3 apps requires 3 VMs
 - Each VM takes time to start
 - Consumes CPU, RAM, storage
 - Needs OS license
 - Requires admins, patching
 - You just want to run 3 apps!

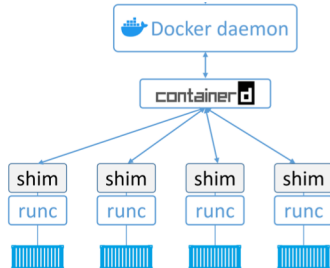
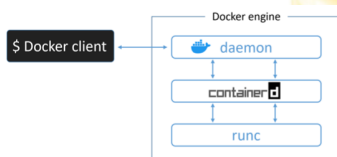
- **Containers perform OS virtualization**

- It's like having multiple OSes



Docker: Client-Server

- Docker relies on a **client-server architecture**
- **Docker client**
 - Command line interface
 - Communicate with server through IPC socket
 - E.g., `/var/run/docker.sock` or IP port
- **Docker engine**
 - Run and manage containers
 - Modular, built from OCI-compliant sub-systems
 - E.g., docker daemon, containerd, runc, plug-ins for networking and storage



Docker Architecture

- **Docker run-time**

- runc: start and stop containers
- containerd
 - Pull images
 - Create containers, volumes, network interfaces

- **Docker engine**

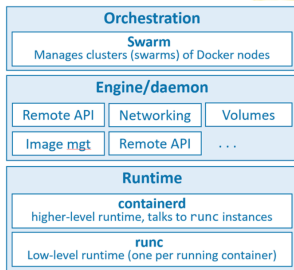
- dockerd
 - Expose remote API
 - Manage images, volumes, networks

- **Docker orchestration**

- docker swarm
- Manage clusters of nodes
- Replaced by Kubernetes

- **Open Container Initiative (OCI)**

- Standardize low-level components of container infrastructure
- E.g., image format, run-time API
- “Death” of Docker



Docker Container

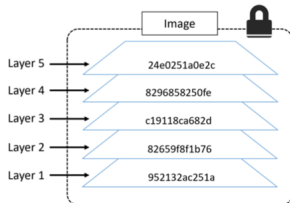
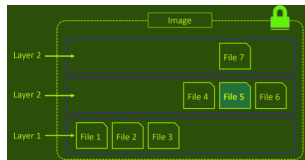
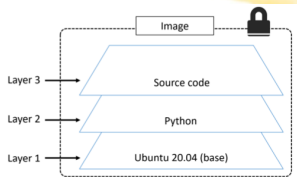
- **Docker Container**
 - Unit of computation
 - Lightweight, stand-alone, executable software package
- **Includes everything needed to run**
 - Code
 - Runtime/system libraries
 - Settings
- **Run-time object**
 - Docker images are build-time objects
 - Like program running (container) vs program code (image)

Docker Image

- **Docker Image**
 - Unit of deployment
- **Contains everything needed to run an app**
 - Application code
 - Application dependencies
 - Minimal OS support
 - File system
- Users can
 - Build images from Dockerfiles
 - Pull pre-built images from a registry
- Multiple layers stacked
 - Typically few 100s MBs

Docker Image Layers

- **Docker image** is a configuration file listing layers and metadata
 - Composed of read-only layers
 - Each layer is independent
 - Each layer comprises many files
- **Docker driver**
 - Stacks layers as a unified filesystem
 - Implements copy-on-write behavior
 - Files from top layers can obscure files from bottom layers
- **Layer hash**
 - Each layer has a hash based on content
 - Layers are pulled and pushed compressed
- **Image hash**
 - Each image has a hash
 - Hash is a function of the config file and layers
 - Image changes generate a new hash

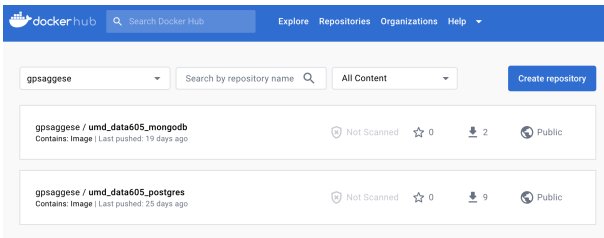
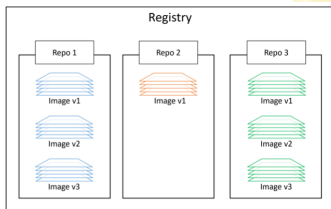


Docker: Container Data

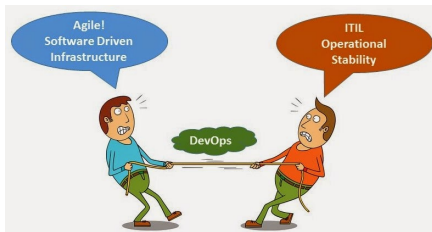
- A container has access to different data
- **Container storage**
 - Copy-on-write layer in the image
 - Ephemeral (temporary data)
 - Data persists until the container is killed
 - Stopping or pausing a container doesn't lose data
 - Containers are immutable
 - Avoid writing persistent data into containers
- **Bind-mount a local dir**
 - Mount a local dir to a dir inside a container
- **Docker volumes**
 - Volumes exist separately from the container
 - E.g., store Postgres DB content
 - State is permanent across container invocations
 - Shareable across containers

Docker Repos

- **Docker Repo (Registry)**
 - E.g., DockerHub, AWS ECR
 - Store Docker images
 - `<registry>/<repo>:<tag>`
 - E.g., `docker.io/alpine:latest`
 - Some repos are vetted by Docker
 - Unofficial repos shouldn't be trusted
 - E.g., <https://hub.docker.com/>



Devops = Devs + Ops



- **Devs**

- Implement app
 - Python, virtual env
- Containerize app
 - Create Dockerfile
 - Instructions to build image
- Build image
- Run app as container
- Test locally

- **Ops**

- Download container images
 - Filesystem, application, dependencies
- Start / destroy containers
- Reproduce issues easily
 - “Here is the log”
 - Run command line
 - Deploy on test system and debug