




UMD DATA605 - Big Data Systems

DevOps with Docker

Instructor: Dr. GP Saggese - gsaggese@umd.edu

Docker - Resources

- We will use Docker during the class project and most tutorials
- Concepts in the slides
- Class tutorials:
 - [tutorial_docker](#)
 - [tutorial_docker_compose](#)
- Web resources:
 - [Docker Tutorial for beginners](#)
 - <https://labs.play-with-docker.com/>
 - <https://training.play-with-docker.com>
 - [A Beginner-Friendly Introduction to Containers, VMs](#)
 - [Official Docker Getting Started Tutorial](#)
- Mastery:
 - Poulton,
[Docker Deep Dive: Zero to Docker in a single book,](#)
2020

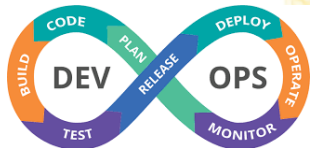


Application Deployment

- **For (almost all) Internet companies, the application is the business**
 - If the application breaks, the business stops working
 - E.g., Amazon, Google, Facebook, on-line banks, travel sites (e.g., Expedia), ...
- **Problem**
 - How to release / deploy / manage / monitor applications?
- **Solutions**
 - Before 2000s: “bare-metal era”
 - 2000s-2010s: “virtual machine era”
 - After ~2013: “container era”

DevOps

- **DevOps** = set of practices that combines:
 - Software development (*dev*)
 - IT operations (*ops*)
- **Containers revolutionized DevOps**
 - Enable true independence between application development and IT ops
 - One team creates an application
 - Another team deploys and manages the applications
 - Create a model for better collaboration (fewer conflicts) and innovation
 - IT: "It doesn't work!"
 - Devs: "What? It works for me"



- Plan
- Code
- Build
- Test
- Release
- Deploy
- Operate
- Monitor

Run on bare metal

- **< 2000s**
 - Running one or few applications on each server (without virtualization)
- **Pros**
 - No virtualization overhead
- **Cons**
 - Not safe / not secure since no separation between applications
 - Expensive
- **Expensive / low efficiency**
 - IT would buy a new server for each application
 - Difficult to spec out the machine → buy “big and fast servers”
 - Overpowered servers operating at 5-10% of capacity
 - Tons of money in the 2000 DotCom boom was spent on machines and networks
- It kind of came back in 2020 but with different use cases in Cloud Computing
- **Winners:** Cisco, Sun, Microsoft



Virtual Machine Era

- **Circa 2000-2010: Virtual Machine**

- Virtual machine technology = run multiple copies of OSes on the same hardware

- **Pros**

- VM runs safely and securely multiple applications on a single server
- IT could run apps on existing servers with spare capacity

- **Cons**

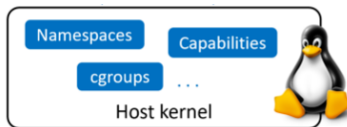
- Every VM requires an OS (waste of CPU, RAM, and disk)
- Buy an OS license
- Monitor and patch each OS
- VMs are slow to boot

- **Winners:** VMWare, RedHat, Citrix



Containers Era

- **Circa 2013: Docker becomes ubiquitous**
- **Docker**
 - Didn't invent containers
 - Made containers simple and mainstream
- Linux supported containers
 - Kernel namespaces
 - Control groups
 - Union filesystems
- **Pros**
 - Containers are fast and portable
 - Don't require full OS
 - All run on a single host
 - Reduce OS licensing cost
 - Reduce OS patching and maintenance
- **Cons**
 - CPU overhead
 - Toolchain to learn / use
- **Winners**: AWS, Microsoft Azure, Google (not Docker Inc.)



Serverless Computing

- **Containers run in an OS, OS runs on a host**
 - **Where is the host running?**
 - Local (your laptop)
 - On premise (your computer in a rack)
 - Cloud instance (e.g., AWS EC2)
 - **How is the host running?**
 - On bare-metal server
 - On a virtual machine
 - On a virtual machine running a virtual machine
- **Serverless computing**
 - Application runs without concern for “*how*” or “*where*”
 - E.g., AWS Lambda

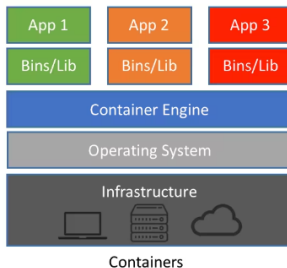
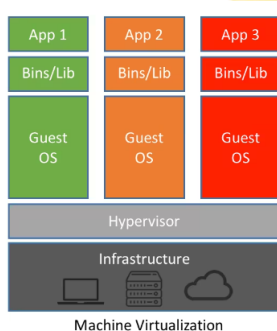
HW vs OS Virtualization

- **Hypervisor performs HW virtualization**

- Carves physical hardware into VMs
- Allocates CPUs, RAM, storage to a VM
- Like having multiple computers
- “Virtual machine tax”
 - Running 3 apps requires 3 VMs
 - Each VM takes time to start
 - Consumes CPU, RAM, storage
 - Needs OS license
 - Requires admins, patching
 - You just want to run 3 apps!

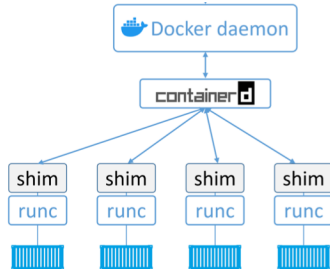
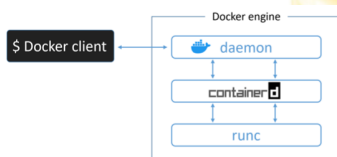
- **Containers perform OS virtualization**

- It's like having multiple OSes



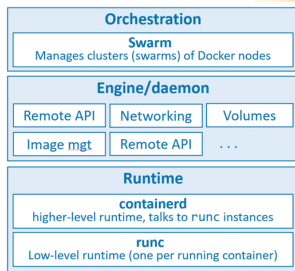
Docker: Client-Server

- Docker relies on a **client-server architecture**
- **Docker client**
 - Command line interface
 - Communicate with server through IPC socket
 - E.g., `/var/run/docker.sock` or IP port
- **Docker engine**
 - Run and manage containers
 - Modular, built from OCI-compliant sub-systems
 - E.g., docker daemon, containerd, runc, plug-ins for networking and storage



Docker Architecture

- **Docker run-time**
 - runc: start and stop containers
 - containerd
 - Pull images
 - Create containers, volumes, network interfaces
- **Docker engine**
 - dockerd
 - Expose remote API
 - Manage images, volumes, networks
- **Docker orchestration**
 - docker swarm
 - Manage clusters of nodes
 - Replaced by Kubernetes
- **Open Container Initiative (OCI)**
 - Standardize low-level components of container infrastructure
 - E.g., image format, run-time API
 - “Death” of Docker



Docker Container




- **Docker Container**
 - Unit of computation
 - Lightweight, stand-alone, executable software package
 - Includes everything needed to run, e.g.,
 - Code
 - Runtime/system libraries
 - Settings
 - Run-time object
 - Docker images are build-time objects
 - Like program running (container) vs program code (image)

Docker Image

- **Docker Image**
 - Unit of deployment
 - Contains everything needed to run an app
 - Application code
 - Application dependencies
 - Minimal OS support
 - File system
 - Users can
 - Build images from Dockerfiles
 - Pull pre-built images from a registry
 - Multiple layers stacked
 - Typically few 100s MBs

Docker Image Layers

- **Docker image** is a configuration file listing layers and metadata
 - Composed of read-only layers
 - Each layer is independent
 - Each layer comprises many files
- **Docker driver**
 - Stacks layers as a unified filesystem
 - Implements copy-on-write behavior
 - Files from top layers can obscure files from bottom layers
- **Layer hash**
 - Each layer has a hash based on content
 - Layers are pulled and pushed compressed
- **Image hash**
 - Each image has a hash
 - Hash is a function of the config file and layers
 - Image changes generate a new hash

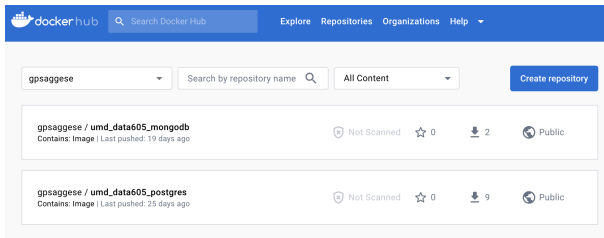
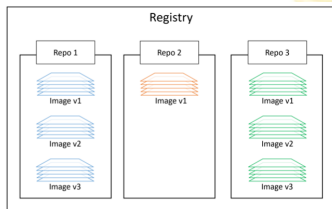


Docker: Container Data

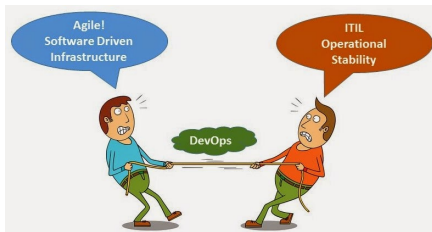
- A container has access to different data
- **Container storage**
 - Copy-on-write layer in the image
 - Ephemeral (temporary data)
 - Data persists until the container is killed
 - Stopping or pausing a container doesn't lose data
 - Containers are immutable
 - Avoid writing persistent data into containers
- **Bind-mount a local dir**
 - Mount a local dir to a dir inside a container
- **Docker volumes**
 - Volumes exist separately from the container
 - E.g., store Postgres DB content
 - State is permanent across container invocations
 - Shareable across containers

Docker Repos

- **Docker Repo (Registry)**
 - E.g., DockerHub, AWS ECR
 - Store Docker images
 - `<registry>/<repo>:<tag>`
 - E.g., `docker.io/alpine:latest`
 - Some repos are vetted by Docker
 - Unofficial repos shouldn't be trusted
 - E.g., <https://hub.docker.com/>



Devops = Devs + Ops



- **Devs**

- Implement app
 - Python, virtual env
- Containerize app
 - Create Dockerfile
 - Instructions to build image
- Build image
- Run app as container
- Test locally

- **Ops**

- Download container images
 - Filesystem, application, dependencies
- Start / destroy containers
- Reproduce issues easily
 - "Here is the log"
 - Run command line
 - Deploy on test system and debug

Containerizing an App

- **Containerizing an app** means creating a container with your app inside
- Develop application code with dependencies
 - Install dependencies
 - Inside a container
 - Inside a virtual env
- Create a Dockerfile describing:
 - App
 - Dependencies
 - How to run it
- Build image with `docker image build`
- (Optional) Push image to Docker image registry
- Run/test container from image
- Distribute app as a container (no installation required)

Building a Container

- **Dockerfile**
 - Describe how to create a container
- **Build context**
 - `docker build -t web:latest .` where `.` is the build context
 - Send directory containing the application to Docker engine to build the application
 - Typically the Dockerfile is in the root directory of the build context

Dockerfile: Example

```
FROM python:3.8-slim-buster
```

```
LABEL maintainer="gsaggese@umd.edu"
```

```
WORKDIR /app
```

```
COPY requirements.txt requirements.txt
```

```
RUN pip3 install -r requirements.txt
```

```
COPY . .
```

```
CMD ["python3", "-m", "flask", "run", "--host=0.0.0.0"]
```

Docker: Commands

- Show all the available images

```
> docker images
```

| REPOSITORY | TAG | IMAGE ID | CREATED | SIZE |
|--------------------|--------|--------------|----------------|--------|
| counter_app-web-fe | latest | 4bf6439418a1 | 17 minutes ago | 54.7MB |
| ... | | | | |

- Show a particular image

```
> docker images counter_app-web-fe
```

| | | | | |
|--------------------|--------|--------------|----------------|--------|
| counter_app-web-fe | latest | 4bf6439418a1 | 17 minutes ago | 54.7MB |
| ... | | | | |

- Delete an image

```
> docker rmi ...
```

- Show the running containers

```
> docker container ls
```

| CONTAINER ID | IMAGE | COMMAND | CREATED | STATUS | PORTS | NAMES |
|--------------|--------------------|------------------------|---------------|--------------|------------------------|-------------|
| 505541bcfe8b | counter_app-web-fe | "python app.py" | 7 minutes ago | Up 7 minutes | 0.0.0.0:5001->5000/tcp | counter_app |
| c1889540cfd2 | redis:alpine | "docker-entrypoint.sh" | 7 minutes ago | Up 7 minutes | 6379/tcp | counter_app |

Docker: Commands

- Show running containers

```
> docker container ls
CONTAINER ID   IMAGE                                COMMAND                  CREATED        STATUS
PORTS          NAMES
281d654f6b8d   counter_app-web-fe                 "python app.py"         5 minutes ago  Up 5 minutes
0.0.0.0:5001->5000/tcp   counter_app-web-fe-1

de55ae4104da   redis:alpine                       "docker-entrypoint.s..." 5 minutes ago  Up 5 minutes
6379/tcp        counter_app-redis-1
```

- Show volumes and networks

```
> docker volume ls
DRIVER      VOLUME NAME
local      counter_app_counter-vol
```

```
> docker network ls
NETWORK ID     NAME                        DRIVER    SCOPE
b4c1976d7c27   bridge                     bridge    local
33ff702253b3   counter-app_counter-net    bridge    local
```

Docker: Delete State

- Commands:

```
> docker container ls
> docker container rm $(docker container ls -q)

> docker images
> docker rmi $(docker images -q)

> docker volume ls
> docker volume rm $(docker volume ls -q)

> docker network ls
> docker network rm $(docker network ls -q)
```

Docker Tutorial

- tutorial_docker.md

Docker Compose

- **Manage multi-container apps running on a single node**
 - Describe app in a single *declarative* YAML file
 - Avoid long Docker command scripts
 - Compose interacts with Docker API to achieve the requested state
 - Examples:
 - Client app and Postgres DB
 - Microservices: Web front-end, Ordering, Back-end DB
- In 2020, Docker Compose became an open standard for “code-to-cloud” process
- **Manage multi-container apps running on multiple hosts**
 - Docker Stacks / Swarm
 - Kubernetes

Docker Compose: Tutorial Example

- Default name for a Compose file is `docker-compose.yml`
 - Specify `-f` for custom filenames
- **Top-level keys:**
 - `version:`
 - Mandatory first line for API version
 - Use latest version, typically 3 or higher
 - `services:`
 - Define microservices
 - `networks:`
 - Create new networks
 - Default bridge network connects containers on the same Docker host
 - `volumes:`
 - Create new volumes
- **Key in services** describes a “service” in terms of container
 - **Inner keys** specify Docker run command params

```
version: "3.8"
services:
  web-fe:
    build: .
    command: python app.py
    ports:
      - target: 5000
        published: 5001
    networks:
      - counter-net
    volumes:
      - type: volume
        source: counter-vol
        target: /code
  redis:
    image: "redis:alpine"
    networks:
      counter-net:
networks:
  counter-net:
volumes:
  counter-vol:
```

Docker Compose: Commands

```
> docker compose --help
```

```
Usage: docker compose [OPTIONS] COMMAND
```

```
Options:
```

| | |
|---------------------------|--|
| --env-file string | Specify an alternate environment file. |
| -f, --file stringArray | Compose configuration files |
| -p, --project-name string | Project name |

```
Commands:
```

| | |
|---------|---|
| build | Build or rebuild services |
| convert | Converts the compose file to platform's canonical format |
| cp | Copy files/folders between a service container and the local filesystem |
| create | Creates containers for a service. |
| down | Stop and remove containers, networks |
| events | Receive real time events from containers. |
| exec | Execute a command in a running container. |
| images | List images used by the created containers |
| kill | Force stop service containers. |
| logs | View output from containers |
| ls | List running compose projects |
| pause | Pause services |
| port | Print the public port for a port binding. |
| ps | List containers |
| pull | Pull service images |
| push | Push service images |
| restart | Restart containers |
| rm | Removes stopped service containers |
| run | Run a one-off command on a service. |
| start | Start services |
| stop | Stop services |
| top | Display the running processes |
| unpause | Unpause services |
| up | Create and start containers |
| version | Show the Docker Compose version information |

Docker Compose: Commands

- Build the containers for the services
`> docker compose build`
- Pull the needed images for the services
`> docker compose pull`
- Show the running services
`> docker compose ps`
- Show the status of the service
`> docker compose ls`
- Bring up the entire service
`> docker compose up`
- Rebuild after trying out some changes in dockerfile/compose file
`> docker-compose up --build --force-recreate`

Docker Compose: Commands

- Show the processes inside each container

```
> docker compose top
```

```
counter_app-redis-1
```

| UID | PID | PPID | C | STIME | TTY | TIME | CMD |
|-----|-------|-------|---|-------|-----|----------|---------------------|
| 999 | 49590 | 49549 | 0 | 10:40 | ? | 00:00:02 | redis-server *:6379 |

```
counter_app-web-fe-1
```

| UID | PID | PPID | C | STIME | TTY | TIME | CMD |
|------|-------|-------|---|-------|-----|----------|------------------------------------|
| root | 49614 | 49574 | 0 | 10:40 | ? | 00:00:00 | python app.py |
| root | 49734 | 49614 | 1 | 10:40 | ? | 00:00:08 | /usr/local/bin/python /code/app.py |

Docker Compose: Commands

- Build the containers for the services

```
> docker compose down
```

```
[+] Running 3/2
```

| | |
|---------------------------------|---------|
| Container counter_app-redis-1 | Removed |
| Container counter_app-web-fe-1 | Removed |
| Network counter_app_counter-net | Removed |
- Shutdown service removing the volume (i.e., resetting state)

```
> docker-compose down -v
```
- Shutdown service removing images and volume

```
> docker-compose down $\downarrow$ --rmi all
```

Docker Compose: Tutorial

- Example taken from <https://github.com/nigelpoulton/counter-app>
- [tutorial_docker_compose](#)
 - > `cd tutorials/tutorial_docker_compose`
 - > `vi tutorial_docker_compose.md`