

NoSQL Taxonomy

Instructor: Dr. GP Saggese - gsaggese@umd.edu

- Concepts in slides
- Silberschatz Chapter 23.6
- High-level view:
 - Seven Databases in Seven Weeks, 2e

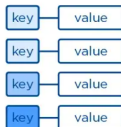


DB Taxonomy

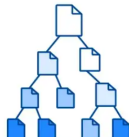
- **At least five DB genres**

- *Relational* (e.g., Postgres)
- *Key-value* (e.g., Redis)
- *Document* (e.g., MongoDB)
- *Columnar* (e.g., Parquet)
- *Graph* (e.g., Neo4j)

Key-Value



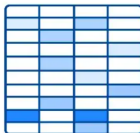
Document



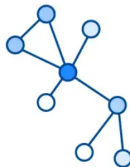
- **Criteria to differentiate DBs**

- Data model
- Trade-off with CAP theorem
- Querying capability
- Replication scheme

Wide-column



Graph



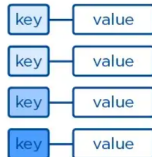
Relational DB

- E.g., *Postgres*, MySQL, Oracle, SQLite
- **Data model**
 - Set-theory and relational algebra
 - Data as tables with rows and columns
 - Many attribute types (e.g., numeric, strings, dates, arrays, blobs)
 - Strictly enforced attribute types
 - SQL query language
 - ACID consistency
- **Application**
 - Relational tabular data
- **Good for**
 - Known data layout, unknown access pattern
 - Schema complexity for query flexibility
 - Regular data
- **Not so good for**
 - Hierarchical data (not a nice row in tables)
 - Variable/dishomogeneous data (record-to-record variation)

Key-Value Store

- E.g., redis, DynamoDB, Git, AWS S3, filesystem
- **Data model**
 - Map keys (e.g., strings) to complex values (e.g., binary blob)
 - Support get, put, delete operations on a primary key
- **Application**
 - Caching data
 - Store users' session data in a web application
 - Store the shopping cart in an e-commerce application
- **Good for**
 - Data not “related” (e.g., no joins)
 - Fast lookups
 - Easy to scale horizontally using partitioning
- **Not so good for**
 - Data queries needed
 - Lacking secondary indexes and scanning

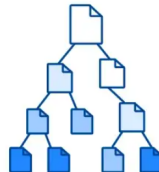
Key-Value



Document Store

- E.g., *MongoDB*, *CouchDB*
- **Data model**
 - Like key-value but value is a document (nested dict)
 - Each document has a unique ID (e.g., hash)
 - Any number of fields per document, even nested
 - E.g., JSON, XML, dict data
- **Application**
 - Semi-structured data
- **Good for**
 - Unknown data structure
 - Maps well to OOP models (less impedance mismatch)
 - Easy to shard and replicate over distributed servers
- **Not so good for**
 - Complex join queries
 - Denormalized form is the norm

Document



Columnar Store


- E.g., *HBase*, *Cassandra*, *Parquet*
- **Data model**
 - Store data by columns, not rows
 - Similar to key-value and relational DBs
 - Use keys to query values
 - Values are groups of columns
- **Application**
 - Store web pages
 - Store time series data
 - OLAP workloads
- **Good for**
 - Horizontal scalability
 - Enable compression and versioning
 - Sparse tables without extra storage cost
 - Inexpensive to add columns
- **Not so good for**
 - Design schema based on query plans
 - No native joins; applications handle joins

Wide-column

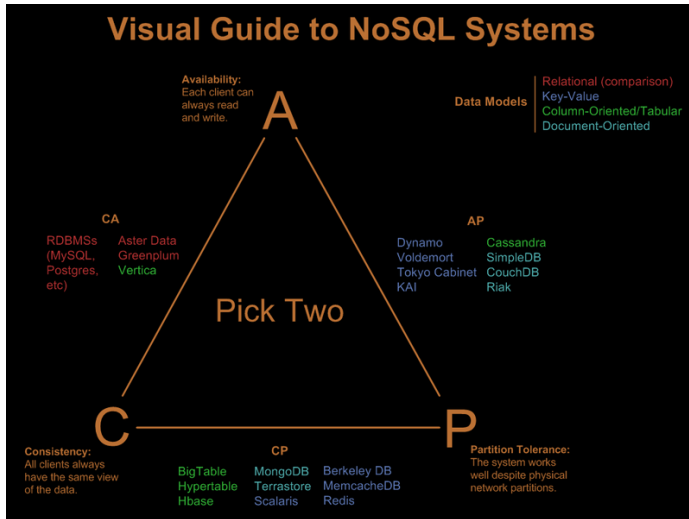
Graph DB

- E.g., *Neo4J*, *GraphX*
- **Data model**
 - Interconnected data, storing nodes and relationships
 - Nodes and edges have properties (key-value pairs)
 - Queries traverse nodes and relationships
- **Applications**
 - Social data
 - Recommendation engines
 - Geographical data
- **Good for**
 - Networked data, difficult to model with relational model
 - Matches OO systems
- **Not so good for**
 - Poor scalability, difficult to partition graph on different nodes
 - Store graph in graph DB, relations in key-value store

Graph



Taxonomy by CAP

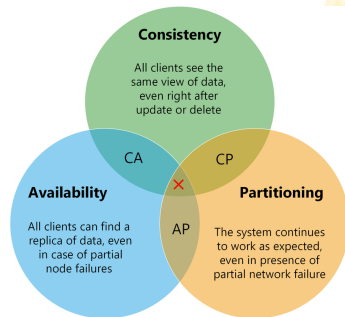


From

<http://blog.nahurst.com/visual-guide-to-nosql-systems>

Taxonomy by CAP

- **CA (Consistent, Available) systems**
 - Struggle with partitions, use replication
 - Traditional RDBMSs (e.g., PostgreSQL, MySQL)
- **CP (Consistent, Partition-Tolerant) systems**
 - Struggle with availability, maintain consistency across partitions
 - BigTable (column-oriented/tabular)
 - HBase (column-oriented/tabular)
 - MongoDB (document-oriented)
 - Redis (key-value)
 - MemcacheDB (key-value)
 - Berkeley DB (key-value)



- **AP (Available, Partition-Tolerant) systems**
 - Achieve “eventual consistency” via replication and verification
 - Dynamo (key-value)
 - Cassandra (column-oriented/tabular)
 - CouchDB (document-oriented)