

# NoSQL Document Stores MongoDB CouchDB

**Instructor:** Dr. GP Saggese - [gsaggese@umd.edu](mailto:gsaggese@umd.edu)\*\*

**TAs:** Krishna Pratardan Taduri, [kptaduri@umd.edu](mailto:kptaduri@umd.edu) Prahar

Kaushikbhai Modi, [pmodi08@umd.edu](mailto:pmodi08@umd.edu)

**with thanks to** Prof. Alan Sussman Prof. Amol

Deshpande Prof. Nguyen Vo (Utah State U.)

Prof. Kathleen Durant (Northeastern U.)

**v1.1**

UMD DATA605 - Big Data Systems

# Key-Value Store vs Document DBs

- **Key-value stores**

- Basically a map or a dictionary
  - E.g., HBase, Redis
- Typically only look up values by key
  - Sometimes can do search in value field with a pattern
- Uninterpreted value (e.g., binary blob) associated with a key
- Typically one namespace for all key-values

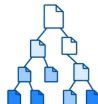
- **Document DBs**

- Collect sets of key-value pairs into *documents*
  - E.g., MongoDB, CouchDB
- Documents represented in JSON, XML, or BSON (binary JSON)
- Documents organized into *collections*
  - Similar to *tables* in relational DBs
  - Large collections can be partitioned and indexed

Key-Value



Document



# Resources

- All concepts in slides
- MongoDB tutorial
- Web
  - <https://www.mongodb.com/>
  - Official docs
  - pymongo
- Book
  - Seven Databases in Seven Weeks, 2e

The  
Pragmatic  
Programmers

## Seven Databases in Seven Weeks

Second Edition

A Guide to Modern  
Databases and the  
NoSQL Movement

Luc Perkins

with Eric Redmond and Jim R. Wilson

Series editor: Bruce A. Tate

Development editor: Jacquelyn Carter



# MongoDB



- Developed by MongoDB Inc.
  - Founded in 2007
  - Based on DoubleClick experience with large-scale data
  - Mongo comes from “hu-mongo-us”
- One of the most used NoSQL DBs (if not the most used)
- **Document-oriented NoSQL database**
  - Schema-less
    - No Data Definition Language (DDL), like for SQL
    - You can store maps with any keys and values
    - Application tracks the schema, mapping between documents and their meaning
  - Keys are hashes stored as strings
    - Document Identifiers id created for each document (field name reserved by Mongo)
  - Values use BSON format
    - Based on JSON (B stands for Binary)
- Written in C++
- Supports APIs (drivers) in many languages
  - E.g., JavaScript, Python, Ruby, Java, Scala, C++,

# MongoDB: Example of Document

- **A document is a JSON data structure**
- It corresponds to a row in a relational DB
  - Without schema
  - Primary key is \_id
  - Values nested to an arbitrary depth

```
{
  "_id" : ObjectId("4d0b6da3bb30773266f39fea"),
  "country" : {
    "$ref" : "countries",
    "$id" : ObjectId("4d0e6074deb8995216a8309e")
  },
  "famous_for" : [
    "beer",
    "food"
  ],
  "last_census" : "Sun Jan 07 2018 00:00:00 GMT -0700 (PDT)",
  "mayor" : {
    "name" : "Ted Wheeler",
    "party" : "D"
  },
  "name" : "Portland",
  "population" : 582000,
  "state" : "OR"
}
```


# MongoDB: Functionalities

---

- **Design goals**
  - Performance
  - Availability / scalability
  - Rich data storage (not rich querying!)
- **Dynamic schema**
  - No DDL (Data Definition Language)
  - Secondary indexes
  - Query language via an API
- **Several levels of data consistency**
  - E.g., atomic writes and fully-consistent reads (at document level)
- **No joins nor transactions across multiple documents**
  - Makes distributed queries easy and fast
- **High availability through replica sets**
  - E.g., primary replication with automated failover
- **Built-in sharding**
  - Horizontal scaling via automated range-based partitioning of data
  - Reads and writes distributed over shards

# MongoDB: Hierarchical Objects

- A Mongo **instance** has:
  - Zero or more “databases”
  - Mongo instance ~ Postgres instance
- A Mongo **database** has:
  - Zero or more “collections”
    - Mongo collection ~ Postgres tables
  - Mongo database ~ Postgres database
- A Mongo **collection** has:
  - Zero or more “documents”
    - Mongo document ~ Postgres rows
- A Mongo **document** has:
  - One or more “fields”
    - It has always primary key `_id`
    - Mongo field ~ Postgres columns



From

<https://www.mongodb.com/docs/manual/core/data-modeling-introduction>

# Relational DBs vs MongoDB: Terms and Concepts

RDBMS Concept	MongoDB Concept	Meaning in MongoDB
database	database	Container for collections
relation / table / view	collection	Group of documents
row / instance	document	Group of fields
column / attribute	field	A name-value pair
index	index	Automatic
primary keys	_id field	Always the primary key
foreign key	reference	Pointers
table joins	embedded documents	Nested name-value pairs

```
{
  "_id" : ObjectId("4d0b6da3bb30773266f39fea"),
  "country" : {
    "$ref" : "countries",
    "$id" : ObjectId("4d0e6074deb8995216a8309e")
  },
  "famous_for" : [
    "beer",
    "food"
  ],
  "last_census" : "Sun Jan 07 2018 00:00:00 GMT -0700 (PDT)",
  "mayor" : {
    "name" : "Ted Wheeler",
    "party" : "D"
  },
  "name" : "Portland",
  "population" : 582000,
  "state" : "OR"
}
```



# Relational vs Document DB: Workflows

---

- **Relational DBs**


- E.g., PostgreSQL
- Know what you want to store
  - Tabular data
- Do not know how to use it
  - Static schema allows query flexibility (e.g., joins)
- Complexity is at insertion time
  - Decide how to represent the data (i.e., schema)

- **Document DBs**

- E.g., MongoDB
- No assumptions on what to store
  - E.g., irregular JSON data
- Know a bit how to access data
  - You want to access the data by key
  - E.g., it's a nested key-value map
- Complexity is at access time
  - Get the data from the server
  - Process data on the client side

# Why Use MongoDB?

- Simple to query
  - Do the work on client side
- It's fast
  - 2-10x faster than Postgres
- Data model / functionalities suitable for most web applications
  - Semi-structured data
  - Quickly evolving systems
- Easy and fast integration of data
- Not well suited for heavy and complex transactions systems
  - E.g., banking system



# MongoDB: Data Model

- **Documents** are composed of field and value pairs
  - **Field names** are strings
  - **Values** are any BSON type
    - Arrays of documents
    - Native data types
    - Other documents
- E.g.,
  - `_id` holds an ObjectId
  - `name` holds a document that contains the fields `first` and `last`
  - `birth` and `death` are of Date type
  - `contribs` holds an array of strings
  - `views` holds a value of the NumberLong type

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```

← field: value  
← field: value  
← field: value  
← field: value


# MongoDB: Data Model

- Documents can be nested
  - Embedded sub-document
- **Denormalized data models**
  - Store multiple related pieces of information in the same record
  - Conceptually is the result of a join operation
- **Normalized data models**
  - Eliminate duplication
  - Represent many-to-many relationships

```
{
  _id: <ObjectId>,
  username: "123xyz",
  contact: {
    phone: "123-456-7890",
    email: "xyz@example.com"
  },
  access: {
    level: 5,
    group: "dev"
  }
}
```

Embedded sub-document

Embedded sub-document



# Schema Free

- MongoDB does not need any pre-defined data schema
- Every **document** in a **collection** can have different fields and values
  - No need for NULL values / union of fields like in relational DBs
- E.g., dishomogeneous data instances

Document

Document

Collection

```
{name: "will",  
  eyes: "blue",  
  birthplace: "NY",  
  aliases: ["bill", "la ciacco"],  
  loc: [32.7, 63.4],  
  boss: "ben"}
```

```
{name: "jeff",  
  eyes: "blue",  
  loc: [40.7, 73.4],  
  boss: "ben"}
```

```
{name: "brendan",  
  aliases: ["el diablo"]}
```

```
{name: "ben",  
  hat: "yes"}
```

```
{name: "matt",  
  pizza: "DiGiorno",  
  height: 72,  
  loc: [44.6, 71.3]}
```



mongoDB

# JSON Format

---

- JSON = JavaScript Object Notation
- Data is stored in field / value pairs
- A field / value pair consists of:
  - A field name (always a string)
  - Followed by a colon :
  - Followed by a typed value “name”: “R2-D2”
- **Data in documents is separated by commas ,**

“name”: “R2-D2”, race : “Droid”

- Curly braces {} hold documents {“name”: “R2-D2”, race : “Droid”, affiliation: “rebels”}
- An array is stored in brackets [] [{“name”: “R2-D2”, race: “Droid”, affiliation: “rebels”}, {“name”: “Yoda”, affiliation: “rebels”}]
- Supports:
  - Embedding of nested objects within other objects

Just references

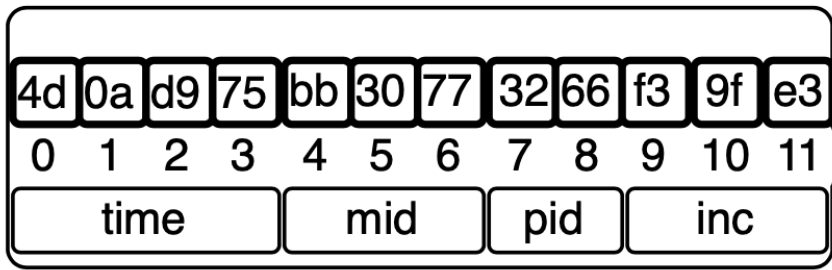
# BSON Format

---

- Binary-encoded serialization of JSON-like documents
  - <https://bsonspec.org>
- Zero or more key/value pairs are stored as a single entity
  - Each entry consists of:
    - a field name (string)
    - a data type
    - a value
- Similar to Protocol Buffer, but more schema-less
- Large elements in a BSON document are prefixed with a length field to facilitate scanning
- MongoDB understands the internals of BSON objects, even nested ones
  - Can build indexes and match objects against query expressions for BSON keys

# ObjectId

- Each JSON data contains an `_id` field of type `ObjectId`
  - Same as a SERIAL constraint incrementing a numeric primary key in PostgreSQL
- An `ObjectId` is 12 bytes, composed of:
  - a timestamp
  - client machine ID
  - client process ID
  - a 3-byte auto-incremented counter
- Each Mongo process can handle its own ID generation without colliding
  - Mongo has a distributed nature
- Details here





# Indexes

---

- **Primary index**
  - Automatically created on the `_id` field
  - B+ tree indexes
- **Secondary index**
- Users can create secondary indexes to:
  - Improve query performance
  - Enforce unique values for a particular field
- Single field index and compound index (like SQL)
  - Order of the fields in a compound index matters
- Sparse property of an index
  - The index contains only entries for documents that have the indexed field
  - Ignore records that do not have the field defined
- Reject records with duplicate key value if an index is unique and sparse
- Details at <https://www.mongodb.com/docs/manual/indexes/>

# CRUD Operations

---

- CRUD = Create, Read, Update, Delete
- **Create** `db.collection.insert(<document>)`  
`db.collection.update(<query>, <update>, upsert: true)` Upsert = update (if exists) or insert (if it doesn't)
- **Read** `db.collection.find(<query>, <projection>)`  
`db.collection.findOne(<query>, <projection>)`
- **Update** `db.collection.update(<query>, <update>, <options>)`
- **Delete** `db.collection.remove(<query>, <justOne>)` Details at <https://www.mongodb.com/docs/manual/crud/>

# Create Operations

---

- **db.collection** specifies the collection (like an SQL table) to store the document **db.collection.insert(<document>)**
  - Without **\_id** field, MongoDB generates a unique key **db.parts.insert(type: "screwdriver", quantity: 15)**
  - Use **\_id** field if it has a special meaning **db.parts.insert(\_id: 10, type: "hammer", quantity: 1)**
- Update 1 or more records in a collection satisfying **query**  
**db.collection.update(<query>, <update>, upsert: true)**
- Update an existing record or create a new record  
**db.collection.save(<document>)**
- A more modern OOP-like syntax than the COBOL / FORTRAN-inspired SQL

# Read Operations

---

- **find** provides functionality similar to **SQL SELECT** command  
**db.collection.find(<query>, <projection>).cursor** with:
  - = WHERE condition
  - = fields in result set
- **db.parts.find(parts: "hammer").limit(5)**
  - Return cursor to handle a result set
  - Can modify the query to impose limits, skips, and sort orders
  - Can specify to return the 'top' number of records from the result set
- **db.collection.findOne(<query>, <projection>)**

# More Query Examples

- Mongo has a functional programming flavor
  - E.g., composing operators, like `\textcolor{red}{$or}`

## SQL

```
SELECT * FROM users WHERE age>33
SELECT * FROM users WHERE age!=33
SELECT * FROM users WHERE name
LIKE "%Joe%"
SELECT * FROM users WHERE a=1 and b='q'
SELECT * FROM users WHERE a=1 or b=2
SELECT * FROM foo
WHERE name='bob' and (a=1 or b=2 )
SELECT * FROM users
WHERE age>33 AND age<=40
```

## Mongo

```
db.users.find({age: {$gt: 33}})
db.users.find({age: {$ne: 33}})
db.users.find({name: /Joe/})
db.users.find({a: 1, b: 'q'})
db.users.find({$or: [{a: 1}, {b: 2}]})
db.foo.find({name: "bob",
$or: [{a: 1}, {b: 2}]})
db.users.find({'age':
{$gt: 33, $lte: 40}})
```

# Query Operators

Command	Description
\$regex	Match by any PCRE-compliant regular expression string (or just use the // delimiters as shown earlier)
\$ne	Not equal to
\$lt	Less than
\$lte	Less than or equal to
\$gt	Greater than
\$gte	Greater than or equal to
\$exists	Check for the existence of a field
\$all	Match all elements in an array
\$in	Match any elements in an array
\$nin	Does not match any elements in an array
\$elemMatch	Match all fields in an array of nested documents
\$or	or
\$nor	Not or
\$size	Match array of given size
\$mod	Modulus
\$type	Match if field is a given datatype
\$not	Negate the given operator check

# Update Operations

---

- **db.collection.insert(<document>)**
  - Omit the **\_id** field to have MongoDB generate a unique key
  - **db.parts.insert(type: "screwdriver", quantity: 15)**
  - **db.parts.insert(\_id: 10, type: "hammer", quantity: 1)**
- **db.collection.save(<document>)**
  - Updates an existing record or creates a new record
- **db.collection.update(<query>, <update>, **upsert: true**)**
  - Will update 1 or more records in a collection satisfying query
- **db.collection.findAndModify(<query>, <sort>, <update>, <new>, <fields>, <upsert>)**
  - Modify existing record(s)
  - Retrieve old or new version of the record

# Delete Operations

---

- `db.collection.remove(<query>, <justone>)`
  - Delete all records from a collection or matching a criterion
  - `<justone>` specifies to delete only 1 record matching the criterion
- Remove all records in **parts** with **type** starting with h
  - `db.parts.remove(type: /h/ )`
- Delete all documents in the **parts** collection
  - `db.parts.remove()`



# MongoDB Features

---

- Document-oriented NoSQL store
- Rich querying
  - Full index support (primary and secondary)
- Fast in-place updates
- Agile and scalable
  - Replication and high availability
  - Auto-sharding
  - Map-reduce functionality
- Scale horizontally over commodity hardware
  - Horizontally = add more machines
  - Commodity hardware = relatively inexpensive servers

# MongoDB vs Relational DBs

---

- Keep the functionality that works well in RDBMSs
  - Ad-hoc queries
  - Fully featured indexes
  - Secondary indexes
- Do not offer RDBMS functionalities that don't scale up
  - Long running multi-row transactions
  - ACID consistency
  - Joins

# MongoDB Tutorial

---

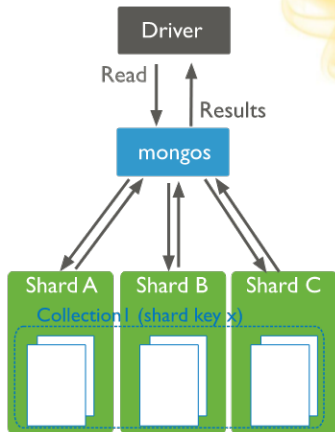
Tutorial is at GitHub The instructions are here:

```
> cd $GIT\_REPO/tutorials/tutorial\_mongodb
```

```
> vi tutorial\_mongo.md
```

# MongoDB Processes and Configuration

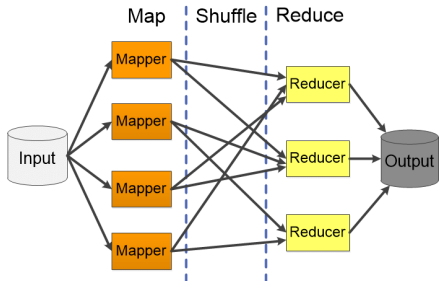
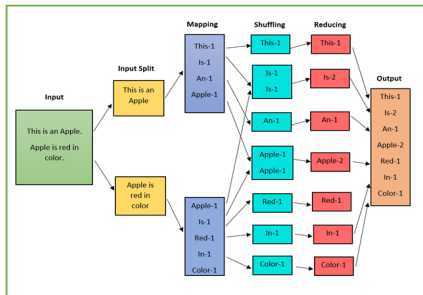
- **mongod**: database instance (i.e., a server process)
- **mongosh**: interactive shell (i.e., a client)
  - Fully functional JavaScript environment for use with a MongoDB
- **mongos**: database router
  - Process all requests
  - Decide how many and which **mongod** instances should receive the query (sharding / partitioning)
  - Collate the results
  - Send result back to the client
- You should have:
  - One **mongos** (router) for the whole system no matter how many **mongods** you have; or
  - One local **mongos** for every client if you wanted to minimize network latency



# MapReduce Functionality

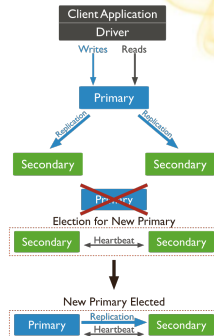
- Perform map-reduce computation given a collection of (keys, value) pairs
- Must provide at least a map function, reduction function, and the name of the result set

```
db.collection.mapReduce(
  <map\_function>,
  <reduce\_function>,
  {
    out: <collection>,
    query: <document>,
    sort: <document>,
    limit: <number>,
    finalize: <function>,
    scope: <document>,
    jsMode: <boolean>,
    verbose: <boolean>
  }
)
```



# Data Replication

- **Data replication** ensure:
  - Redundancy
  - Backup
  - Automatic failover
- Replication occurs through groups of servers known as **replica sets**
  - **Primary set**: set of servers that client asks direct updates to
  - **Secondary set**: set of servers used for duplication of data
  - Different properties can be associated with a secondary set,
    - E.g., secondary-only, hidden delayed, arbiters, non-voting
- If the primary fails the secondary sets “vote” to elect the new primary set



# Sync vs Async Replication


- **Synchronous replication:** updates are propagated to other replicas as part of a single transaction

- Implementations

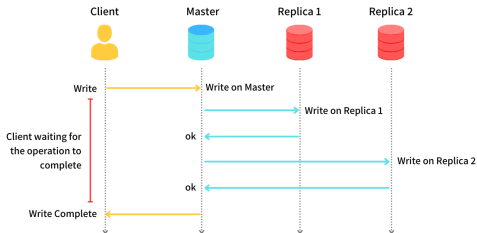
- 2-Phase Commit (2PC)
- Paxos
- Both solutions are complex / expensive

- **Asynchronous replication**

- The primary node propagates updates to replicas
- The transaction is completed before replicas are updated (even if there are failures)
- Commits are quick at cost of consistency



Asynchronous Replication



Synchronous Replication

# Data Consistency

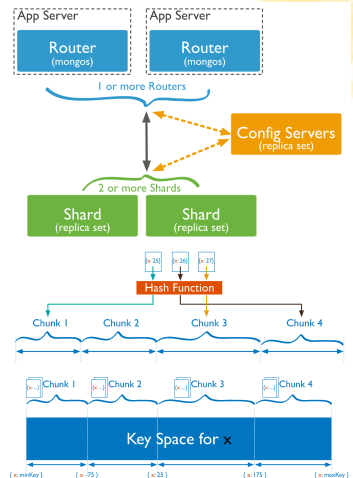
---

- **Client decides how to enforce consistency for reads**
- Reads to a primary have **strict consistency**
  - Reads reflect the latest changes to the data
  - All writes and *consistent* reads go to the primary
- Reads to a secondary have **eventual consistency**
  - Updates propagate gradually
  - Client may read a previous state of the database
  - All *eventually consistent* reads are distributed among the secondaries



# MongoDB: Sharding

- **Shard** = subset of data
  - A collection is split in pieces based on the shard key
  - Data distributed based on shard key or intervals [a, b)
- **Sharding** = method for distributing data across different machines
- **Horizontal scaling** can be achieved through sharding
  - Divide data and workload over multiple servers
  - Complexity in infrastructure and maintenance
- **mongos** acts as a query router interfacing clients and sharded cluster
  - Each shard can be deployed as a replica set
  - Config servers store metadata and configuration settings for cluster



# RDMBS Internals

---

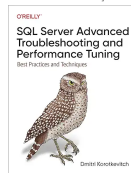
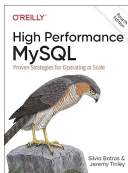
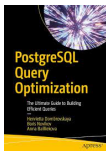
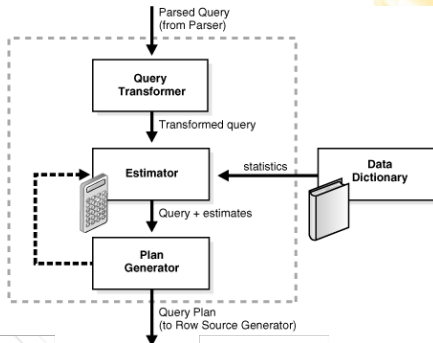
**Storage hierarchy** - How are tables mapped to files? - How are tuples mapped to disk blocks?

**Buffer Manager** - Bring pages from disk to memory - Manage the limited memory

**Query Processing Engine** - Given a user query, decide how to “execute” it - Specify sequence of pages to be brought in memory - Operate upon the tuples to produce results

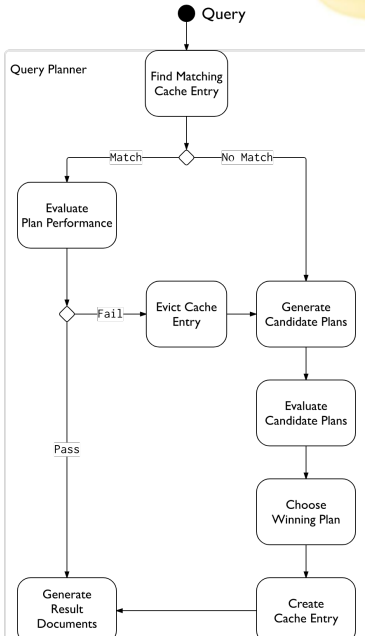
# Query Optimizer

- **RDBMSs: query optimizer is static**
  - Assign a cost to each query plan
  - Estimate some cost params (e.g., time to access data)
  - Search for the best query
  - At least traditional RDBMSs



# Query Optimizer

- **MongoDB: query optimizer is dynamic**
  - Try different query plans and learn which ones perform well
  - The space of query plans is not so large, because there are no joins
  - When testing new plans
    - Execute multiple query plans in parallel
    - As soon as one plan finishes, terminate the other plans
  - Cache the result
  - If a plan that was working well starts performing poorly try again different plans
    - E.g, data in the DB has changed, parameter values to a query are different



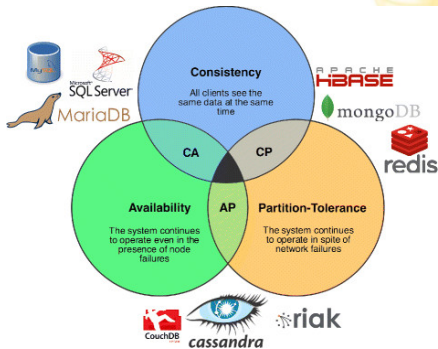
# MongoDB: Strengths

---

- Provide a flexible and modern query language
- High-performance
  - Implemented in C++
- Very rapid development, open source
  - Support for many platforms
  - Many language drivers
- Built to address a distributed database system
  - Sharding
  - Replica sets of data
- Tunable consistency
- Useful for working with a huge quantity of data not requiring a relational model
  - The relationships between the elements does not matter
  - What matters is the ability to store and retrieve great quantities of data

# MongoDB: Limitations

- No referential integrity
  - Aka foreign key constraint
- Lack of transactions and joins
- High degree of denormalization
  - Need to update data in many places instead of one
- Lack of predefined schema is a double-edged sword
  - You must have a data model in your application
  - Objects within a collection can be completely inconsistent in their fields
- CAP Theorem: targets consistency and partition tolerance, giving up on availability

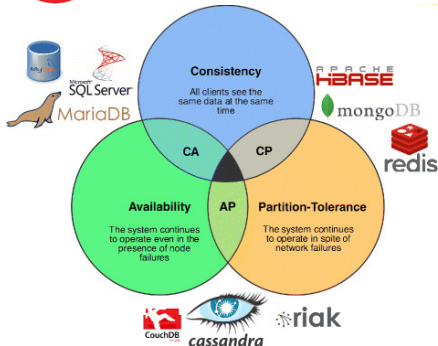


# Couchbase

- NoSQL document-oriented DB (like MongoDB)
- Couchbase = merge of CouchDB and membase
  - *CouchDB*
    - Open source document store
    - HTTP RESTful API to add, update, delete documents
    - Support all 4 ACID properties
  - *membase*
    - Distributed key-value store (like Redis)
    - Designed to scale both up and down
    - Highly available and partition tolerant
- Uses HTTP protocol to query and interact with objects in the DB
  - No query language
  - Objects stored in *buckets*
  - Collection of JSON docs, with no special relation to

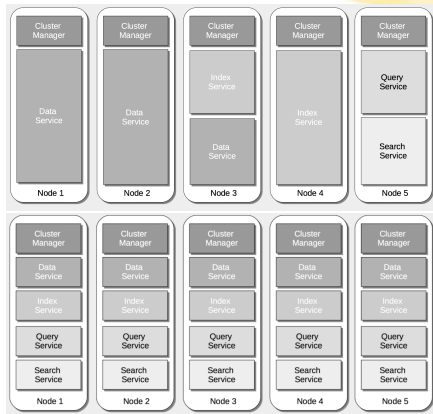


# Couchbase



# Architecture

- Every Couchbase node consists of **different services**:
  - Data service
  - Index service
  - Query service
  - Cluster manager component
- Services can run on separate nodes of the cluster, if needed
- **Data replication**
  - Across nodes of a cluster
  - Across data centers
- **Data service**
  - Writes data *asynchronously* to disk after acknowledging to the client
  - Optionally *synchronous*: ensure data is written to more than one server before acknowledging a write





# Queries

---

- **Can create multiple views over documents**
  - Views are optimized / indexed by Couchbase for fast queries
  - Re-indexed when underlying documents changes
  - Can do full-text searches using the indexes
- **Perform well when:**
  - There are infrequent changes to the structure of documents
  - Know in advance what kinds of queries you want to execute
- **Query**
  - Uses a custom query language called N1QL (“nickel”)
  - Extends SQL to JSON documents
  - Queries over multiple documents using (server-side) joins
- **Map-reduce support**
  - (Map) First define a view with the columns of the document your are interested in
  - (Reduce) Optionally define aggregate functions over the data