

UMD DATA605 - Big Data Systems

Python Dask

Instructor: Dr. GP Saggese - gsaggese@umd.edu**

TAs: Krishna Pratardan Taduri, kptaduri@umd.edu Prahar
Kaushikbhai Modi, pmodi08@umd.edu

v1.1

Dask: Resources

- Web resources:
 - Dask project
 - Dask examples
- Tutorial
 - Dask_tutorial
 - Dask_advanced_tutorial
- Class project
- Mastery
 - Data science with Python and Dask, 2019
 - Amazon



Dataset Size Issues

- **Small datasets (< 1 GB)**
 - Fits into RAM
 - Manipulation doesn't require paging to disk
- **Medium dataset (< 1TB)**
 - Doesn't fit into RAM
 - Fits into local disk
 - Performance penalty imposed by using local disk
 - Need multiple CPU cores
 - Difficult to take advantage of parallelism with Python / Pandas
- **Large dataset (> 1TB)**
 - Doesn't fit into RAM
 - Doesn't fit into local disk
 - Need multiple servers
 - Python / Pandas were not built to operate on distributed datasets
 - Use frameworks for massive datasets
 - E.g., Hadoop, Spark, Dask, Ray



Dataset Size Issues

- **Small datasets**
 - < 1 GB
- **Medium dataset**
 - < 1TB
- **Large dataset**
 - *1TB*
- **The thresholds are fuzzy and changing over time**
 - E.g., you can scale the computer 10x and get 10x bigger data sets
- **Problem with scaling datasets**
 - Long run times
 - Rewriting code in different language / API for datasets of different size
 - Need to think about *what to do* it and *how to do* it efficiently
 - Cumbersome framework (Pandas easy, Hadoop difficult)

Dask

- **Dask is written in Python**

- It scales natively Numpy, Pandas, sklearn
- Dask objects are wrappers (don't just mirror the interface) objects from the respective libraries (e.g., Pandas DataFrame, numpy array)
- Parallel parts are called “chunks” or “partitions”
 - Are queued to be worked on
 - Shipped between machines
 - Worked locally on a machine

- **Pros**

- Users don't need to learn a new language, but can use familiar interfaces
- Can focus on writing code that is optimized for parallelism
 - Dask does the heavy lifting

- **Scaling Dask is easy**

- Users can write a prototype task on a local machines and use a cluster when needed
- No need to refactor existing code
- No need to handle cluster-specific issues
 - E.g., resource management, data recovery, data movement
- Dask runs on multi-core
- Dask can use cluster managers
 - E.g., Yarn, Mesos, Kubernetes, AWS ECS

Dask Layers




Figure 3: alt_text

Scaling Up vs Scaling Out

- **Scaling up**

- = replace equipment with larger, faster equipments
 - E.g., buy a larger pot, replace knife with food processor
- **Pros**
 - You got better hardware, nothing else needs to change (e.g., code)
- **Cons**
 - There will be a time where you exceed the capacity of the current machines
 - Cost: more powerful machines are expensive

- **Scaling out**

- = divide the work between many workers in parallel
 - E.g., buy more pots and hire more cooks
- **Pros**
 - Task scheduler organizes computation, assigning workers to each task
 - More cost-effective solution since no specialized hardware is needed
- **Cons**
 - Need to write code to expose parallelism
 - Costs of maintaining a cluster



Dask: Computation

- **Lazy computations**

- User defines the transformations on the data
- No need to wait for one computation to finish before defining the next
- Avoid loading the entire data in memory by operating in chunks
- E.g.,
 - Split a 2GB file into 32 64MB chunks
 - Operate on 8 chunks at a time on each server
 - The max memory consumption doesn't exceed $512\text{MB} = (8 \times 32)$
- Each task tracks object dimensions and data types
 - No code is executed

- **compute()**


- Running a computation (aka materializing) `''' python missing_count_pct = missing_count.compute()'''`

- **persist()**

- As soon as a node in the graph emits results, its intermediate work is
- discarded to minimize memory usage
- If we need to do additional computation on intermediate nodes we need to re-run the graph
- **persist()** tells Dask to keep the intermediate result in memory
- This speeds up a large and complex DAG that needs to be reused many

Dask: Data Structures

- **Dask DataFrame** implements Pandas DataFrame
 - Tabular / relational data
- **Dask Array** implements numpy ndarray
 - Multidimensional array**
- **Dag Bag** coordinates Python lists of objects
 - Parallelize computations on unstructured or semi-structured data



Task Reading Data

- `dask.dataframe.read_csv()`
 - Doesn't load the data in memory with
 - Tries to infer the types of the columns
 - By randomly sampling some data
 - Best to set the data types
 - Even better is to use Parquet since it stores data and types together
- Partitions = chunks of data that can be worked independently
 - E.g., 33 partitions
 - Graph is composed of 99 tasks
 - Each partition reads data, splits data, initializes df object

Column names →

Summons number	Plate ID	Registration state	Plate type	Issue date	Violation code
----------------	----------	--------------------	------------	------------	----------------

Number of partitions → npartitions=33

Column datatypes →

int64	object	object	object	object	int64
...
...
...
...

Low Level APIs: Delayed


- Handle computations that don't fit in native Dask data structures (e.g., Dask DataFrame)
- In the example below there is parallelism that can be exploited

```
def inc(x):  
    return x + 1  
  
def double(x):  
    return x * 2  
  
def add(x, y):  
    return x + y  
  
data = [1, 2, 3, 4, 5]  
  
output = []  
for x in data:  
    #  $(x + 1) + (x * 2) = 3x + 1$   
    a = inc(x)  
    b = double(x)  
    c = add(a, b)  
    # 1 -> 4  
    # 2 -> 7  
    # 3 -> 10
```



Low Level APIs: Futures

- In parallel programming, a “future” encapsulates the asynchronous execution of a callable, representing the eventual result of the operation
- Futures is the most general way of specifying concurrency in Dask
 - Everything can be expressed in terms of futures
 - User can specify what's blocking and what's not blocking
- Python **concurrent.futures**
 - High-level interface for asynchronously executing callables
 - Thread pool or Process pool (same interface **Executor**)
- Dask extends **concurrent.futures**
 - Dask client can be used anywhere **concurrent.futures** can be used




Different Types of Parallel Workload

- Break program in medium-size tasks of computation
 - E.g., a function call

Embarrassingly Parallel

Hadoop/Spark/Dask/Airflow/Prefect



Encoding Task Graph

- Dask encodes tasks in terms of Python dicts and functions

```
def inc(i):  
    return i + 1
```

```
def add(a, b):  
    return a + b
```

```
x = 1
```

```
y = inc(x)
```

```
inc(x)
```

Task Scheduling

- Data collections (Bags, Arrays, DataFrame) and their operations create task graphs
 - Nodes in the task graph are Python functions
 - Edges are dependencies (e.g., output from one task used as input in another task)
- Task graphs are scheduled for execution
- Single-machine scheduler
 - Use local process or thread pool
 - Simple but it can only run on a single machine
- Distributed scheduler
 - It can run locally or distributed across a cluster

Collections

(create task graphs)

Dask Array

Dask DataFrame

Dask Bag




Task Graph



Schedulers

(execute task graphs)



Single-machine
(threads, processes,
synchronous)

Distributed

Task Scheduling

- **Dask task scheduler orchestrates the work dynamically**
 - Not a static scheduling of operations like a relational DB
 - When the computation takes place, Dask dynamically assesses:
 - What tasks has been completed
 - What tasks is left to do
 - What resources (CPUs) are free
 - Where the data is located
- **This dynamic approach handles a variety issues:**
 - Worker failure
 - Just re-run
 - Workers completing work at different speeds because of:
 - Different computation
 - Different hardware
 - Different workloads on the servers
 - Slower access to the data
 - Network unreliability
 - Just re-run or remove the isolated nodes



Dask vs Spark

- Spark has
- **Pros**
 - Popular framework for analyzing large datasets
 - In-memory alternative to MapReduce / Hadoop
- **Cons**
 - Spark is a Java library, supporting Python through PySpark API
 - Python code is executed on JVM through py4j
 - Difficult to debug since execution occurs outside Python
 - Different DataFrame API than Pandas
 - Learn how to do things “the Spark way”
 - You might need to implement things twice to go from exploratory analysis to large experiments / production
 - Optimized for MapReduce operations over a collection
 - Difficult to set-up and configure

Tutorial

Tutorial - From the official documentation

<https://docs.dask.org/en/stable/10-minutes-to-dask.html>