

## Knowledge Representation

**Instructor:** Dr. GP Saggese - [gsaggese@umd.edu](mailto:gsaggese@umd.edu)

**References:**

- Mostly papers and Internet
- AIMA 7: Logical agents
- AIMA 8, First-order logic
- AIMA 9: Inference in first-order logic
- AIMA 10, Knowledge representation

- ***Knowledge Representation***
  - Basics of Knowledge Representation
  - Examples of Logic
  - Logical Agents
  - Ontologies
  - Reasoning in Ontologies
- Propositional logic
- First-order Logic
- Non-classical Logics

- Knowledge Representation
  - *Basics of Knowledge Representation*
  - Examples of Logic
  - Logical Agents
  - Ontologies
  - Reasoning in Ontologies
- Propositional logic
- First-order Logic
- Non-classical Logics

# What is Knowledge Representation?

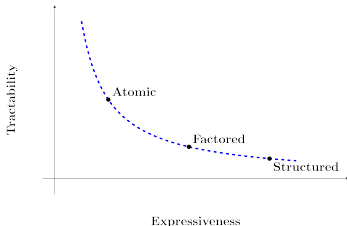
---

- **Knowledge Representation (KR)** is the study of how to formally encode information so that machines can reason with it
  - E.g., rules, logic, ontologies, semantic networks
  - It is at the heart of symbolic AI and complements learning-based approaches
- Defines:
  - **structure** (how knowledge is organized)
  - **semantics** (what it means)
- Serves as a bridge between perception (data) and reasoning (logic)
  - Essential for explainability and transparency in intelligent systems
- Enables machines to:
  - Draw conclusions
  - Perform planning
  - Answer queries
  - ...

# Expressiveness vs. Tractability

- **Tradeoff in AI / ML**

- **Expressiveness**: richness of concepts that can be captured
- **Tractability**: whether reasoning can be performed efficiently
- More expressive languages lead to harder computation



- Choosing the right knowledge representation formalism depends on the application needs

- **Atomic**

- Treats each state as a single, indivisible entity
- E.g., depth-first search algorithms (e.g., E3 in Chess)
- Simple and fast but limited in capturing complex relationships

- **Factored**

- E.g., propositional logic
- E.g.,  $P_{1,1}$ : "Pit in square (1,1)",  $B_{1,2} \Leftrightarrow (P_{1,1} \vee P_{2,2} \vee P_{1,3})$
- Captures relationships between variables but can't express complex structures

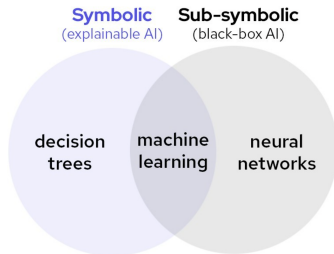
- **Structured**

- E.g., first-order logic
- $\forall x \forall y \text{ Father}(x, y) \Rightarrow \text{Parent}(x, y) = \text{"A father of a person is their parent"}$
- More expressive but undecidable in general

# Symbolic vs. Sub-symbolic Representation

---


- **Symbolic knowledge** representation uses discrete, human-readable symbols
  - E.g., logic, knowledge graphs
  - Interpretable and suitable for rule-based reasoning
  - Struggle with ambiguity
- **Sub-symbolic knowledge** representation uses learned, distributed representations
  - E.g., vector embeddings
  - E.g., deep learning excels at perception and pattern recognition
  - Lack transparency
- **Neuro-symbolic approaches** blends the two approaches
  - Reason over learned concepts using structured logic



# Neuro-symbolic Approach: Conceptual Spaces

- **Conceptual spaces** are frameworks for representing knowledge using geometric structures

- A concepts is a region in a multidimensional space defined by quality dimensions
- Similarity between objects is modeled by spatial distance
- Each dimension represents an interpretable feature



- **Example**

- Dimensions: Color, Size, Shape
- “Apple” occupies a region that is typically red / medium-sized / round
- “Banana” occupies a different region: yellow / medium / curved and long

- Differences from symbolic representations

- Symbolic systems use discrete symbols without structure
  - E.g., Apple vs Banana

- **Benefits**

- Natural modeling of similarity and vagueness
- Useful for grounding symbols in perception (link between sensory inputs and symbolic language)

# Natural languages

---

- **Natural languages** (e.g., English, Italian) are:
  - **Expressive**
    - Medium for communication rather than representation
  - **Ambiguous**
    - E.g., “spring” is both a “season” and “something that goes boing”
  - **Context-dependent**
    - Meaning depends on the sentence and context
    - E.g., “Look!”
- **Sapir-Whorf hypothesis**
  - The language you speak shapes how you perceive, think about, and experience the world
    - Even through arbitrary grammatical features, such as gender of nouns
  - Some languages lack words for certain concepts (e.g., direction)
  - Some languages have many words for the same concepts
    - E.g., Arctic languages have many words for snow (fresh, hard)
  - Newspeak from 1984 (Orwell)
    - You can't think certain concepts since you don't have words for it



# Procedural vs Declarative Approaches

---

- **Procedural approach**
  - Focuses on *how* a task is done
  - Encodes desired behavior directly into the program
  - E.g., a robot programmed with specific steps to navigate a maze
- **Declarative approach**
  - Specifies *what* the goal is, *not how* to achieve it
  - Describes relationships between actions and goals
  - Leaves solution search to the system
  - E.g., describing the goal “reach the exit” and letting the system find the path
- **Comparison**
  - Procedural: more control, less flexibility
  - Declarative: more abstraction, easier to modify or extend
- **Integration of approaches**
  - Many successful AI systems use a hybrid approach
  - Declarative knowledge can be compiled into procedural code
  - E.g., a planner generates procedures (plans) from declarative goals

# Programming Languages

---

- A **programming language** (e.g., C++, Python) is a formal language
  - Data structures represent facts
  - Code updates data structures in a domain-specific way
  - Programming is procedural (vs declarative)
- **Limitations:**
  - Programming languages lack:
    1. A general mechanism for deriving facts from other facts
      - Code updates data structures based on programmer's domain knowledge
    2. Expressiveness to handle partial information
      - A variable represents a single value or unknown
      - Can't easily handle partial information or quantify uncertainty
      - E.g., "A white knight is in b1 or in f6"
- **Declarative language**
  - E.g., propositional logic, first order logic
  - Knowledge and inference are separate:
    1. Knowledge represents the domain-specific problem
    2. Inference is domain independent
  - Compositional semantics
    - The meaning of a sentence is a function of the meaning of its parts

- Knowledge Representation
  - Basics of Knowledge Representation
  - *Examples of Logic*
  - Logical Agents
  - Ontologies
  - Reasoning in Ontologies
- Propositional logic
- First-order Logic
- Non-classical Logics

# Propositional Logic

---

- Uses atomic statements (propositions) and logical connectives
  - **Syntax**
    - Atomic formulas:  $P, Q$
    - Connectives: NOT ( $\neg$ ), AND ( $\wedge$ ), OR ( $\vee$ ), IMPLIES ( $\implies$ )
  - **Semantics**
    - Based on truth tables
    - Each proposition has a binary truth value: true or false
  - **Inference mechanisms**
    - Modus ponens: from  $P$  and  $P \implies Q$ , infer  $Q$
    - Resolution: derive contradictions to infer conclusions
- **Applications:** best used in closed and well-defined environments
  - Digital circuit design
  - Rule-based systems
  - Simplified AI models
- **Limitations**
  - Cannot represent objects, relations, or quantifiers
  - Not suitable for open or dynamic domains

# First-Order Logic (FOL)

---

- **Extension of propositional logic**

- Introduces predicates, variables, and quantifiers
  - Variables  $x$
  - Predicate  $Human(x)$
  - Universal quantifier “for all”  $\forall$
  - Existential quantifier “there exists”  $\exists$
- E.g.,  $\forall x(Human(x) \implies Mortal(x)) = \text{“All humans are mortal”}$
- Represents more complex and structured knowledge than propositional logic
- Can model properties, relationships, and quantification over objects

- **Inference mechanisms**

- Unification: matches predicates with variables
- Resolution: deduces new facts from known statements
- Model checking: verifies truth of statements under specific interpretations

- **Computational properties**

- Inference is semi-decidable: valid conclusions may require infinite time
- More powerful but computationally more complex than propositional logic

- **Applications**

- Knowledge representation
- Automated theorem proving
- Semantic web and ontologies

# Reasoning and Inference in Logic

---

- **Logical inference** is the process of deriving new facts from known ones using formal rules
  - Used to make decisions and answer questions based on a Knowledge Base
- **Knowledge base (KB):**
  - A structured set of facts and rules used for logical reasoning
- **Inference engine:**
  - Mechanism that applies logical rules to a KB to derive conclusions or answer queries
  - **Forward chaining:**
    - Starts with known facts and applies inference rules to extract more data
    - E.g., given  $A \rightarrow B$  and  $A$ , infer  $B$
  - **Backward chaining:**
    - Begins with a goal and works backward to find supporting facts
    - E.g., to prove  $B$ , check if  $A \rightarrow B$  and then prove  $A$
  - **Resolution:**
    - A complete inference rule for propositional and first-order logic
    - Useful in automated theorem proving
  - **Entailment** ( $KB \models \alpha$ ):
    - Sentence  $\alpha$  is entailed by KB if it is true in all models where KB is true

# Rule-Based Systems (1/2)

---

- A rule-based system uses “if-then” rules to derive conclusions or make decisions
  - It mimics human decision-making by applying logical rules to a set of facts
- **Key components**
  - Knowledge base: stores facts and rules
  - Inference engine: applies rules to known facts to infer new facts or take actions
  - Working memory: holds current facts being considered
- **How it works**
  - *Match*: find rules whose conditions match current facts
  - *Conflict resolution*: decide which rule to apply if multiple rules match
  - *Act*: apply the chosen rule to modify facts or trigger actions
  - *Repeat*: continue until no more rules can be applied
  - **E.g.,**
    - *Rule*: If a patient has a fever and a rash, then suggest measles
    - *Fact*: Patient has a fever and a rash
    - *Conclusion*: Suggest measles

# Rule-Based Systems (2/2)

---

- **Pros**
  - Easy to modify and update rules
  - Transparent and explainable reasoning
  - Good when expert knowledge can be clearly articulated
- **Cons**
  - Hard to scale to very large or complex domains
  - Cannot handle uncertainty without extensions (e.g., probabilistic reasoning)
  - Rule conflicts and maintenance can become challenging
- **Applications**
  - Expert systems
    - E.g., medical diagnosis, technical troubleshooting
  - Business rule engines
  - Game AI
  - Tools for legal reasoning

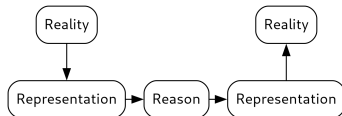


# Grounding

---

- **Grounding**

- Connect abstract symbols to real-world entities or observations
- E.g., link Apple to the fruit “apple”
- Make representations meaningful beyond syntax
  - Enable agents to act meaningfully in the real world
  - Avoid purely symbolic manipulation without real-world relevance



- **Challenges**

- Noisy, incomplete sensory data
- Complex, context-dependent mapping from inputs to concepts

- **Applications**

- Robotics: object recognition, manipulation
- Natural language understanding
- Autonomous agents, cognitive systems

- Knowledge Representation
  - Basics of Knowledge Representation
  - Examples of Logic
  - **Logical Agents**
  - Ontologies
  - Reasoning in Ontologies
- Propositional logic
- First-order Logic
- Non-classical Logics

# Reflex Agents

---

- **Reflex agents** act based on the current percept, ignoring percept history
  - Operate using a condition-action rule: *"if condition, then action"*
  - Rely on predefined rules
  - Have no internal state or memory
    - E.g., a thermostat: *"if temperature < threshold, turn on the heater"*
- **Pros**
  - Fast and efficient in well-defined environments
- **Cons**
  - Struggle with complex or partially observable environments
  - Cannot plan ahead or learn from experience
- **Application**
  - Simple or fully observable environments where quick reactions are sufficient

# Knowledge-based Agents

---

- **Intelligence** is achieved by *reasoning* on an internal *representation of knowledge*
- **Knowledge-based agents:**
  - Form representations of a complex world
  - Use inference to derive new representations
  - Deduce actions from new representations
  - Accept tasks as goal descriptions
  - Achieve competence by learning new knowledge
  - Adapt to changes by updating knowledge
  - Utilize a knowledge base to store information
  - Explain actions based on knowledge
    - E.g., medical diagnosis system infers diseases, suggests treatments
    - E.g., chess program uses move database to plan strategy
  - Handle incomplete or uncertain information through probabilistic reasoning

# Logic / Knowledge Base (1/2)

---

- **Knowledge base (KB)** is a set of:
  - **Sentences**  $\alpha$  expressing assertions (observed, assumed or derived) about the world
    - E.g., “it rains”, “the ground is dry”, “the ground is wet”
  - **Rules**
    - E.g., “If it rains, the ground gets wet”
- **Knowledge representation language** is a formal way of creating sentences about the world
- **Syntax** specifies all the sentences  $\alpha$  that are well-formed in a logic / knowledge base
  - E.g., in arithmetic the sentence:
    - “ $x + y = 4$ ” is well-formed
    - “ $x4y+ =$ ” is not well-formed
- **Semantics** is the meaning of sentences (i.e., their truth) with respect to each possible world
  - E.g., the sentence  $x + y = 4$ 
    - Is true in the world (model) in which  $x = 2, y = 2$
    - Is false in the world  $x = 1, y = 1$

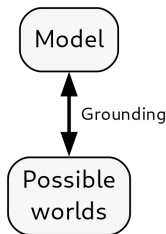
## Logic / Knowledge Base (2/2)

---

- **Axiom** is a sentence taken as given
  - Not derived from other sentences
- **Inference** is the process of deriving new sentences from old ones
  - It should be done in a “logical” way
- **Truth values** of a sentence
  - In most logics every sentence is either true or false
  - **Fuzzy logic** allows sentences to have different degrees of truth
    - $\text{Belief}(\alpha) = 0.5$
  - **Probabilistic logic** allows sentences to have different probability of being true
    - $\text{Pr}(\alpha) = 0.3$

# Model and Possible Worlds: Examples

- **Example:** world where there is rain and wet ground
  - In each possible world/model, values are assigned to all relevant variables
  - “Possible worlds” can be thought of as real the environments
  - Model  $m$  is a mathematical abstraction of “possible world”
    - E.g.,  $m$  is  $(Rain = F, WetGround = T)$
  - Each possible world is a complete assignment of truth values to all relevant propositions
    - World 1:  $(Rain = T, WetGround = T)$
    - World 2:  $(Rain = T, WetGround = F)$
    - World 3:  $(Rain = F, WetGround = T)$
    - World 4:  $(Rain = F, WetGround = F)$
- **Example:** represent worlds with “men and women sitting at a table”
  - Model represents all possible worlds as  $(x \text{ men}, y \text{ women})$
  - Sentence  $x + y = 4$  is true in certain worlds, false in others
  - In worlds with  $x = 2$  men and  $y = 2$  women,  $x + y = 4$  is true



# Satisfaction of a Sentence in a Model

---

- A model  $m$  fixes all the variables  $x_1, \dots, x_n$  used in sentences
  - E.g.,  $(Rain = T, WetGround = T)$
- If a sentence  $\alpha$  is true in model  $m$ , we say “the model  $m$  **satisfies** the sentence  $\alpha$ ”
  - E.g., the model  $(Rain = T, WetGround = F)$  satisfies  $\alpha : Rain = T$
  - Note: this seems backwards, since in our common way of reasoning, the world is fixed and sentences are evaluated as true or false
- $M(\alpha)$  is the set of all the models in which  $\alpha$  is true
  - E.g.,
    - $\alpha : Rain = T$
    - $M(Rain = T) = \{(Rain, WetGround), (Rain, \neg WetGround)\}$



# Logical Entailment

---

- **Logical entailment** between sentences is the fact that a sentence follows logically from another sentence in a KB
- “ $\alpha$  **entails**  $\beta$ ” (written  $\alpha \models \beta$ ) iff (by def) in every model in which  $\alpha$  is true,  $\beta$  is also true
  - Equivalent to  $M(\alpha) \subseteq M(\beta)$
- E.g., in the “rain and wet ground” world
  - $\alpha : \text{“Rain} \implies \text{WetGround”}$  entails  $\beta : \text{“}(Rain = T, WetGround = T)\text{”}$
- E.g., in the “sitting table” world
  - $\alpha: \text{“}x = 0\text{”}, \beta: \text{“}x \cdot y = 0\text{”}$
  - $\alpha$  entails  $\beta$  since in any model in which  $x = 0$  is true, also  $x \cdot y = 0$  is true, regardless of the value of  $y$
- **Intuition:**
  - Entailment is not related to a proof, it just “preserves truth” across all models
  - “If you believe your KB, you must believe the entailed sentences”

# Logical Entailment vs Implication

---

- Entailment and implication are related but distinct
  - **Logical entailment** is about truth following from known facts
  - **Implication** is about a relationship between two statements
- **Logical entailment** ( $KB \models \alpha$ ):
  - Means  $\alpha$  is always true in any world where KB is true
  - E.g.,
    - KB: “It is raining”, “If it rains, the ground is wet”
    - Entailed: “The ground is wet”
- **Implication** ( $A \implies B$ ):
  - A statement in logic that says: “If A is true, then B is true”
    - Doesn’t guarantee A or B is true by itself
    - Implication is true unless A is true and B is false
  - E.g.,
    - A: “It is raining”, B: “The ground is wet”
    - $A \implies B$  is the statement “If it is raining, then the ground is wet”
    - This statement can be true even if it’s not raining
- **Intuition:**
  - Entailment is “meta-level truth-following”
  - Implication is “within the logic”

# Model Checking Procedure

---

- $M(KB)$  represents all the models / possible worlds that are true given our  $KB$
- **Problem:**
  - We want to verify whether “a sentence  $\alpha$  is entailed by  $KB$ ” ( $KB \models \alpha$ )
- **Solution:**
  - According to the definition, we need to verify that  $\alpha$  is true in all the models in which  $KB$  is true
    - I.e.,  $M(KB) \subseteq M(\alpha)$
- E.g., **model checking procedure** (brute force)
  1. Enumerate all the models / possible worlds
  2. Find which models are possible given the  $KB$ , i.e.,  $M(KB)$
  3. Check whether the sentence  $\alpha$  is true in all the models that are compatible with the  $KB$

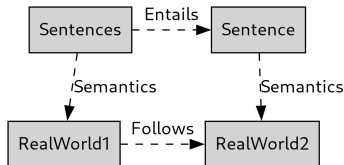
# Sound and Complete Inference Algorithm

---

- Inference: a syntactic process of deriving new sentences from others, using formal rules of a proof system (e.g., modus ponens, resolution, etc.)
  - You know: "If it rains, the ground gets wet."
  - You see: "It is raining."
  - You infer: "The ground must be wet."
- The ideal inference algorithm is both sound and complete
- **Sound** inference algorithm
  - Derives only sentences entailed from  $KB$
  - "Whatever the inference algorithm finds, it's correct", i.e., no false positives
  - E.g., model checking is sound
    - It works only when the space of models is finite
    - When it works, it is truth preserving
- **Complete** inference algorithm
  - Can derive any sentence entailed from  $KB$
  - "The inference algorithm doesn't miss anything," i.e., no false negatives

# Isomorphism between Model and Possible Worlds

- A sound and complete inference algorithm should yield conclusions guaranteed to be true in any world where the premises are true
  -
- In other words, even if the inference operates on “syntax” (the internal representation):
  - “Sentences in the representation” correspond to “aspects of the real world”
  - “Entailment between sentences in the representation” corresponds to “implication between aspects of the real world”



# Entailment vs Inference vs Implication

---

- **Logical entailment**

- $A$  entails  $B$ : if the fact  $A$  is true and that automatically guarantees that fact  $B$  must also be true
- E.g., *Rain* entails *WetGround* iff in every possible world where *Rain* is true, *WetGround* is also true
  - $Rain = T, WetGround = F \rightarrow$  violation
  - Since there is at least one counterexample, *Rain* does not entail *WetGround*

- **Inference**

- This is what you (a person or a computer) figure out based on what you know
- You start with some truths, then reason your way to new truths
- It's "reasoning inside the logic system"

- **Implication**

- "If  $A$ , then  $B$ "
- It doesn't say whether  $A$  is true; it just says, if it happens, then  $B$  follows
- It's a "statement inside the logic system"

# Grounding

---

- **Grounding** is the operation of linking abstract symbols to reality
  - E.g., words, variables in the representation ... → objects, entities, or situations in the real world
  - It is the bridge between representation in a KB and the world
- How **can we know** that a KB accurately reflects the real world?
  - We can't be sure!
  - Do we live in a simulation? What is reality?
- We **assume** that is correct
  - Agent's sensors create a sentence in the KB when something happens in the real world  
IF smell = burning THEN food\_is\_burning
  - Agent learns rules and acts  
IF food\_is\_burning THEN turn\_off\_stove
- We assume that “learning” (going from particular cases to general cases) is typically correct
  - Learning is still fallible
  - E.g., smell = burning because maybe somebody is cooking on a grill

- Knowledge Representation
  - Basics of Knowledge Representation
  - Examples of Logic
  - Logical Agents
  - *Ontologies*
  - Reasoning in Ontologies
- Propositional logic
- First-order Logic
- Non-classical Logics



# Ontologies (in computer science)

---

- **Ontology:**

- Is a formal, explicit representation of a domain
- Describes the types of things that exist and how they relate to each other
  - Classes: types of things
  - Individuals: specific objects
  - Properties: how things are related

- **Examples:**

- A medical ontology defines relationships between diseases, symptoms, and treatments
- A geographical ontology describes cities, states, and countries
- Semantic web (an extension of the current web to give meaning to information)

- **Goal**

- Provide a vocabulary for a domain of knowledge
- Enable machines and humans to understand and share information consistently
- Enable reasoning about entities and their relationships

- **Related Concepts**

- Schema: database-oriented structure, often more rigid than ontologies
- Taxonomy: simpler hierarchical tree-like classification
- Knowledge base: a collection of facts and rules, sometimes built from an ontology

# Ontologies: Components


---

- **Classes / Concepts:**
  - Represent general concepts in a domain
  - E.g., Person, City, Car
- **Individuals / Instances:**
  - Specific, concrete examples of classes
  - E.g., GP (an instance of Person), Rome, Ferrari 458
- **Properties / Relations:**
  - Describe interactions or associations between classes or instances
  - E.g., isMortal, locatedIn, hasAge
- **Attributes / Data values**
  - Specify data associated with instances
  - E.g., (GP, hasAge, <your\_guess>)
- **Constraints**
  - Rules that restrict the kinds of values a property can take
  - E.g., (Ferrari 458, mustBe, red)
- **Axioms:**
  - Logical statements that define rules and constraints
  - E.g., all humans are mortal:  $\forall x (Person(x) \implies Mortal(x))$
- **Hierarchies:**
  - Organize classes and properties into parent-child relationships
  - E.g., Student is a subclass of Person

# Ontology: Example University

- **Classes**

- Student
- Professor
- Course
- Department



- **Properties:** relationships between Classes

- takesCourse (Student  $\rightarrow$  Course)
- teachesCourse (Professor  $\rightarrow$  Course)
- belongsToDepartment (Student, Professor  $\rightarrow$  Department)

- **Individuals:** examples of Classes

- Student: Alice, Bob
- Professor: GP, DrNo
- Course: DATA605, MSML610

- **Axioms:** logical rules that must be true

- Every Course must be taught by exactly one Professor
- Every Student must belong to exactly one Department

- Knowledge Representation
  - Basics of Knowledge Representation
  - Examples of Logic
  - Logical Agents
  - Ontologies
  - *Reasoning in Ontologies*
- Propositional logic
- First-order Logic
- Non-classical Logics

# Example of Reasoning Tasks (1/4)

---

- **Subsumption**

- “Is class A a subclass of B?”
- Check whether one concept is more general than another
- E.g., if Person subsumes Student, every Student is necessarily a Person
- Important for building taxonomies and ontologies

- **Satisfiability**

- “Can an instance of a concept exist?”
- Test if a concept is logically consistent (i.e., without contradiction)
- E.g., if the concept FlyingPenguin requires flying but is also defined as a penguin (which cannot fly), it might be not satisfiable

- **Classification**

- Organize concepts into a hierarchy
- Automatically organize concepts into a hierarchy by checking subsumption relationships
- E.g., given definitions of Animal, Bird, and Penguin, classification places Penguin under Bird, and Bird under Animal

## Example of reasoning tasks in KR (2/4)

---

- **Instance Checking**
  - “Is a specific individual an instance of a concept?”
  - E.g., is GP an instance of Student?
- **Consistency Checking**
  - “Is the entire knowledge base free of contradictions?”
  - E.g., no Person is both Alive and Dead at the same time
- **Realization**
  - “What is the most specific class an instance belongs to?”
  - E.g., discovering that GP is a Professor rather than just a Human
- **Retrieval**
  - Find all individuals that satisfy a certain condition
  - E.g., retrieve all instances classified as TeachingAssistant

## Example of reasoning tasks in KR (3/4)

---

- **Query Answering**
  - Answer complex queries about the knowledge base
  - E.g., “Find all Person that study at the university and are not Student”
- **Abduction**
  - Given an observation, infer the best explanation
  - E.g., seeing a Person carrying a backpack and wearing flip-flops in the snow and infer that is likely a Student
- **Deduction**
  - Infer consequences that logically follow from facts and rules
  - E.g., if John is a Student in ComputerScience then he can attend MSML610

# Example of reasoning tasks in KR (4/4)

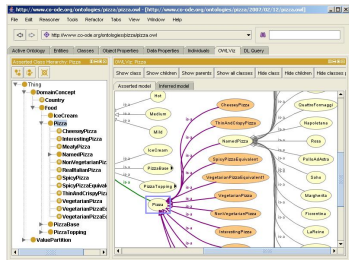
---

- **Belief Revision**
  - Update the knowledge base when new, possibly conflicting, information arrives
  - E.g., learning that not every student in ComputerScience can take MSML610 and revise a previous rule
- **Temporal Reasoning**
  - Reason about events over time
  - E.g., If EventA happens before EventB, then EventB cannot Cause EventA
- **Causal Reasoning**
  - Infer causes and effects among entities or events
  - E.g., inferring that (Storm, Cause, Flooding) based on temporal and physical knowledge



# Ontologies tools: Protege Example

- **Protégé** is a free, open-source platform for building ontologies
  - Developed at Stanford
- **Construct and visualize ontologies**
  - Users can define classes, properties, individuals, and relationships
- **Enable reasoning** over ontologies using plugins
  - E.g., checking consistency, inferring new knowledge
- **Supports**
  - Major ontology languages
    - OWL (Web Ontology Language)
    - RDF (Resource Description Framework)
  - Multiple serialization formats
    - RDF/XML, Turtle, OWL Functional Syntax
- **Use cases**
  - Domain-specific knowledge modeling (e.g., biomedicine, law)
  - Semantic Web applications
  - AI systems that require structured knowledge



- Knowledge Representation
- ***Propositional logic***
  - Syntax
  - Semantics
- First-order Logic
- Non-classical Logics

- Knowledge Representation
- Propositional logic
  - **Syntax**
  - Semantics
- First-order Logic
- Non-classical Logics

# Propositional Logic

---

- **Propositional logic** is a formal system for reasoning about statements that can be true or false
  - **Syntax** defines the allowable sentences
    - Consists of proposition symbol and logical connectives
    - E.g.,  $P \wedge Q$
  - **Semantics** is the way in which the truth of sentences is determined
    - Truth tables or deduction rules evaluate the truth value of complex sentences
    - E.g., if  $P$  is true and  $Q$  is false then  $P \wedge Q$  is false
- **Atomic representation**
  - No internal structure within atomic propositions
- **Uses**
  - SAT solvers
    - Tools for determining if a propositional logic formula can be satisfied
    - E.g., used in hardware verification and scheduling problems
  - Expert systems
    - Use logic rules to mimic human decision-making
    - E.g., medical diagnosis systems
  - Rule-based agents
    - Agents that operate based on a set of predefined rules
    - E.g., automated customer service chatbots

# Proposition Symbol

---

- **Proposition symbol**
  - Is an atomic sentence consisting of a single symbol
    - E.g.,  $P$ ,  $Q$ ,  $North$
  - Doesn't have truth value, it is just a symbol for a real-world statement
    - Needs grounding
  - Stands for a proposition that can be true or false
    - E.g.,  $K_{E,5}$  = "the Knight is in E5"
    - $K_{E,5}$  is not composed of any other symbol, it is an atomic symbol
- *True* and *False* are proposition symbols with inherent truth values

# Sentences

---

- **Atomic sentence**
  - Is a sentence composed of a single proposition symbol
  - E.g.,  $P$
- **Complex sentence**
  - Is constructed from simpler (sentences) using parentheses and logical connectives
  - E.g.,  $(P \wedge Q) \vee R$
  - It is a recursive definition that allow to build more complex sentences
- Common logical connectives
  - Not:  $\neg$
  - And:  $\wedge$  (looks like an “A” for “and”)
  - Or:  $\vee$  (comes from Latin “vel” which means “or”)
  - Implies:  $\implies$
  - If and only if:  $\iff$
- Each sentence (atomic or complex) can be only true or false

# Proposition Logic: Weather Example

---

- Proposition symbols are:
  - *Rain* = “it’s raining”
  - *Cold* = “it’s cold”
  - *Sunny* = “it’s sunny”
  - *Snow* = “it’s snowing”
  - *Cloudy* = “it’s cloudy”
- **Atomic sentence** can be positive or negated
  - E.g., *Rain*,  $\neg \textit{Rain}$  = “it’s not raining”
- **Negation**
  - E.g.,  $\neg(\textit{Rain} \vee \textit{Cloudy})$  = “it’s not the case that it’s raining or cloudy”
- **Conjunction / Disjunction**
  - E.g.,  $\textit{Rain} \wedge \textit{Cold}$  = “it’s raining and it’s cold”
  - E.g.,  $\textit{Rain} \vee \textit{Snow}$  = “it’s either raining or snowing”
- **Implication** is a sentence containing a premise  $\implies$  conclusion
  - Aka “if-then statements”, “rules”
  - E.g.,  $\textit{Rain} \implies \neg \textit{Snow}$  = “if it’s raining, it’s not snowing”
- **Biconditional**:  $A \implies B \wedge B \implies A$ 
  - E.g.,  $\textit{Sunny} \iff \neg \textit{Cloudy}$  = “it’s sunny if and only if it’s not cloudy”

# Grammar in BNF form

---

- **Backus Normal Form** formally represents the grammar of propositional logic
- Ambiguous, i.e., the same sentence can be parsed in multiple ways
  - E.g.,  $\neg A \vee B = (\neg A) \vee B$  or  $\neg(A \vee B)$  ?
- To eliminate ambiguity define the precedence for each operator
  - E.g.,  $\neg$  has higher precedence than  $\wedge, \vee$  so:  
 $\neg A \vee B$  means  $(\neg A) \vee B$

```
Sentence → AtomicSentence | ComplexSentence
AtomicSentence → True | False | P | Q | R | ...
ComplexSentence → ( Sentence )
                  | ¬ Sentence
                  | Sentence ∧ Sentence
                  | Sentence ∨ Sentence
                  | Sentence ⇒ Sentence
                  | Sentence ⇔ Sentence
```

OPERATOR PRECEDENCE :  $\neg, \wedge, \vee, \Rightarrow, \Leftrightarrow$



- Knowledge Representation
- Propositional logic
  - Syntax
  - *Semantics*
- First-order Logic
- Non-classical Logics

# Semantics of Propositional Logic

---

- **Semantics** are rules for determining the truth of a sentence  $\alpha$  with respect to a model  $m$ 
  - Determine if a sentence is true or false, given a possible world
- In propositional logic, a model  $m$  fixes the truth value (true or false) for every proposition symbol/atomic sentence, e.g.,
- The models are **abstractions** of the real world and have no a-priori connection to a specific world, e.g.,
  - $P_{1,2}$  is just a symbol and can mean:
    - “There is a pit in  $[1, 2]$ ” or
    - “I’m in Paris today and tomorrow”
  - Need grounding

# Computing the Truth Value of a Sentence

---

- The **truth value of a sentence** is derived from the truth of the proposition symbols (recursively from the model  $m$ )

$$m = \{P_{1,2} = F, P_{2,2} = F, P_{3,1} = T\}$$

- All sentences  $\alpha$  are constructed from atomic sentences (assigned by model  $m$ ) and connectives:
  - $\neg P$  is T iff  $P$  is F in  $m$
  - $P \wedge Q$  is T iff  $P$  and  $Q$  are both true in  $m$
  - $P \vee Q$  is T iff  $P$  or  $Q$  are true in  $m$
  - $P \implies Q$  is true unless  $P$  is true and  $Q$  is false in  $m$
  - $P \iff Q$  is true iff  $P$  and  $Q$  are both true or both false in  $m$
- **Truth table** contains truth value of a sentence for each possible model  $m$ 
  - E.g.,  $X = A \wedge B \vee C$

A	B	C	X
F	F	F	F
F	F	F	T
...			

# Interpretation of Implication

---

- In a **logical implication** " $P \implies Q$ " there is **no causation** between  $P$  and  $Q$
- " $P \implies Q$ " says *"If  $P$  is true, I claim that  $Q$  is true. Otherwise I am making no claim"*
  - E.g., *"5 is odd implies that Tokyo is the capital of Japan"* is a true sentence in propositional logic (although very odd)
- **Pathological cases** for implication
  - An implication is true whenever the antecedent is false
  - E.g., *"5 is even implies pigs fly"* is true

# Model Checking is Sound and Complete

---

- **Model checking algorithm**
  - Enumerate all the models (truth tables)
  - Check if  $\alpha$  is true for every model where  $KB$  is true
- The model checking algorithm is:
  - **Sound**
    - “Any inference made by the algorithm is correct”
    - Implements the definition of entailment
  - **Complete**
    - “Any true sentence is inferred correctly by the algorithm”
    - Works for any  $KB$  and  $\alpha$
    - Always terminates (finite number of models)
- Complexity of model checking with  $n$  variables
  - Time complexity is  $O(2^n)$  (NP-complete)
    - Worst case is exponential
    - Average case is better than exponential
  - Space complexity is  $O(n)$  since enumeration is depth-first

# Inference in Propositional Logic

---

- Inference are the rules of reasoning
  - **Modus Ponens**: if  $p \implies q$  and  $p$ , infer  $q$ 
    - If it rains, the ground will be wet. It rains.
    - Therefore, the ground is wet.
  - **Modus Tollens**: if  $p \implies q$  and  $\neg q$ , infer  $\neg p$ 
    - If it rains, the ground will be wet. The ground is not wet.
    - Therefore, it did not rain.
  - **Syllogism (Transitivity)**
    - If  $p \implies q$  and  $q \implies r$ , then  $p \implies r$
  - **Disjunctive Syllogism**
    - If  $p \vee q$  and  $\neg p$ , infer  $q$
  - **Addition**: if  $p$ , then  $p \vee q$
  - **Simplification**: from  $p \wedge q$ , infer  $p$  (or  $q$ )
  - **Conjunction**: from  $p$  and  $q$ , infer  $p \wedge q$
  - **Resolution Rule**
    - From  $(p \vee q)$  and  $(\neg p \vee r)$ , infer  $(q \vee r)$

# Propositional Theorem Proving

---

- To **prove** a desired sentence  $\alpha$  under a knowledge base KB
  - Apply rules of inference to construct a proof of  $\alpha$
  - Any sentence can have only one of the following truth values:
    1. True
    2. False
    3. Undecidable under the KB
- **Theorem proving vs model checking**
  - Model checking involves enumerating all models to show the sentence is true/false in all models where KB is true
  - Propositional theorem proving builds a proof
  - If the proof is short, theorem proving can be more efficient than model checking

# Logical Equivalence of Sentences

---

- Two sentences  $\alpha$  and  $\beta$  are logically equivalent  $\alpha \equiv \beta$ 
  - Iff they are true in the same set of models:

$$M(\alpha) = M(\beta)$$

- Iff they entail each other:

$$\alpha \models \beta \wedge \beta \models \alpha$$

- E.g.,  $P \vee Q \equiv Q \vee P$



# Logical Equivalences (1/2)

---

- Commutativity of  $\wedge$  and  $\vee$

$$(\alpha \wedge \beta) \equiv (\beta \wedge \alpha)$$

$$(\alpha \vee \beta) \equiv (\beta \vee \alpha)$$

- Associativity of  $\wedge$  and  $\vee$

$$(\alpha \wedge \beta) \wedge \gamma \equiv \alpha \wedge (\beta \wedge \gamma) \equiv \alpha \wedge \beta \wedge \gamma$$

$$(\alpha \vee \beta) \vee \gamma \equiv \alpha \vee (\beta \vee \gamma) \equiv \alpha \vee \beta \vee \gamma$$

- Distributivity of  $\wedge$  over  $\vee$

$$\alpha \wedge (\beta \vee \gamma) \equiv (\alpha \wedge \beta) \vee (\alpha \wedge \gamma)$$

- Distributivity of  $\vee$  over  $\wedge$

$$\alpha \vee (\beta \wedge \gamma) \equiv (\alpha \vee \beta) \wedge (\alpha \vee \gamma)$$

- Double negation elimination:

$$\neg(\neg\alpha) \equiv \alpha$$

## Logical Equivalences (2/2)

---

- Contraposition:

$$(\alpha \implies \beta) \equiv (\neg\beta \implies \neg\alpha)$$

- Implication elimination:

$$(\alpha \implies \beta) \equiv (\neg\alpha \vee \beta)$$

- Biconditional elimination:

$$(\alpha \iff \beta) \equiv (\alpha \implies \beta) \wedge (\beta \implies \alpha)$$

- De Morgan equivalence:

$$\neg(\alpha \wedge \beta) \equiv (\neg\alpha \vee \neg\beta)$$

$$\neg(\alpha \vee \beta) \equiv (\neg\alpha \wedge \neg\beta)$$

# Deduction theorem

---

- A **valid sentence**  $\alpha$  is true for all the models
  - E.g.,  $P \vee \neg P$
  - Aka “tautology”
  - Every tautology is equivalent to the sentence *True*
- **Contradiction** is a sentence  $\alpha$  that is false for all the models
  - E.g.,  $P \wedge \neg P$
  - Every contradiction is equivalent to the sentence *False*
- **Deduction theorem**
  - The sentence  $\alpha$  entails  $\beta$  (written  $\alpha \models \beta$ ) iff the sentence  $\alpha \implies \beta$  is a tautology
- The deduction theorem is like a **bridge** between two different but closely related ideas:
  - Entailment  $\models$  is a semantic notion
    - It's about truth across all models
    - $\alpha \text{ entails } \beta$  means in every possible world where  $\alpha$  is true,  $\beta$  is also true
  - Implication  $\implies$  is a syntactic notion
    - $\alpha \implies \beta$  is just another formula inside the logic

# Satisfiability

---

- A sentence  $\alpha$  is **satisfiable** iff  $\alpha$  is true for some model
- **SAT problem** is about determining satisfiability of sentence in propositional logic
  - One can enumerate all the possible models until one is found to satisfy the sentence  $\alpha$
  - It is NP-complete
- A sentence  $\alpha$  is **un-satisfiable** iff  $\alpha$  is never true (i.e., a contradiction)
- **Validity and satisfiability**
  - $\alpha$  is valid (i.e., a tautology) iff  $\neg\alpha$  is un-satisfiable
  - By contrapositive,  $\alpha$  is satisfiable iff  $\neg\alpha$  is not valid ( $\neg\alpha$  is not a tautology)

# Proof by Contradiction

---

- The sentence  $\alpha \models \beta$  is true iff the sentence “ $\alpha \vee \neg\beta$ )” is un-satisfiable (i.e., a contradiction)
- In other words in a proof by contradiction:
  - Assume  $\alpha$
  - Assume that the sentence  $\beta$  is false
  - Prove that this leads to a contradiction
  - Thus  $\beta$  must be true

# Propositional Logic

---

- E.g.,  $P \wedge Q$
- **Pros**
  - Declarative
    - Semantics is based on relation between sentences and possible worlds
  - Can deal with partial information
    - E.g., “A white knight is in b1 or in f6” is represented with  $WK1_{b1} \vee WK2_{f6}$
  - Compositional semantics
    - The meaning of a sentence is a function of the meaning of its parts
  - Context independent
  - Unambiguous
- **Cons**
  - Can't concisely describe environment with many objects, e.g.,
    - In English “*The pawn is in a cell around b1*” requires all the possible states to be enumerated
  - Can't represent uncertainty
    - “*There is 50% of probability that the pawn is in b1*”

- Knowledge Representation
- Propositional logic
- ***First-order Logic***
  - Syntax
  - Semantics
- Non-classical Logics

- Knowledge Representation
- Propositional logic
  - *Syntax*
  - Semantics
- First-order Logic
- Non-classical Logics



# First-Order Logic (FOL)

---

- First-order logic (FOL) extends propositional logic by:
  - Introducing **quantifiers** (e.g.,  $\forall$ ,  $\exists$ )
  - Using **predicates** to represent properties and relations
- Combines pros of propositional logic with pros of natural language
  - Built around objects and relations
  - Allows to express facts about some or all objects, e.g.,
    - *"Some humans have green eyes"*
    - *"Chess pieces around the Queen are at risk"*
- FOL provides expressive power to represent structured, relational knowledge

# First-Order Logic: Syntax

- **Constants**: represent specific objects
  - E.g., *Socrates*
- **Predicates**: describe properties or relations
  - E.g.,  $Human(x) = "x \text{ is human}"$
- **Functions**: map tuples of objects to objects
  - E.g.,  $Mother(x) = "the \text{ mother of } x"$
- **Variables**: act as placeholders
  - E.g.,  $x, y$
- **Quantifiers**
  - E.g.,  $\forall x$  (for all  $x$ ),  $\exists x$  (there exists an  $x$ )

```
Sentence → AtomicSentence | ComplexSentence
AtomicSentence → Predicate | Predicate(Term,...) | Term = Term
ComplexSentence → ( Sentence )
                | ¬ Sentence
                | Sentence ∧ Sentence
                | Sentence ∨ Sentence
                | Sentence ⇒ Sentence
                | Sentence ⇔ Sentence
                | Quantifier Variable,... Sentence
```

```
Term → Function(Term,...)
      | Constant
      | Variable
```

```
Quantifier → ∀ | ∃
Constant → A | X1 | John | ...
Variable → a | x | s | ...
Predicate → True | False | After | Loves | Raining | ...
Function → Mother | LeftLeg | ...
```

OPERATOR PRECEDENCE :  $\neg, =, \wedge, \vee, \Rightarrow, \Leftrightarrow$

# Sentences

---

- **Term** is a logical expression that refers to an object
  - E.g., *Richard*
- **Atomic sentence** is a predicate symbol (i.e., corresponds to relations) followed by a list of terms in parenthesis *Predicate(Term1, Term2, ...)*
  - E.g., *Brother(Richard, John)* means “*Richard is the brother of John*”, under a given grounding / interpretation,
  - E.g., *Married(Father(Richard), Mother(John))* means “*the father of Richard and the mother of John are married*”
- **Complex sentence** is a sentence using logical connectives complex, with the same syntax and semantics as in propositional logic
- **Variable** is a term that represents a possible object
  - Typically represented as lowercase letter (e.g.,  $x$ ,  $y$ ,  $z$ )
  - Can be used as argument of a function, e.g., *LeftLeg*( $x$ )
- **Equality symbol** signifies that two terms refer to the same object
  - E.g., *Father(John) = Henry*

# Quantifiers

---

- **Quantifiers** express properties of entire collections of objects, instead of enumerating objects by name (like in propositional logic)
- **Universal quantifier**:  $\forall x P(x)$ 
  - Universal quantifier makes a statement about *every* object
  - Statement is true if  $P(x)$  is true for all  $x$
- **Existential quantifier**:  $\exists x P(x)$ 
  - Existential quantifier makes a statement about *some* object (without naming it)
  - Statement is true if  $P(x)$  is true for at least one  $x$
- Variables are **bound** by quantifiers or **free** (unbound)
  - E.g.,  $\forall x (Cat(x) \rightarrow Mammal(x))$

# Nested Quantifiers

---

- More complex sentences can be expressed using **multiple quantifiers**
  - The order of quantifiers is important, so one can use parentheses to clarify
- **Example:**
  - “*Brothers are siblings*”:  $\forall x, y \text{ Brother}(x, y) \implies \text{Sibling}(x, y)$
  - Siblings is a symmetric relationship:  $\forall x, y \text{ Sibling}(x, y) \iff \text{Sibling}(y, x)$
  - “*Everybody loves somebody*”:  $\forall x \exists y \text{ Loves}(x, y)$
  - “*There is someone loved by everyone*”:  $\exists y \forall x \text{ Loves}(x, y)$

## Connection between $\forall$ and $\exists$

---

- The universal and existential quantifiers are **connected** through negation and De Morgan rules

$$\forall x \neg P(x) \iff \neg \exists x : P(x)$$

$$\neg(\forall x P(x)) \iff \exists x \neg P(x)$$

$$\forall x P(x) \iff \neg \exists x \neg P(x)$$

$$\exists x P(x) \iff \neg(\forall x \neg P(x))$$

- Knowledge Representation
- Propositional logic
  - Syntax
  - *Semantics*
- First-order Logic
- Non-classical Logics

# First-order Logic: Semantics

---

- **Semantics** define how sentences are interpreted in a domain
- Symbols represent entities, relationships, and functions in the domain
  - **Constant symbols** represent specific objects
    - E.g., *Alice*, *GP*, *CS101*
  - **Predicate symbols** represent relationships among objects
    - E.g., *EnrolledIn(Student, Class)*, *Teaches(Professor, Class)*, *IsStudent(x)*, *IsProfessor(x)*
  - **Function symbols** represent mappings between objects
    - E.g., *AdvisorOf(Student)*, *DepartmentOf(Professor)*
- An **interpretation** maps the world to its mathematical description, and vice versa
  - There are many possible interpretations
  - The intended interpretation is the one that is the most natural
  - E.g., map the symbol *GP*  $\rightarrow$  me
  - $\forall x (Human(x) \rightarrow Mortal(x))$  means that every human is mortal



# Representing Knowledge in FOL

---

- FOL enables representation of:
  - General rules:  $\forall x (Bird(x) \rightarrow CanFly(x))$
  - Specific facts:  $Bird(Tweety)$
- Complex relations captured through predicates:
  - $Loves(Romeo, Juliet)$ ,  $GreaterThan(3, 2)$
- Functions express object construction:
  - $FatherOf(John)$
- Knowledge base is built from axioms and facts
- Enables reasoning about objects, properties, and their relationships

# First-Order Logic: Inference

---

- **Goal:** derive new sentences from existing ones using sound rules
  - **Universal Instantiation:**
    - From  $\forall x P(x)$  infer  $P(c)$  for any constant  $c$
  - **Existential Instantiation:**
    - From  $\exists x P(x)$  infer  $P(c)$  with a new constant  $c$
  - Standard propositional logic rules such as Modus Ponens, Modus Tollens, Resolution, etc., also apply within FOL
    - Modus ponens: from  $P$  and  $P \implies Q$ , infer  $Q$
- FOL inference is semi-decidable:
  - If a sentence is entailed, a proof can be found
  - If not entailed, proof search may not terminate

- Knowledge Representation
- Propositional logic
- First-order Logic
- ***Non-classical Logics***
  - Intro and Examples
  - Description Logics
  - Semantic Web

- Knowledge Representation
- Propositional logic
- First-order Logic
- Non-classical Logics
  - *Intro and Examples*
  - Description Logics
  - Semantic Web

# Ontological Commitment

---

- **Ontological commitments** are assumptions about reality made by a language
- Different formal models make **different assumptions** on how the truth of sentences is defined
  - Propositional logic:
    - The world consists of facts that are either true or false
    - E.g.,  $P \vee Q$
  - First-order logic:
    - The world consists of objects with relations among them that hold or do not hold
    - $\forall x : \text{Human}(x) \implies \text{Mortal}(x)$
  - Temporal logic:
    - Facts about objects and relations hold at particular times or intervals
  - Higher-order logic:
    - Relations of first-order logic are objects themselves
    - E.g., can make assertions about relations (e.g., “all relations are transitive”)

# Epistemological Commitment

---

- **Epistemological commitment** is a possible states of knowledge by an agent with respect to each fact (i.e., belief), e.g.,
  - Propositional logic, first-order logic
    - 3 possible states of belief regarding any sentence: true, false, or unknown
  - Probability theory
    - There is a degree of belief in  $[0, 1]$  about each sentence
    - $\Pr(X = 6) = 0.3$
- Ontological commitment = what exists in the world
- Epistemological commitment = what an agent believes about facts

# Non-monotonic Logic

---

- **Non-monotonic logic** is a type of logic where adding new information can invalidate previous conclusions
- Monotonic (classical) vs non-monotonic logic
  - In classical logic, once something is proven, it stays proven even if more information is added
  - In non-monotonic logic, conclusions can change as new facts are learned
- E.g.,
  - Initial knowledge: *"Birds typically fly"*
  - Fact: *"Tweety is a bird"*
  - Conclusion: *"Tweety can fly"*
  - New fact: *"Tweety is a penguin"*
  - Fact: *"penguins are birds that can't fly"*
  - Revised conclusion: *"Tweety cannot fly"*
- E.g., "Swans are white birds" and then black swans are discovered
- Real-world situations often involve incomplete or evolving knowledge
  - Non-monotonic logic allows systems to reason flexibly and adapt to new circumstances

# Default reasoning

---

- **Default reasoning** is reasoning where assumptions are made by default in the absence of contrary evidence
  - It allows conclusions based on typical situations unless exceptions are found
  - Relates to “common sense”
- **Key idea**
  - Assume the most likely case unless specified otherwise
  - If new information contradicts the assumption, revise the conclusion
- **E.g.,**
  - Default rule: *“Typically, birds can fly”*
  - Fact: *“Tweety is a bird”*
  - Conclusion: *“Tweety can fly”*
  - New fact: *“Tweety is a penguin”*
  - Revised conclusion: *“Tweety cannot fly”*
- **Why it is useful**
  - Default reasoning allows systems to function reasonably without knowing everything



# Non-Monotonic Logic: University Example

---

- **Initial facts**
  - *Alice* is a *Student*
  - *Alice* belongs to the *ComputerScience* department
  - *CS101* is a *Course* offered by the *ComputerScience* department
  - Default rule: *Students* in the *ComputerScience* department take classes in their department
- **Initial reasoning**
  - Since *Alice* is a *Student* in *ComputerScience*, by default *Students* take *CS101*
  - Conclusion: *Alice takesCourse CS101*
- **New information**
  - *Alice* is a student who does not meet the prerequisites for *CS101*
- **Revised reasoning**
  - New conclusion: *Alice* does not *takeCourse CS101*

# Common Sense Reasoning

---

- **Common sense reasoning** is the ability to make assumptions, draw conclusions based on everyday knowledge about the world
  - Involves typical, unstated knowledge that humans take for granted, e.g.,
    - *"If you drop a glass, it will likely break"*
    - Knowing that *"people eat food when they are hungry"* without being explicitly told
- **Characteristics**
  - Deals with incomplete, uncertain, or ambiguous information
  - Relies on defaults, heuristics, and typical patterns rather than strict logical proofs
  - Often flexible and tolerant of exceptions
- **Challenges**
  - Common sense knowledge is vast, informal, and often not precisely defined
  - Difficult to encode all of it explicitly in a machine-readable form
  - Handling exceptions and contradictions is complex
- **Techniques**
  - Knowledge graphs
  - Non-monotonic logic
  - Probabilistic reasoning
  - Machine learning models trained on large, diverse data

# Common Sense Reasoning: University Example

---

- **Initial facts**
  - *Alice is a Student*
  - *Bob is a Student*
  - *CS101 is a Course offered by the ComputerScience department*
- **Common sense knowledge**
  - *Students typically Enroll in Courses offered by their Department*
  - *Students usually attend Course they are enrolled in*
  - *Professors usually teach the Courses they are assigned*
- **Reasoning steps**
  - *Alice belongs to the ComputerScience department*
  - *CS101 is offered by the ComputerScience department*
  - *Common sense suggests Alice is likely enrolled in CS101, even if enrollment is not explicitly stated*
  - *Therefore, it is reasonable to assume "Alice takesCourse CS101"*
- **New information**
  - *Alice is pursuing research only and not taking courses*
  - *The assumption that Alice takesCourse CS101 must be revised*

# Open World vs Closed World Assumptions

---

- **Closed World Assumption (CWA)**

- Missing information is false
- E.g.,
  - Fact: “*Alice takes CS101*” is known
  - Nothing is said about Bob
  - Under CWA: Conclude Bob does not take CS101

- **Open World Assumption (OWA)**

- Missing information is unknown, not false
- E.g.,
  - Fact: “*Alice takes CS101*” is known
  - Nothing is said about Bob
  - Under OWA: Can’t conclude if Bob takes CS101 or not it is unknown

- **Applications**

- CWA
  - Traditional relational databases (SQL)
  - Logic programming
  - Business rules and systems requiring complete data
- OWA
  - Semantic Web (RDF, OWL)
  - Knowledge representation with incomplete or growing data

# Inductive Logic Programming

---

- **Inductive logic programming**

- Learns logical rules from examples and commonsense knowledge
- Given positive and negative examples, and background facts, infer logical rules that explain the examples

- **Example**

- Commonsense knowledge: *"Birds have wings"*
- Positive example: *"Tweety is a bird that can fly"*
- Negative example: *"Penguin cannot fly"*
- Learned rule: *"Birds can fly unless they are penguins"*

- **Features**

- Produces human-readable logical rules
- Integrates learning with symbolic reasoning
- Supports background knowledge integration

- **Challenges**

- Computational complexity with large datasets
- Handling noisy, incomplete, or ambiguous data

- Knowledge Representation
- Propositional logic
- First-order Logic
- Non-classical Logics
  - Intro and Examples
  - *Description Logics*
  - Semantic Web

# Description Logic

---

- **Description logic**
  - Represents structured knowledge about a domain
  - Balances expressivity and computational efficiency
  - More expressive than propositional logic, less than first-order logic
- **Core building blocks:**
  - Classes: abstract groups
    - E.g., *Person*, *Animal*
  - Properties: binary relations between individuals
    - E.g., *hasChild*, *ownsPet*
  - Instances: specific objects
    - E.g., *GP*, *Nuvolo*
- Supports reasoning tasks such as:
  - Concept subsumption: *"Is class A a subset of class B?"*
  - Instance checking: *"Does instance a belong to class A?"*
- **Syntax often combines:**
  - Atomic concepts and roles
  - Logical constructors ( $\sqcap$ ,  $\sqcup$ ,  $\neg$ ,  $\forall$ ,  $\exists$ )
  - E.g.,
    - $Father \equiv Man \sqcap \exists hasChild. Person$
- Widely used in ontologies, e.g., OWL (Web Ontology Language)

- **Attributive Concept Language with Complements (ALC)** is a basic but expressive description logic
  - Concepts can be combined using logical operators, e.g.,
    - $\sqcap$  means “and”
    - $\sqcup$  means “or”
    - $\neg$  means “not”
  - Allows for existential and universal quantification, e.g.,  $\exists R.C$ ,  $\forall R.C$
  - Interpretation is set-theoretic
    - Classes as sets, properties as binary relations
- **Examples**
  - “All students take some course”:  $\text{Student} \equiv \exists \text{takes.Course}$
  - “A mother is a woman who has at least one child”:  
 $\text{Mother} \equiv \text{Woman} \sqcap \exists \text{hasChild.T}$
- ALC:
  - Is decidable
  - Balances expressiveness and computational complexity
  - Is the basis for more complex logics used in OWL



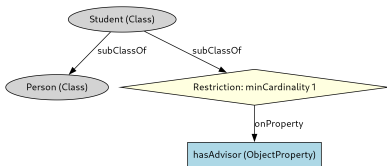
# SHOIN

---

- **SHOIN** is a description logic more expressive than ALC
- Components:
  - $\mathcal{S}$ : Allows transitive properties
    - E.g., *ancestorOf* is transitive
  - $\mathcal{H}$ : Supports role hierarchies
    - E.g., *hasSon*  $\sqsubseteq$  *hasChild*
  - $\mathcal{O}$ : Introduces specific individuals
    - E.g., *John* is a nominal class
  - $\mathcal{I}$ : Enables roles to be navigated backward
    - E.g., *isChildOf* is inverse of *hasChild*
  - $\mathcal{N}$ : Sets cardinality constraints
    - E.g., “*has exactly 1 children*”
- **Examples**
  - “*Exactly two children*”: *Person*  $\sqsubseteq$  ( $= 2$  *hasChild*. $\top$ )
- **Characteristics**
  - More powerful but reasoning is harder (exponential complexity)
  - Model richer real-world scenarios
  - Foundation for OWL DL reasoning capabilities

- **Web Ontology Language (OWL):**
  - “OWL” easier to pronounce than “WOL”
  - Semantic web language designed to represent complex knowledge about things and their relationships
  - Enables rich knowledge representation on the web (based on SHOIN)
  - Supports formal semantics for machine reasoning
  - Key constructs:
    - Classes, properties, individuals, axioms
- **Example**
  - “Every cat is a mammal”:  $Cat \sqsubseteq Mammal$
- **OWL variants:**
  - OWL Lite: simpler, for classification hierarchies
  - OWL DL: full expressiveness with decidable reasoning
  - OWL Full: maximum expressiveness, but undecidable
- **Applications**
  - Semantic search
  - Biomedical data

# Example of OWL in RDF



```
1.  ```\n2.  <rdf:RDF xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#" \n3.      xmlns:owl="http://www.w3.org/2002/07/owl#" \n4.      xmlns:rdfs="http://www.w3.org/2000/01/rdf-schema#" \n5.      xmlns:ex="http://example.org/">\n6.    <owl:Class rdf:about="http://example.org/Person"/>\n7.    <owl:Class rdf:about="http://example.org/Student">\n8.      <rdfs:subClassOf rdf:resource="http://example.org/Person"/>\n9.      <rdfs:subClassOf>\n10.        <owl:Restriction>\n11.          <owl:onProperty rdf:resource="http://example.org/hasAdvisor"/>\n12.          <owl:minCardinality rdf:datatype="http://www.w3.org/2001/XMLSchema#nonNegat\n13.            1\n14.          </owl:minCardinality>\n15.        </owl:Restriction>\n16.      </rdfs:subClassOf>\n17.    </owl:Class>\n18.    <owl:ObjectProperty rdf:about="http://example.org/hasAdvisor"/>\n19.  </rdf:RDF>\n20.  ```\n
```

# RDF (Resource Description Framework)

- **Resource Description Framework (RDF)** is a standard model for data interchange on the web
  - Represent structured information in a machine-readable way

- **Basic building block** is a triple:

- *Subject*: the entity being described, e.g., NuvoIo
- *Predicate*: the property or relationship, e.g., isA
- *Object*: the value or another entity, e.g., Dog

Subject	Predicate	Object
Book123	hasTitle	"The Great Gatsby"
Book123	hasAuthor	Author456
Author456	hasName	"F. Scott Fitzgerald"
Book123	publishedYear	"1925"
Book123	belongsToGenre	"Fiction"

- **Key Features:**
  - Statements are directed graphs of nodes and edges
  - Components are URIs (Uniform Resource Identifiers) to ensure global uniqueness or literals
    - E.g., `http://example.org/NuvoIo`
- **Use Cases:**
  - Building knowledge graphs
  - Enabling semantic search
  - Supporting ontologies (e.g., OWL)

# SPARQL

---

- SPARQL is the query language for RDF data
  - Allows users to retrieve and manipulate data stored in RDF format
- **Key Concepts:**
  - *Triple Patterns*: Query fragments that match triples in an RDF graph
  - *Basic Graph Pattern*: A set of triple patterns combined
  - *Variables*: Stand in for unknown parts of the triples
    - E.g., ?person, ?animal
- **Main Query Types:**
  - **SELECT**: Retrieve specific variables from the data
  - **CONSTRUCT**: Create new RDF triples based on query results
  - **ASK**: Return a boolean indicating whether a pattern exists
  - **DESCRIBE**: Return an RDF graph describing resources
- **Example:**
  - *"Find all resources that are of type Bird"*  
`SELECT ?animal WHERE { ?animal rdf:type ex:Bird }`

- Knowledge Representation
- Propositional logic
- First-order Logic
- Non-classical Logics
  - Intro and Examples
  - Description Logics
  - *Semantic Web*

# Semantic Web

---

- The **Semantic Web** extends the current Web by enabling machines to understand and interpret data
  - HTML is human-readable but lacks semantic structure for computers
  - The Semantic Web adds meaning / semantics to data
  - Allow better data integration, automation, and discovery across sites
- **Key Technologies**
  - RDF (Resource Description Framework): base data model
  - SPARQL: query language for RDF data
  - OWL (Web Ontology Language): define rich ontologies
- **Current Status**
  - Some core ideas (e.g., structured data and ontologies) are widely adopted
  - Full vision remains only partially realized
- **Challenges**
  - Complexity of widespread adoption
  - Issues around privacy, data ownership, and feasibility
  - Need for standardization and tools
- **Criticism**
  - Skepticism about practicality and scalability
  - Concerns about centralization and censorship

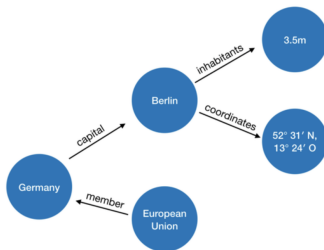
# WikiData

---

- **WikiData** is a free, open, collaborative knowledge base
  - Stores structured data for Wikipedia
  - Accessible via APIs using SPARQL queries
- **Graph-based data model**
  - *Item*: represents an entity or concept
    - Q42 → Douglas Adams
  - *Property*: describes a relationship or attribute
    - P31 (instance of), P27 (country of citizenship)
  - *Value*: specific data linked to an item via a property
    - Q42 (Douglas Adams) → P31 (instance of) → Q5 (human)
    - Q42 → P106 (occupation) → Q36180 (science fiction writer)
  - *Reference*: supports a claim by citing a source
    - Stating Douglas Adams's citizenship with a reference to a biography
  - *Qualifier*: adds context or additional information to a statement
    - Q90 (Paris) → P1082 (population) → "2,165,423"
    - With qualifier: P585 (point in time) → "2021"
    - Meaning: "The population of Paris was 2,165,423 in the year 2021"
- **Applications:**
  - Knowledge graph
  - Semantic search
  - AI reasoning
  - Data enrichment



- **DBpedia** extracts structured content from Wikipedia
  - Creates a large-scale, multilingual knowledge graph for querying
  - Data is extracted as RDF triples (*Subject, Predicate, Object*)
    - E.g., “Berlin” entity linked with properties like `dbo:country` Germany, `dbo:populationTotal` 3.5M
  - Enables semantic queries over Wikipedia data via SPARQL endpoints
- **Applications**
  - Semantic Web
  - Enhancing AI models with real-world knowledge




# Semantic Networks

---

- **Semantic Networks** represent knowledge as graphs of concepts and relations
  - Nodes represent concepts
  - Edges represent relations (e.g., “is-a”, “part-of”)
    - E.g., if a Dog is an Animal, it inherits Animal traits
  - Examples: WordNet, ConceptNet
- **Pros**
  - Easy to visualize and traverse
  - Support reasoning
  - Common in early AI systems and current KG applications

# WordNet


- **WordNet** is a large lexical database of English words
  - Models semantic relationships between words
  - Manually curated for high-quality semantic relations
  - Incomplete for domain-specific language



- **Graph** with synsets as nodes and relations as edges
  - *Synsets* are sets of synonyms expressing a distinct concepts
    - E.g., {car, automobile} share the same synset
  - *Relations* between synsets
    - Is-a relationships (e.g., Dog is a type of Animal)
    - Part-whole relationships (e.g., Wheel is a part of Car)
    - Opposite meanings
- **Applications**
  - Word sense disambiguation: choose correct word meaning in context
  - Semantic similarity measures: assess concept closeness
  - Information retrieval and question answering systems

# ConceptNet

- **ConceptNet** is a large knowledge graph
  - Connects words and phrases with labeled semantic relationships
  - Represents commonsense knowledge about the world
- **Structure**
  - Nodes: concepts (words or phrases)
  - Edges: semantic relationships between concepts, e.g.,
    - *IsA*: (dog, animal)
    - *PartOf*: (wheel, car)
    - *UsedFor*: (knife, cutting)
    - *CapableOf*: (bird, fly)
    - *Causes*: (fire, smoke)
  - E.g., (bicycle, UsedFor, transportation)
- **Applications**
  - Natural language understanding
  - Question answering and chatbots
  - Commonsense reasoning in AI
  - Semantic search and recommendation systems



# Frame-Based Representations

---

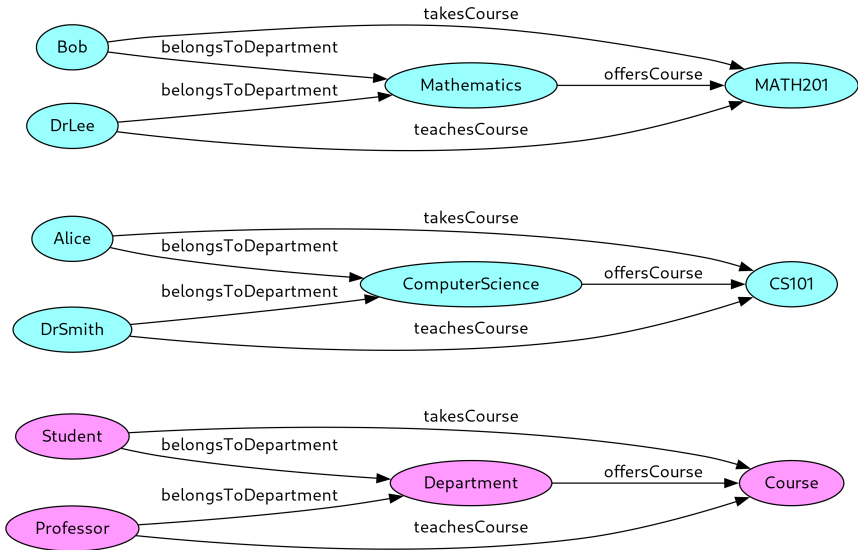
- Frame-based systems represent structured knowledge about objects, events, or situations
- Key Concepts:
  - Frame: A data structure for a concept or entity
    - E.g., a frame for Dog might include properties like `hasLegs`, `hasFur`, `barks`
  - Slots: attributes or relationships associated with the frame
    - E.g., slot `hasLegs` with value 4
  - Slot fillers: values or links to other frames that fill the slots
    - E.g., slot `eats` might link to another frame `Meat`
- Example:
  - Frame: Dog
  - Slots:
    - `isA: Animal`
    - `hasLegs: 4`
    - `sound: Bark`
    - `canDo: [Run, Fetch]`
- Features:
  - Inheritance: frames can inherit slots and slot values from more general frames, e.g.,

# Knowledge Graphs (KGs)

---

- KGs represent entities and their relationships as a graph structure
  - Nodes = entities
  - Edges = relations
  - E.g., “Paris  $\rightarrow$  isCapitalOf  $\rightarrow$  France”
- Query languages like SPARQL allow expressive information retrieval
- KGs support reasoning via path traversal and schema inference
- Applications:
  - Question answering
  - Recommendation
  - Semantic search
- Widely used by Google, Facebook, and academic search engines

# Knowledge Graph: University Example



# Technologies

---

- **TransE (Translation Embedding)**
  - Embedding model for knowledge graph completion
  - Represents relationships as translations in vector space:  $h + r \approx t$
  - Good for 1-to-1 relations, less effective with complex patterns
- **RotatE**
  - Embeds entities in complex space
  - Models relations as rotations:  $t = h \circ r$  where  $\circ$  is complex multiplication
  - Captures symmetry, antisymmetry, inversion, and composition
- **DeepProbLog**
  - Combines ProbLog (probabilistic logic) with deep learning
  - Supports neural predicates in logic programs
  - Learns probabilistic facts and neural components jointly
- **PyMLN**
  - Python-based Markov Logic Network (MLN) system
  - MLNs combine first-order logic with probabilistic graphical models
  - Allows reasoning with weighted logical rules
- **ProbLog**
  - Probabilistic logic programming language
  - Extends Prolog by attaching probabilities to facts
  - Computes success probabilities of queries
- **Tuffy**