



RBF: Learning Models (2/2)

- Use iterative approach (similar to EM algorithm):
 - Fix γ , solve for $\underline{\mathbf{w}}$ (one-step learning)
 - Fix $\underline{\mathbf{w}}$, solve for γ (gradient descent)
- **Step 1**
 - Assume γ is known and fixed
 - Learn $\underline{\mathbf{w}}$
- Impose perfect interpolation:

$$E_{in} = \frac{1}{n} \sum (h(\underline{\mathbf{x}}_i) - y_i)^2 = 0$$

- **Problem:**

$$h(\underline{\mathbf{x}}_j) = \sum_i w_i \exp(-\gamma \|\underline{\mathbf{x}}_i - \underline{\mathbf{x}}_j\|^2) = \sum_i w_i \phi_{i,j} = \underline{\phi}_j^T \underline{\mathbf{w}} = y_i$$

- N equations (one per point) and N unknowns $\underline{\mathbf{w}}$
- $\underline{\Phi}$ is known, function of data set and γ

Learning RBF Models

- The problem in matrix form is:

$$\underline{\underline{\Phi}} \cdot \underline{\underline{w}} = \underline{\underline{y}}$$

- If $\underline{\underline{\Phi}}$ is invertible, then $\underline{\underline{w}} = \underline{\underline{\Phi}}^{-1} \underline{\underline{y}}$
 - Desired values on training points
 - Exponential interpolates other points
- If $\underline{\underline{\Phi}}$ is not invertible, optimize in least square sense:

$$\operatorname{argmin}_{\underline{\underline{w}}} E_{in} = \sum_i (h(\underline{\underline{x}}_j) - y_i)^2$$

- Compute pseudo-inverse (assuming $\underline{\underline{\Phi}}^T \underline{\underline{\Phi}}$ is invertible)
- Assign weights:

$$\underline{\underline{w}} = (\underline{\underline{\Phi}}^T \underline{\underline{\Phi}})^{-1} \underline{\underline{\Phi}}^T \underline{\underline{y}}$$

- Step 2**

- Assume $\underline{\underline{w}}$ is known and fixed
- Learn γ

RBF Network vs Neural Networks

- The regression model for Neural Networks and RBF model is **similar**:

- **RBF**:

$$h(\underline{x}) = \sum_i w_i e^{-\gamma \|\underline{x} - \underline{x}_i\|^2} = \underline{w}^T \underline{\phi}$$

- **Neural networks**:

$$h(\underline{x}) = \Theta(\underline{w}^{(L)T} \underline{x}^{(L)}) = \Theta(\underline{w}^{(L)T} \underline{\Theta}(\underline{W}^{(L-1)} \dots))$$

- **Difference**:

- RBF has a single layer
- Neural networks have multiple layers

- **Similarities**:

- Combine features with weights using dot product
- Extract features from inputs
 - RBF features: $e^{-\gamma \|\underline{x} - \underline{x}_i\|^2}$, always > 0
 - NN hidden layers: features can be > 0 or < 0

RBF Network vs SVM

- The **model form** is the same:

- RBF:

$$h(\underline{x}) = \text{sign}\left(\sum_i w_i e^{-\gamma \|\underline{x} - \underline{x}_i\|^2}\right)$$

- SVM:

$$h(\underline{x}) = \text{sign}(\underline{w}^T \underline{x} + b)$$

- The interpretation is completely different (interpolation vs large margin)
 - RBF: all vectors (or centers of few clusters) contribute to the model
 - SVM: only support vectors contribute to the model


K-Nearest Neighbor (KNN) Model

- The **model form** is like:

$$h_{\underline{w}}(\underline{x}) = \frac{1}{n} \sum_{\underline{x}_i \text{ closest to } \underline{x}} w_i$$


- Idea:**

- Closeness implies a distance (e.g., euclidean metric) or similarity (e.g., a kernel)
- Consider the k closest points to the evaluation point \underline{x}
- Take an average of their response



KNN: Intuition of Number Degrees of Freedom

- Nearest neighbor model ($k = 1$)
 - Use response of closest point to \underline{x}
 - Similar to Voronoi tessellations: each point has a region where it is the closest and assigns its output to that region
- k is **the only parameter** for KNN
 - For $k = 1 \implies N$ neighborhoods, one around each training point
 - For $k = N \implies$ single neighborhood
 - Effective number of parameters: $\frac{N}{k}$, imagining N/k non-overlapping neighborhoods



KNN: Assumptions on the Data

- KNN makes **no assumption on data**
 - Opposite of linear model with strong data assumption
- KNN assumes **locality in parameter space**
 - Model is constant in example's neighborhood
 - E.g., $k = 1$: Voronoi tessellation (low-bias/high-variance)
 - E.g., $k = N$: Average value (high-bias/low-variance)

KNN: Training and Test Error

- For $k = 1$
 - No error on training set (low bias / high variance) assuming non-noisy target
 - E_{out} larger than E_{in}
- Increasing k
 - Training error E_{in} increases
 - Test error E_{out} decreases, then increases
 - Typical model complexity behavior in bias-variance diagrams

KNN vs RBF Models

- **Similarities**

- K-Nearest Neighbor is a *discrete* version of the RBF model

- **Differences:**

- Consider only the k *closest examples* to the point \underline{x} (not all examples in the training set)
- Use a *constant kernel* (responses are not weighted by distance)

- *Clustering*
- Anomaly Detection

K-Means Clustering: Problem Formulation

- N unlabeled points $\{\underline{x}_1, \underline{x}_2, \dots, \underline{x}_N\}$
- Partition points into K clusters S_1, \dots, S_K
 - Each cluster defined by center $\underline{\mu}_k$
 - Each point \underline{x}_i assigned to cluster $c(\underline{x}_i)$
 - Unknowns are $c(\underline{x}_1), \dots, c(\underline{x}_N), \underline{\mu}_1, \dots, \underline{\mu}_K$
- Minimize distance between each \underline{x}_i and assigned center $\underline{\mu}_k$ where $k = c(\underline{x}_i)$

$$\begin{aligned} J(c_1, \dots, c_N, \mu_1, \dots, \mu_K) &= \sum_{k=1}^K \sum_{\underline{x}_n \in S_k} \|\underline{x}_n - \underline{\mu}_k\|^2 \text{ (scanning the clusters)} \\ &= \sum_{i=1}^N \|\underline{x}_i - \underline{\mu}_{c(\underline{x}_i)}\|^2 \quad \text{(scanning the points)} \end{aligned}$$

- **K-means clustering is NP-hard** (combinatorial) and thus intractable
 - In fact there are K^N possible assignments

K-Means Clustering: Lloyd's Algorithm



- **Start with a random assignment** of N points to K clusters
 - Better than picking random centroids
- **Each iteration** has 2 steps
 - **Step 1:** Move centroid
 - Move each cluster's centroid to the mean point
 - Iterate over K clusters
 - $\underline{\mu}_k \leftarrow \frac{1}{|S_k|} \sum_{\underline{x}_n \in S_k} \underline{x}_n$
 - **Step 2:** Cluster assignment
 - Assign each \underline{x}_n to the closest cluster center
 - Iterate over N points
 - $S_k \leftarrow \{\underline{x}_n : \|\underline{x}_n - \underline{\mu}_k\| \leq \|\underline{x}_n - \underline{\mu}_l\| \forall l \neq k\}$

K-Means Clustering: Convergence

- **K-means algorithm converges** since:
 - Finite number of possible partitions (and values of objective functions)
 - Objective function $J(\cdot)$ always decreases
- **Objective function always decreases**
 - Cost function $J(\underline{\mu}_1, \dots, \underline{\mu}_K, c_1, \dots, c_N)$ depends on:
 - Centroids c_1, \dots, c_N
 - Point assignments $\underline{\mu}_1, \dots, \underline{\mu}_K$
 - K-means minimizes J by:
 - Adjusting centroids (fixed assignments)
 - Adjusting assignments (fixed centroids)
 - Similar to coordinate descent
- Generally converges to a **local minimum**
 - Run K-means multiple times with different random initializations
 - Choose best result


K-Means Clustering: Non-Separable Clusters

- For simplicity, you tend to imagine a clear separation between clusters
 - Clusters, especially in high dimensions, are not obviously separable
- Using K-means on data not obviously separated
 - E.g., market segmentation
 - E.g., t-shirt sizing
 - Collect height and width of customers
 - Run K-means
 - Find optimal way to split population into 3 sizes (S, M, L)



Choosing the Number of Clusters

- Often unclear how many clusters K exist
 - Visual analysis can be inconclusive in 2D or 3D
 - More difficult in high dimensions
- **Elbow Method**
 - Vary clusters K
 - Compute optimal cost function $J(\cdot)$
 - Choose K at “elbow” point if visible
 - Elbow absent if curve resembles hyperbole $\approx 1/K$
- **End-to-end approach**
 - Choose K to optimize later stages
 - E.g., more t-shirt sizes (more clusters) \implies
 - Satisfy customers
 - Complicates manufacturing
 - Increases inventory management



Interpretation of Clusters

- Cluster meaning **often interpreted manually** (difficult to automate)
 - Examine cluster centroids
 - Centroid values show “typical” point in each cluster
 - High, low, or zero feature values highlight key characteristics
 - Analyze distribution of features per cluster
 - Plot histograms or boxplots for each feature
 - Identify features that vary sharply across clusters
 - Visualize clusters in 2D or 3D
 - E.g., PCA, t-SNE, UMAP
 - Understand separation and internal structure
 - Compare clusters to external labels if available
 - See if clusters align with known real-world groups
 - Train a classifier like decision tree
 - Important features for predicting cluster reveal their meaning
- **Example: Customer Segmentation**

Cluster	Age	Annual Income	Spending Score	Label
Cluster 1	25 yrs	30K	90	Young Big Spenders
Cluster 2	50 yrs	80K	40	Comfortable Mid-Lifers
Cluster 3	35 yrs	120K	20	High Income, Low Spending Customers

- Clustering
- *Anomaly Detection*

Anomaly Detection: Problem Formulation

- **Problem:**

- $\{\underline{x}_1, \dots, \underline{x}_N\}$ examples with features $\underline{x} \in \mathbb{R}^P$ for good instances
- Detect bad/anomalous instances

- **Algorithm:**

- Unknown characteristics of “*bad instances*”
- Learn common traits of “*good instances*” using unsupervised learning
 - Find distribution for \underline{x}_i : $\Pr(\underline{x} \text{ is good})$
- Pick features
 - Find “sensitive” features, e.g., ratio between CPU load and network traffic
- Estimate distribution $\Pr(\underline{x} \text{ is good})$
- Choose threshold ε
- For new instance \underline{x}_{new} , if $\Pr(\underline{x}_{new} \text{ is good}) \leq \varepsilon$ flag as anomaly

Anomaly Detection: Example of Aircraft Engines

- **Problem**

- Test aircraft engines to identify anomalies in a new engine

- **Solution:**

- Features \underline{x}_i can be:
 - Heat generated
 - Vibration intensity
 - ...
- Collect data for all engines
- Model $\Pr(\underline{x} \text{ is good})$
- Decide if a new engine is acceptable $\Pr(\underline{x}_{new} \text{ is good}) \leq \varepsilon$ or needs more testing

Anomaly Detection: Example of Hacked Account

- **Problem**

- Find if an account for a given user i was hacked

- **Solution:**

- Model features that represent “user i activity”
- Features \underline{x}_i can be:
 - How many times s/he logs a day
 - How many times s/he fails to enter the password
 - How fast s/he types
 - How many pages s/he visits
 - How many times s/he posts comments
 - ...
- Model $\Pr(\underline{x} \text{ is good})$
- Identify unusual users by checking $\Pr(\underline{x}_{new} \text{ is good}) \leq \varepsilon$

Anomaly Detection: Example of Data Center

- **Problem**

- Monitor servers in a data center to find malfunctioning or hanged servers

- **Solution:**

- Features \underline{x}_i can be:
 - Memory in use
 - CPU load
 - Network traffic
 - Number of reads/writes per sec
 - CPU load / network activity
 - ...
- Model $\Pr(\underline{x} \text{ is good})$
- Identify unusual systems by checking $\Pr(\underline{x}_{new} \text{ is good}) \leq \varepsilon$

Using a Gaussian Model for Anomaly Detection

- Aka “density estimation”
- Given N examples $\underline{x}_1, \dots, \underline{x}_N \in \mathbb{R}^p$
- Ensure that the **features have a Gaussian distribution**
 - If not, apply some transformations, e.g., $\log(x_i + k)$
- **Estimate the parameters** of the Gaussian model $f_X(\underline{x})$
- Given a new example \underline{x}_{new} **compute** $\Pr(\underline{x}_{new} \text{ is good}) \leq \varepsilon$ to flag an anomaly

Estimate Univariate Gaussian Model

- You have N (scalar) examples $\underline{x}_1, \dots, \underline{x}_N \in \mathbb{R}$ representing “good instances”
- Assume the data is generated by a Gaussian distribution

$$X \sim \mathcal{N}(\mu, \sigma)$$


which has a PDF:

$$f_X(x; \mu, \sigma) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

- Estimate mean and sigma with maximum likelihood:

$$\mu = \frac{1}{N} \sum_i x_i$$

$$\sigma^2 = \frac{1}{N-1} \sum_i (x_i - \mu)^2$$



Estimate Multivariate Independent Gaussian Model


- You have N examples $\underline{x}_1, \dots, \underline{x}_N \in \mathbb{R}^p$ for “good instances”
- Assume independence of the features, the PDF of a multi-variate Gaussian X is:

$$f_X(\underline{x}; \underline{\mu}, \underline{\sigma}) = \prod_{i=1}^p f_{X_i}(x_i; \mu_i, \sigma_i)$$

- Infer the parameters μ_i and σ_i using discrete formulas to get the complete model

Estimate a Multi-Variate Gaussian Model

- **Problem:**
 - Often features vary together (e.g., network use and CPU load), causing mis-classifications with independent assumptions
 - Components of \underline{x}_{new} are within range but nonsensical together
 - E.g., low network use with high CPU load
- **Solution 1:**
 - Engineer features to highlight anomalies
 - Address variable correlation not modeled in independent Gaussian models
- **Solution 2:**
 - Estimate the entire multivariate model instead of assuming independence



Estimate a Multi-Variate Gaussian Model

- The PDF of a multi-variate Gaussian is:

$$f_{\mathbf{X}}(\mathbf{x}; \underline{\boldsymbol{\mu}}, \underline{\boldsymbol{\Sigma}}) = \frac{1}{(2\pi)^{\frac{n}{2}} |\underline{\boldsymbol{\Sigma}}|^{\frac{1}{2}}} \exp\left(-\frac{1}{2}(\mathbf{x} - \underline{\boldsymbol{\mu}})^T \underline{\boldsymbol{\Sigma}}^{-1}(\mathbf{x} - \underline{\boldsymbol{\mu}})\right)$$

- Estimate:

$$\underline{\boldsymbol{\mu}} \in \mathbb{R}^d = \frac{1}{N} \sum_{k=1}^N \mathbf{x}_k$$

$$\underline{\boldsymbol{\Sigma}} \in \mathbb{R}^{d \times d} = \{s_{ij}\} = \frac{1}{N-1} \sum_{k=1}^N (\mathbf{x}_k - \underline{\boldsymbol{\mu}})(\mathbf{x}_k - \underline{\boldsymbol{\mu}})^T$$

- Model requires more examples to train due to more parameters
- Independence between variables decomposes multivariate Gaussian into product of Gaussian distributions

Evaluate Anomaly Detection Systems

- To evaluate models one needs to:
 - Compare different models
 - Tune hyperparameters (e.g., ϵ) of models
 - Estimate out-of-sample error
- As always we should use a single real number for comparison
 - Use any classification metric, e.g.,
 - True/false positive/negative rate
 - Precision or recall
 - F-score
- *Labeled* data is still needed to rate models

$y = 0$ good

$y = 1$ anomalous

Evaluate Anomaly Detection Systems

- Often, anomalous examples $y = 1$ are much fewer than good examples $y = 0$
 - E.g., 10,000 good vs 20 bad examples
 - Address class imbalance for accurate model performance
- **Algorithm:**
 - Pick 60% of $y = 0$ data to train (only on good examples)
 - Split remaining $y = 0$ and $y = 1$ into validation and test sets
 - Ensure both sets represent the overall dataset
 - Train, validation, and test sets should not overlap but have the same characteristics
 - Helps evaluate model performance accurately
 - Use validation set to compare models, estimate hyperparameters
 - E.g., ϵ is the threshold for anomaly detection
 - Use test set to evaluate final model
 - Train on normal data, test on both normal and anomalous data
- **Aircraft engine example**

Dataset	$y = 0$ (Good Engines)	$y = 1$ (Bad Engines)
Total Example	10,000	20
Train Set	6,000	0
Validation Set	2,000	10
Test Set	2,000	10

Anomaly Detection vs Supervised Learning

- Even in unsupervised learning, you need labeled data for model evaluation
- Difference with supervised learning:
 - Anomaly detection/unsupervised learning: Train only on good examples
 - Supervised learning: Train on both good and bad examples
- Use:
 - Anomaly detection/unsupervised learning:
 - Learn from good examples due to few anomalous examples
 - Strong prior on the model
 - Future anomalous examples unknown (no prior)
 - Supervised learning: Less skewed classes in training set
 - Not a clear-cut decision