



Linear Regression: Find Optimal Model

- You want to minimize $E_{in}(\underline{\mathbf{w}}) = (\underline{\mathbf{X}}\underline{\mathbf{w}} - \underline{\mathbf{y}})^T(\underline{\mathbf{X}}\underline{\mathbf{w}} - \underline{\mathbf{y}})$ with respect to $\underline{\mathbf{w}}$

$$\nabla E_{in}(\underline{\mathbf{w}}^*) = \underline{\mathbf{0}}$$

$$\frac{2}{N} \underline{\mathbf{X}}^T (\underline{\mathbf{X}}\underline{\mathbf{w}}^* - \underline{\mathbf{y}}) = \underline{\mathbf{0}}$$

$$\underline{\mathbf{X}}^T \underline{\mathbf{X}}\underline{\mathbf{w}}^* = \underline{\mathbf{X}}^T \underline{\mathbf{y}}$$

- If the square matrix $\underline{\mathbf{X}}^T \underline{\mathbf{X}}$ is invertible:

$$\underline{\mathbf{w}}^* = (\underline{\mathbf{X}}^T \underline{\mathbf{X}})^{-1} \underline{\mathbf{X}}^T \underline{\mathbf{y}} = \underline{\mathbf{X}}^\dagger \underline{\mathbf{y}}$$

- The matrix $\underline{\mathbf{X}}^\dagger \triangleq (\underline{\mathbf{X}}^T \underline{\mathbf{X}})^{-1} \underline{\mathbf{X}}^T$ is called **pseudo-inverse**
 - It generalizes the inverse for non-square matrices, in fact:
 - $\underline{\mathbf{X}}^\dagger \underline{\mathbf{X}} = \underline{\mathbf{I}}$
 - If $\underline{\mathbf{X}}$ is square and invertible, then $\underline{\mathbf{X}}^\dagger = \underline{\mathbf{X}}^{-1}$

Complexity of One-Step Learning

- Learning with the pseudo-inverse is **one-step learning**
 - Contrasts with iterative methods, e.g., gradient descent
- Inverting a square matrix of size P is related to the number of parameters, not examples N
 - Complexity of one-step learning is $O(P^3)$

Linear Models Are Linear in What?

- A **model is linear** when the signal $s = \sum_{i=0}^P w_i x_i = \underline{\mathbf{w}}^T \underline{\mathbf{x}}$ is linear **with variables**
 - Unknown variables: weights w_i
 - Inputs x_i are fixed
- Applying a **non-linear transform to inputs** $z_i = \Phi(x_i)$ keeps the model linear, e.g.,
 - Positive/negative part (e.g., $z_i = \text{RELU}(x_i), \text{RELU}(-x_i)$)
 - Waterfall (conditioning model to different feature ranges)
 - Thresholding (e.g., $z_i = \min(x_i, T)$)
 - Indicator variables (e.g., $z_i = I(x_i > 0)$)
 - Winsorizing (replace outliers with a large constant value)
- Applying a **non-linear transform to weights** $z_i = \Phi(w_i)$ makes the model non-linear

Non-Linear Transformations in Linear Models

- **Transform variables**

- Use $\Phi : \mathcal{X} \rightarrow \mathcal{Z}$
- Transform each point $\underline{x}_n \in \mathcal{X} = \mathbb{R}^d$ into a point in feature space $\underline{z}_n = \Phi(\underline{x}_n) \in \mathcal{Z} = \mathbb{R}^{\tilde{d}}$ with $d \neq \tilde{d}$

- **Learn**

- Learn linear model in \mathcal{Z} , obtaining $\underline{\tilde{w}}$ for separating hyperplane

- **Predict**

- Evaluate model on new point in \mathcal{Z} :

$$y = \text{sign}(\underline{\tilde{w}}^T \Phi(\underline{x})) \text{ or } y = \underline{\tilde{w}}^T \Phi(\underline{x})$$

- Compute **decision boundary**

- In \mathcal{X} if Φ is invertible; or
- By classifying any $\underline{x} \in \mathcal{X}$ in \mathcal{Z}

- *Perceptron*
- Logistic Regression
- LDA, QDA
- Kernel Methods

Example of Classification Problems

- **Binary classification problem**

- $y \in \{0, 1\}$
 - Typically assign 1 to what you want to detect
- Email: spam, not_spam
- Online transaction: fraudulent, valid
- Tumor: malignant, benign

- **Multi-class classification problem**

- $y \in \{0, 1, 2, \dots, K\}$
- Email tagging: work, family, friends
- Medical diagnosis: not_ill, cold, flu
- Weather: sunny, rainy, cloudy, snow

Linear Regression for Classification

- Use **linear regression for classification**
 - Transform outputs into $\{+1, -1\} \in \mathbb{R}$
 - Learn $\underline{\mathbf{w}}^T \underline{\mathbf{x}}_n \approx y_n = \pm 1$
 - Use $\text{sign}(\underline{\mathbf{w}}^T \underline{\mathbf{x}}_n)$ as model (perceptron)
- **Not optimal**: outliers influence fit due to square distance metric
 - Use weights from linear regression to initialize a learning algorithm for classification (e.g., PLA)

Perceptron Learning Algorithm (PLA)

- First machine learning algorithm discovered
- **Algorithm**
 - Training set $\mathcal{D} = \{(\underline{x}_1, y_1), \dots, (\underline{x}_n, y_n)\}$
 - Initialize weights \underline{w}
 - Random values
 - Use linear regression for classification as seed
 - Pick a misclassified point $\text{sign}(\underline{w}^T \underline{x}_i) \neq y_i$ from training set \mathcal{D}
 - Update weights: $\underline{w}(t+1) = \underline{w}(t) + y_i \underline{x}_i$
 - Like stochastic gradient descent
 - Iterate until no misclassified points
- Algorithm **converges** (slowly) for linearly separable data
- **Pocket version of PLA**
 - Idea
 - Continuously update solution
 - Keep best solution “in the back pocket”
 - Have a solution if stopping after max iterations

Non-Linear Transformations for Classifications

- Classification problems have **varying degrees of non-linear boundaries**:
 1. Non-linearly separable data
 - E.g., + in center, - in corners in a 2-feature scatter plot
 2. Mostly linear classes with few outliers
 3. Higher-order decision boundary
 - E.g., quadratic
 4. Non-linear data-feature relationship
 - E.g., variable threshold

- Perceptron
- *Logistic Regression*
- LDA, QDA
- Kernel Methods

Logistic Regression Is a Probabilistic Classifier

- **Logistic regression** learns:
 - The probability of each class $\Pr(y|\underline{x})$ given input \underline{x}
 - Instead of predicting class y directly
- **Parametric approach**: assume $\Pr(y|\underline{x}; \underline{w})$ has a known functional form

$$\Pr(y = 1|\underline{x}; \underline{w}) = \text{logit}(\underline{w}^T \underline{x})$$

- **Optimize parameters** \underline{w} using maximum likelihood

$$\underline{w}^* = \text{argmax}_{\underline{w}} \Pr(\mathcal{D}; \underline{w})$$

- **Predict** by outputting class with highest probability

$$h_{\underline{w}}(\underline{x}) = \begin{cases} +1 & \Pr(y = 1|\underline{x}; \underline{w}) \geq 0.5 \\ -1 & \Pr(y = 1|\underline{x}; \underline{w}) < 0.5 \end{cases}$$

Logistic Regression: Example

- Assume y is:
 - The event y_i *"patient with characteristics \underline{x} had heart attack"*
 - Function of parameters \underline{x} (e.g., age, gender, diet)
- In data set \mathcal{D} :
 - No samples of $\Pr(y|\underline{x})$
 - Have realizations:
 - *"Patient with \underline{x}_1 had a heart attack"*
 - *"Patient with \underline{x}_2 didn't"*
 - ...
- Learn $\Pr(y|\underline{x})$
 - Find best parameters \underline{w} for logistic regression model to explain data \mathcal{D}

Logistic Function

- Aka “sigmoid”
- **Logistic function** $\text{logit}(s)$ is defined as

$$\text{logit}(s) \triangleq \frac{e^s}{1 + e^s} = \frac{1}{1 + e^{-s}}$$

- Varies in $[0, 1]$
- Crosses the origin at 0.5
- Asymptotes at 0 and 1
- It is a soft version of $\text{sign}()$

Logistic Regression vs Linear Classifier

- **Functional form is similar**
 - **Logistic regression**: $h(\underline{\mathbf{w}}) = \text{logit}(\underline{\mathbf{w}}^T \underline{\mathbf{x}})$
 - **Linear classifier** (perceptron): $h(\underline{\mathbf{w}}) = \text{sign}(\underline{\mathbf{w}}^T \underline{\mathbf{x}})$
- **Difference** in probabilistic interpretation and fitting method
 - Logistic regression** lacks samples of probability function to interpolate
 - Uses realizations of random variable
 - Seeks model parameters maximizing data likelihood
 - **Linear classification** assumes class value is linear function of inputs

Error for Probabilistic Binary Classifiers

- For probabilistic binary classification, use **log-probability error** as point-wise error

$$e(h(\underline{\mathbf{x}}), y) \triangleq -\log(\Pr(y = h(\underline{\mathbf{x}})|\underline{\mathbf{x}}; \underline{\mathbf{w}}))$$

- Negate for positive errors: $\log([0, 1]) \in [-\infty, 0)$
- Log probability generalizes 0-1 error
 - Case $y = 1$
 - Output $h(\underline{\mathbf{x}})$ close to 1 \implies probability 1 $\implies \log(1) = 0 \implies e(\cdot) = 0$
 - Output close to 0 $\implies e() = -\log(0) \rightarrow +\infty$
 - Similar behavior for $y = 0$

One-Liner Error for Probabilistic Binary Classifiers

- **Point-wise error** for example (\underline{x}, y) for probabilistic binary classifiers is defined as:

$$\begin{aligned} e(h(\underline{x}), y) &\triangleq -\log(\Pr(h(\underline{x}) = y|\underline{x})) \\ &= \begin{cases} -\log(\Pr(y = 1|\underline{x})) & y = 1 \\ -\log(\Pr(y = 0|\underline{x})) & y = 0 \end{cases} \\ &= \begin{cases} -\log(\Pr(y = 1|\underline{x})) & y = 1 \\ -\log(1 - \Pr(y = 1|\underline{x})) & y = 0 \end{cases} \end{aligned}$$

- Any function of a binary variable:

$$y = \begin{cases} a & x = 1 \\ b & x = 0 \end{cases}$$

can be written as one-liner: $y = x \cdot a + (1 - x) \cdot b$

- Point-wise error can be written independently of $\Pr(y = 1|\underline{x})$:

$$e(h(\underline{x}), y) = -y \log(\Pr(y = 1|\underline{x})) - (1 - y) \log(1 - \Pr(y = 1|\underline{x}))$$

One-Liner Error for Logistic Regression

- The point-wise error for a binary classifier is:

$$e(h(\underline{\mathbf{x}}), y) = -y \log(\Pr(y = 1|\underline{\mathbf{x}})) - (1 - y) \log(1 - \Pr(y = 1|\underline{\mathbf{x}}))$$

- Simplify further **with logit function**:

$$e(h(\underline{\mathbf{x}}), y) \triangleq -\log \Pr(h(\underline{\mathbf{x}}) = y)$$

... a bit of math manipulation ...

$$= -\log \text{logit}(y \underline{\mathbf{w}}^T \underline{\mathbf{x}})$$

$$\text{since } \text{logit}(s) = \frac{1}{1 + e^{-s}}$$

$$= \log(1 + \exp(-y \underline{\mathbf{w}}^T \underline{\mathbf{x}}))$$

- Point-wise error for logistic regression equals **cross-entropy error**

Cross-Entropy Error

- **Point-wise error for logistic regression** has expression:

$$e(h(\underline{\mathbf{x}}), y) = \log(1 + \exp(-y \cdot \underline{\mathbf{w}}^T \underline{\mathbf{x}}))$$

- It is called **cross-entropy error**
- Note: no $-$ before $\log(\cdot)$ but before $y \cdot \underline{\mathbf{w}}^T \underline{\mathbf{x}}$
- Cross-entropy error **generalizes 0-1 error**
 - If $\underline{\mathbf{w}}^T \underline{\mathbf{x}}$ agrees with y in sign and $|\underline{\mathbf{w}}^T \underline{\mathbf{x}}|$ is large \implies error goes to 0
 - If they disagree in sign \implies error goes towards ∞
- Define **in-sample error on training set** as average of point-wise errors:

$$E_{in} \triangleq \frac{1}{N} \sum_n e(h(\underline{\mathbf{x}}_n), y_n)$$

Fitting Logistic Regression

- A plausible error measure of a hypothesis is based on **likelihood of data**

$$\Pr(\mathcal{D} | h = f)$$

- “How likely is the data \mathcal{D} under the assumption that $h = f$?”
- “How likely is that the data \mathcal{D} was generated by h ?”
- **Maximize likelihood** \mathcal{D} generated from logistic regression

$$\Pr(y = 1 | \underline{\mathbf{x}}; \underline{\mathbf{w}})$$

- It can be proved that this is equivalent to **minimizing in-sample error** on training set **using cross-entropy error**

Fitting Logistic Regression (Optional)

- Find $\underline{\mathbf{w}}$ that maximizes likelihood for data set $\mathcal{D} = \{(\underline{\mathbf{x}}_1, y_1), \dots, (\underline{\mathbf{x}}_N, y_N)\}$ generated by model $h(\underline{\mathbf{x}})$:

$$\Pr(D|\underline{\mathbf{w}}) = \Pr(y_1 = h(\underline{\mathbf{x}}_1) \wedge \dots \wedge y_N = h(\underline{\mathbf{x}}_N)) = \Pr(y_1 = y'_1 \wedge \dots \wedge y_N = y'_N)$$

- Model form:

$$y' = h(\underline{\mathbf{x}}) = \begin{cases} +1 & \text{if } \text{logit}(\underline{\mathbf{w}}^T \underline{\mathbf{x}}) > 0.5 \\ -1 & \text{otherwise} \end{cases}$$

- Assuming independence among training examples:

$$\Pr(D|\underline{\mathbf{w}}) = \prod_{i=1}^N \Pr(y_i = y'_i | \underline{\mathbf{x}}_i)$$

- Fold y_n in expression:

- When $y_n = 1$, $\Pr(y_n = y'_n) = \text{logit}(\underline{\mathbf{w}}^T \underline{\mathbf{x}}_n)$
- When $y_n = -1$, $\Pr(y_n = y'_n) = \text{logit}(-\underline{\mathbf{w}}^T \underline{\mathbf{x}}_n)$

$$\Pr(y_n = y'_n) = \text{logit}(y_n \underline{\mathbf{w}}^T \underline{\mathbf{x}}_n)$$

Fitting Logistic Regression (Optional)

- Given:

$$\Pr(D|\underline{\mathbf{w}}) = \prod_{i=1}^N \text{logit}(y_n \underline{\mathbf{w}}^T \underline{\mathbf{x}}_n)$$

- Re-write optimization as minimizing sum of point-wise errors

$$E_{in} = \sum e(h(\underline{\mathbf{x}}_n), y_n)$$

- Maximize $\log(\dots)$ with respect to $\underline{\mathbf{w}}$ since \log argument > 0 and $\log()$ is monotone
- Equivalently, minimize:

$$\begin{aligned} -\frac{1}{N} \log(\dots) &= -\frac{1}{N} \log(\prod(\dots)) = -\frac{1}{N} \sum (\log(\dots)) \\ &= \frac{1}{N} \sum \log\left(\frac{1}{\text{logit}(y_n \underline{\mathbf{w}}^T \underline{\mathbf{x}}_n)}\right) \\ &= \frac{1}{N} \sum \log(1 + \exp(-y_n \underline{\mathbf{w}}^T \underline{\mathbf{x}}_n)) = \frac{1}{N} \sum e(h(\underline{\mathbf{x}}_n), y_n) = E_{in}(\underline{\mathbf{w}}) \end{aligned}$$

Gradient Descent for Logistic Regression

- Gradient descent requires two inputs:
 - Gradient of the cost function $\frac{\partial E}{\partial w_j}$ for all j
 - Cost function $E_{in}(\underline{\mathbf{w}})$

- The cost function is:

$$E_{in}(\underline{\mathbf{w}}) = \frac{1}{N} \sum_i e(h(\underline{\mathbf{x}}_i; \underline{\mathbf{w}}), y_i)$$

- The cost function for logistic regression:

$$E_{in}(\underline{\mathbf{w}}) = \frac{1}{N} \sum_i \log(1 + \exp(-y_i \underline{\mathbf{w}}^T \underline{\mathbf{x}}_i))$$

- Thus gradient descent converges to global minimum
 - It can be shown that $E_{in}(\underline{\mathbf{w}})$ is convex in $\underline{\mathbf{w}}$
 - In fact sum of exponentials and flipped exponentials is convex and log is

One-Vs-All Multi-Class Classification

- Aka “one-vs-rest” classifier
- Assume we have n classes c_1, \dots, c_n to distinguish given \underline{x}
- Learn
 - Create n binary classification problems where we classify c_i vs c_{-i} (everything but i)
 - Learn n classifiers with optimal \underline{w}_i , each estimating $\Pr(y = i | \underline{x}; \underline{w}_i)$
- Predict
 - Evaluate the n classifiers
 - Pick the class $y = i$ with the maximum $\Pr(y = i | \underline{x})$

Cost Function for Multi-Class Probabilistic Classification

- The cost function for logistic regression is:

$$E_{in}(\underline{\mathbf{w}}) = -\frac{1}{N} \sum_{i=1}^N (y_i \log \Pr(y = 1|\underline{\mathbf{x}}_i) + (1 - y_i) \log(1 - \Pr(y = 1|\underline{\mathbf{x}}_i)))$$

- Encode the expected outputs $\underline{\mathbf{y}}_i$ one-hot
 - I.e., the j -th element $\underline{\mathbf{y}}_i|_j$ is 1 iff the correct class is the j -th
 - E.g., for $k = 4$ 1000
- Using $\underline{\mathbf{h}}(\underline{\mathbf{x}})$ as the outputs from the model and $\Pr(y = 1|\underline{\mathbf{x}}) = p(\underline{\mathbf{x}})$:

$$E_{in}(\underline{\mathbf{w}}) = -\frac{1}{N} \sum_{i=1}^N \sum_{k=1}^K \left(\underline{\mathbf{y}}_i \log(\underline{\mathbf{p}}(\underline{\mathbf{x}}_i)) + (1 - \underline{\mathbf{y}}_i) \log(1 - \underline{\mathbf{p}}(\underline{\mathbf{x}}_i)) \right) |_k$$

- In the innermost summation we consider the error on each class / digit
 - E.g., for $k = 4$ 1000 vs 0100
 - The digits that are equal don't contribute to the error
 - The different digits give a positive contribution

- Perceptron
- Logistic Regression
- **LDA, QDA**
- Kernel Methods

Basic Idea of Parametric Models

- We assume a model is generating the data
 - The functional form of the model is known
 - The model is parametrized with unknown parameters to estimate
- **Pros**
 - Utilize structure in the data
 - Models are easy to fit: few parameters to estimate
 - Accurate predictions if model assumptions are correct
- **Cons**
 - Strong assumptions about the data
 - Low accuracy if model assumptions are incorrect

Linear and Quadratic Discriminant Analysis

- Aka LDA and QDA
- Parametric models
 - Assume each class generating process is multivariate Gaussian
 - Classifiers with linear and quadratic decision surface
- **Pros**
 - Closed-form solutions easy to compute (sample mean and covariance)
 - Inherently multiclass
 - No hyperparams to tune
- **Cons**
 - Strong assumptions about the data

LDA / QDA: Model Form

- Both LDA and QDA assume that the class generating process $f_k(\underline{\mathbf{x}}; \underline{\boldsymbol{\mu}}_k, \underline{\boldsymbol{\Sigma}}_k)$ is from a multivariate Gaussian
- Linear discriminant analysis has a model:

$$f_k(\underline{\mathbf{x}}; \underline{\boldsymbol{\mu}}_k, \underline{\boldsymbol{\Sigma}}_k) \sim \mathcal{N}(\underline{\boldsymbol{\mu}}_k, \underline{\boldsymbol{\Sigma}})$$

where:

- Means $\underline{\boldsymbol{\mu}}_k$ are different for all k classes
 - The covariance matrix $\underline{\boldsymbol{\Sigma}}$ is the same for all k classes
 - It can be proven that the classes are separated by linear decision boundaries
- Quadratic discriminant analysis has a model:
 - Classes k have different covariance matrix $\underline{\boldsymbol{\Sigma}}_k$
 - It can be proved that the classes are separated by quadratic boundaries

Bayes Theorem for LDA / QDA

- Consider a classification setup with multi-class output $Y \in \{1, \dots, K\}$
- We need to build a parametric model for the conditional distribution:

$$\Pr(Y = k | X = \underline{x})$$

- Use Bayes theorem:

$$\Pr(Y = k | X = \underline{x}) = \frac{\Pr(X = \underline{x} | Y = k) \cdot \Pr(Y = k)}{\Pr(X = \underline{x})}$$

where:

- $\Pr(X = \underline{x} | Y = k)$: given a class, estimate the probability of \underline{x}
 - $\Pr(Y = k) = \pi_k$: estimate the probability of each class (prior)
 - $\Pr(X = \underline{x})$: estimate the probability of each input
- Estimate probabilities from data

LDA / QDA: Boundary Decision (*)

- Consider the ratio between the probabilities of $Y = k$ vs $Y = j$:

$$r = \frac{\Pr(Y = k|X = \underline{\mathbf{x}})}{\Pr(Y = j|X = \underline{\mathbf{x}})}$$

- Using the model assumption and Bayes theorem:

$$\begin{aligned}\Pr(Y = k|X = \underline{\mathbf{x}}) &\propto \Pr(X = \underline{\mathbf{x}}|Y = k) \cdot \Pr(Y = k) \\ &= f_k(\underline{\mathbf{x}}; \underline{\boldsymbol{\mu}}_k) \cdot \pi_k\end{aligned}$$

where:

$$f_k(\underline{\mathbf{x}}) = c \cdot \exp\left(-\frac{1}{2}(\underline{\mathbf{x}} - \underline{\boldsymbol{\mu}}_k)^T \underline{\boldsymbol{\Sigma}}^{-1}(\underline{\mathbf{x}} - \underline{\boldsymbol{\mu}}_k)\right)$$

- We can apply a $\log(\cdot)$ since it is a monotone transformation

$$r = \log \frac{f_k(\underline{\mathbf{x}})}{f_j(\underline{\mathbf{x}})} + \log \frac{\pi_k}{\pi_j}$$

LDA / QDA: Boundary Decision (*)

- We can apply a $\log(\cdot)$ since it is a monotone transformation

$$r = \log \frac{f_k(\underline{\mathbf{x}})}{f_j(\underline{\mathbf{x}})} + \log \frac{\pi_k}{\pi_j}$$

- The second term does not depend on $\underline{\mathbf{x}}$
- The first term is proportional to:

$$(\underline{\mathbf{x}} - \underline{\boldsymbol{\mu}}_k)^T \underline{\boldsymbol{\Sigma}}^{-1} (\underline{\mathbf{x}} - \underline{\boldsymbol{\mu}}_k) - (\underline{\mathbf{x}} - \underline{\boldsymbol{\mu}}_j)^T \underline{\boldsymbol{\Sigma}}^{-1} (\underline{\mathbf{x}} - \underline{\boldsymbol{\mu}}_j)$$

- By expanding the expression, simplifying $\underline{\mathbf{x}}^T \underline{\boldsymbol{\Sigma}}^{-1} \underline{\mathbf{x}}$ noticing that $\underline{\mathbf{x}}^T \underline{\boldsymbol{\Sigma}}^{-1} (\underline{\boldsymbol{\mu}}_k - \underline{\boldsymbol{\mu}}_j)$ plus its transposed value (because $\underline{\boldsymbol{\Sigma}}$ is symmetrical) is equal to $2\underline{\mathbf{x}}^T \underline{\boldsymbol{\Sigma}}^{-1} (\underline{\boldsymbol{\mu}}_k - \underline{\boldsymbol{\mu}}_j)$ we get:

$$-\frac{1}{2}(\underline{\boldsymbol{\mu}}_k + \underline{\boldsymbol{\mu}}_j)^T \underline{\boldsymbol{\Sigma}}^{-1} (\underline{\boldsymbol{\mu}}_k - \underline{\boldsymbol{\mu}}_j) + \underline{\mathbf{x}}^T \underline{\boldsymbol{\Sigma}}^{-1} (\underline{\boldsymbol{\mu}}_k - \underline{\boldsymbol{\mu}}_j)$$

LDA / QDA: Learn Model

- In practice:
 - We don't care about $\Pr(X = \underline{x})$ since this is common for all classes
 - We assume to know the prior $\Pr(Y = k) = \pi_k$ or we estimate it from the data
 - We need to estimate the conditional probability $\Pr(X = \underline{x} | Y = k)$
- The model assumes that the conditional probability has a Gaussian distribution:

$$\Pr(X = \underline{x} | Y = k) = f_k(\underline{x}; \underline{\mu}_k, \underline{\Sigma}_k)$$

where:

$$f_k(\underline{x}) = \frac{1}{(2\pi)^n |\underline{\Sigma}_k|^{1/2}} \exp \left(-\frac{1}{2} (\underline{x} - \underline{\mu}_k)^T \underline{\Sigma}_k^{-1} (\underline{x} - \underline{\mu}_k) \right)$$

- We can estimate the parameters $\underline{\mu}_k$, $\underline{\Sigma}_k$ from the data using sample mean and covariance

Evaluating LDA / QDA

- When we get a new $\underline{x} = \underline{x}'$ we compute for each class $Y = k$

$$\Pr(Y = k | X = \underline{x}) \propto f_k(\underline{x}; \underline{\mu}_k, \underline{\underline{\Sigma}}_k) \cdot \pi_k$$

and choose the k that maximizes the posterior probability

- Sometimes $f_i(\underline{x})$ is from a multivariate Gaussian distribution where $\underline{\underline{\Sigma}}$ is diagonal because of feature independence
 - Then the expressions can be simplified even more

- Perceptron
- Logistic Regression
- LDA, QDA
- ***Kernel Methods***