




# Bayesian Network: Car Insurance Company (1/2)

- A **car insurance company**:
  - Receive an application from an individual to insure a specific vehicle
  - Analyze information about the individual and its car
  - Decide on appropriate annual premium to charge
  - Pay out a claim, based on the type of claim
- Build a Bayesian network that captures the causal structure of the domain
  - **Input information**:
    - About the applicant: *Age, YearsWithLicense, DrivingRecord, GoodStudent*
    - About the vehicle: *MakeModel, VehicleYear, Airbag, SafetyFeatures*
    - About the driving situation: *Mileage, HasGarage*
  - Some input informations are **important but not available**:
    - *RiskAversion*
    - *DrivingBehavior*
  - **Type of claims**:
    - *MedicalCost*: injuries sustained by the applicant
    - *LiabilityCost*: lawsuits filed by other parties against applicant
    - *PropertyCost*: vehicle damage to either party and theft of the vehicle

## Bayesian Network: Car Insurance Company (2/2)

- **Blue nodes:** information provided by the applicants
- **Brown nodes:** hidden variables (not observable)
- **Violet nodes:** target variables



- *Exact Inference in Bayesian Networks*
- Approximate Inference in Bayesian Networks

# Exact Inference in Bayesian Networks

- **Goal of exact inference**

- Compute the posterior  $P(X|\underline{E} = \underline{e})$  for query variable  $X$  given evidence  $\underline{e}$

- **Variables involved**

- Query variable  $X$
- Evidence variables  $\underline{E} = \{E_1, \dots, E_m\}$
- Hidden variables  $\underline{Y} = \{Y_1, \dots, Y_\ell\}$

- **Inference by Enumeration**

- Use full joint distribution and sum over all hidden variables:

$$P(X|e) = \alpha \sum_Y P(X, e, Y)$$

- **Variable Elimination**

- Improves efficiency by caching intermediate results
- Eliminates variables systematically to avoid redundant sums
- Removing irrelevant variables
  - Variables not ancestors of query or evidence can be ignored

- **Problems**

- Exact inference is efficient  $O(n)$  for trees, but intractable  $O(2^n)$  in general
- It doesn't work for continuous variables

• Basis for approximate methods when exact is impractical

# Exact Inference in Bayesian Networks: Example

- You get a call from both John and Mary, what is the probability of the burglary?

$$P(\text{Burglary} | \text{JohnCalls} = \text{True}, \text{MaryCalls} = \text{True})$$

- A conditional probability can be computed summing terms from the full joint distribution

$$\Pr(X|\underline{e}) = \alpha \Pr(X, \underline{e}) = \alpha \sum_y \Pr(X, \underline{e}, \underline{y})$$

- Terms of the joint distribution can be written as products of conditional probabilities from the Bayesian network

$$\Pr(b|j, m) = \alpha \Pr(B, j, m) = \alpha \sum_e \sum_a \Pr(B, j, m, e, a)$$

- Then the joint probability is written in terms of CPTs of the Bayesian network

$$\Pr(b|j, m) = \alpha \sum_e \sum_a \Pr(b) \Pr(e) \Pr(a|b, e) \Pr(j|a) \Pr(m|a)$$

- 
- Exact Inference in Bayesian Networks
  - ***Approximate Inference in Bayesian Networks***

# Monte Carlo Algorithms

---

- **Monte Carlo algorithms** are randomized sampling algorithms used to estimate quantities that are difficult to calculate exactly
  - E.g., samples from the posterior probability of a Bayes network
- **Pros**
  - The accuracy of the approximation depends on the number of samples generated
  - You can get arbitrarily close to the true probability distribution with enough samples
  - Is used in many branches of science
- **Cons**
  - Difficult to understand how the variables interact
  - Computationally intensive



# Sampling from Arbitrary Distributions


- **Goal:** Sample from a discrete or continuous probability distribution
- **Solution**
  - Start with uniform random numbers  $r \in [0, 1]$
  - Construct CDF (cumulative distribution function)  $F(x)$ 
    - $F(x) = \Pr(X \leq x)$
  - For discrete distributions:
    - Create table of outcomes and cumulative probabilities
    - Find smallest outcome where  $F(x) > r$
  - For continuous distributions:
    - Use inverse transform:  $x = F^{-1}(r)$ , e.g.,

$$F(x) = 1 - e^{-\lambda x} \rightarrow x = F^{-1}(r) = -\frac{1}{\lambda} \ln(1 - r)$$

- If  $F^{-1}$  has no closed form, use numerical methods

# Sampling Bayesian Network Without Evidence

- **Goal:** Generate events from a Bayesian network without evidence (prior sampling)
- **Solution**
  - Sample variables in topological order (to ensure parents have values)
  - Source nodes have known unconditional probability distribution
    - E.g.,  $\Pr(Rain) = 0.5$
  - Conditional variable's probability distribution depends on parent's values
    - E.g.,  $\Pr(WetGrass|Rain = T) = 0.1$
  - Implement Bayesian network semantics, representing joint probability:



$$f_{PS}(x_1, \dots, x_n) = \prod_{i=1}^n \Pr(x_i | \text{parents}(X_i))$$

where *PS* means “Prior Sampling”

# Consistency of Sampling

- **Consistency of estimation:** distribution from prior sampling converges to true probability as  $N \rightarrow \infty$
- If  $N_{PS}$  is the number of times event  $x_1, \dots, x_n$  occurs:

$$\lim_{N \rightarrow \infty} \frac{N_{PS}(x_1, \dots, x_n)}{N} = \Pr(x_1, \dots, x_n)$$


- Estimate probability using:

$$\Pr(x_1, \dots, x_n) \approx \frac{N_{PS}(x_1, \dots, x_n)}{N}$$

- Converges with rate  $\frac{1}{\sqrt{N}}$

# Rejection Sampling

- **Rejection sampling** is a method for sampling from a hard-to-sample distribution
- **Goal:** Compute  $\Pr(X = x|E = e)$  when evidence  $e$  is rare
  1. Generate samples from the prior distribution
    - Estimate  $\Pr(x, e)$
  2. Reject samples not matching evidence, i.e.,  $X \wedge E \neq e$ 
    - Remaining samples  
 $X \wedge E = e$  estimate  
 $\Pr(X, E = e)$
  3. Count occurrences of  $X = x$  in remaining samples  $X \wedge E = e$ 
    - Estimate  $\Pr(X = x|E = e)$
- **Example:**
  - You want to estimate  $\Pr(\text{Rain}|\text{Sprinkler} = T)$
  - Sample 100 times
    - You get 73 samples with  $\neg\text{Sprinkler}$  and they are rejected
    - You are left with 27 samples with *Sprinkler*
    - Out of them only 8 have *Rain* and 19 have  $\neg\text{Rain}$



# Rejection Sampling: Pros and Cons

---

- **Pros**

- Consistent estimate
  - Converges to true value as number of samples increases

- **Cons**

- Many samples are rejected, depending on rarity of  $\Pr(E = e)$
- Fraction of samples matching evidence  $e$  decreases exponentially with more evidence variables
  - Curse of dimensionality
  - Not suitable for complex systems
- Difficult with continuous variables
  - E.g.,  $\Pr(E = e)$  is theoretically 0 due to limited floating-point precision

# Importance Sampling

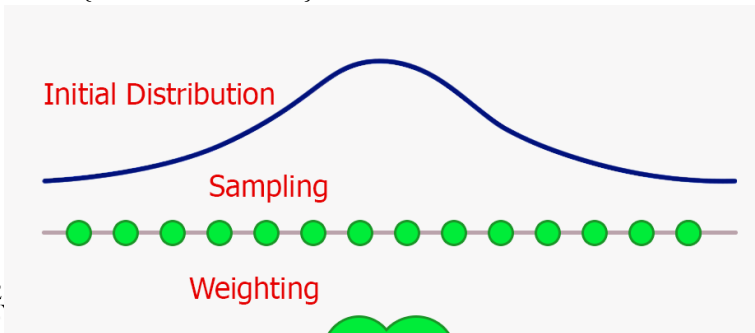
::: columns ::: {.column width=60%}

- **Importance sampling**

- Draw samples from “easier” distribution  $Q(X)$
- Weight each sample by importance weight  $w = \frac{\Pr(X)}{Q(X)}$
- Estimate probability by averaging weighted samples:

$$E[f(X)] \approx \frac{1}{N} \sum_{i=1}^N w_i f(X_i)$$

::: ::: {.column width=35%}



# Markov Chain Monte Carlo

---

- This is one of the top 10 most mind blowing algorithms in history
  - Euclide's GCD
  - Fundamental theorem of calculus
  - Quicksort
  - Fast Fourier Transform
  - Viterbi algorithm
  - MCMC sampling
  - Kalman filter
  - RSA Algorithm
  - ...
- Invented by Ulam, Von Neumann, Metropolis and others during the Manhattan Project (1940)
  - Used to solve high-dimensional integrals, Bayesian inference, ...
- **Purpose:** Approximate inference for Bayesian networks when exact inference is hard
  - MCMC differs from rejection and importance sampling
  - Make random changes to preceding sample instead of generating each sample independently
  - **Magic:** two very different objects (Markov Chains) and Bayesian Networks

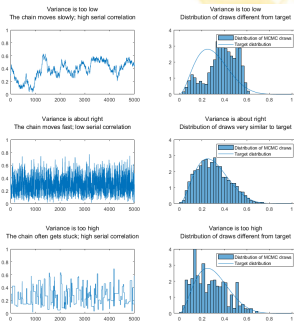
# Markov Chain Construction

- A **Markov chain** is a “random walk” through states, where the future depends only on the present
  - *Sequence of states:*  $\underline{x}^{(0)}, \underline{x}^{(1)}, \underline{x}^{(2)}, \dots$
  - *Initial state:* starting configuration
  - *Transition probabilities:*  $\Pr(\underline{x} \rightarrow \underline{x}')$
  - After  $t$  steps, distribution is  $\pi_t(\underline{x})$
  - When  $\pi_t(\underline{x}) = \pi_{t+1}(\underline{x})$ , chain reaches stationary distribution
- **Transition operator:** moves from one state to another:
  - Gibbs sampling: resample one variable given its Markov blanket
  - Metropolis–Hastings: propose a new state, then accept/reject based on a probability ratio
- There are algorithms generate a Markov Chain from a Bayesian network
- Under certain conditions:
  - Ergodicity: chain can reach any state
  - Aperiodicity: chain does not get stuck in cycles stationary distribution equals the **posterior distribution** over non-evidence variables given



# Markov Chain Monte Carlo: Mixing

- **Mixing** describes how quickly a Markov chain forgets its starting point and explores the whole state space efficiently
  - A well-mixed chain:
    - Moves between different high-probability regions often
    - Has low correlation between successive samples
  - Poor mixing:
    - Chain gets stuck in one mode for a long time
    - Leads to biased estimates and high variance
- In practice:
  - Discard initial samples as a burn-in period (before convergence)
  - After convergence, collected samples approximate the true posterior
- **Example**
  - If sampling from a bimodal distribution:
    - “Poor mixing” means the chain stays in one peak
    - “Good mixing” jumps between both



# Gibbs Sampling in Bayesian Networks

- Special case of Markov Chain Monte Carlo (MCMC) method that samples one variable at a time
- **Algorithm:**
  - Start with an initial complete assignment to all non-evidence variables
  - Keep evidence variables fixed at observed values
  - For each non-evidence variable  $X_i$ :
    - Sample  $X_i$  from  $P(X_i|\text{MB}(X_i))$  where  $\text{MB}(X_i)$  is the Markov blanket, i.e., parents, children, spouse of a node
- **Example:**
  - Weather network:  $P(\text{Cloudy}|\text{Sprinkler}, \text{Rain}, \text{WetGrass})$
  - Fix  $\text{WetGrass} = \text{true}$ ,  $\text{Sprinkler} = \text{true}$
  - Sample  $\text{Cloudy}$  and  $\text{Rain}$  iteratively
- **Pros**
  - Simple to implement for any Bayesian network
  - Handles large, complex graphs with local updates
- **Cons**
  - Can mix slowly if variables are highly correlated

# Metropolis–Hastings Sampling

- More general Markov Chain Monte Carlo method than Gibbs sampling
- **Algorithm**
  - Start at a current state  $\underline{x}$
  - Propose a new state  $\underline{x}'$  from a proposal distribution  $q(x'|x)$ , e.g.,
    - With 95% probability Gibbs sampling
    - Otherwise use importance sampling
  - Compute the acceptance probability:
    - $A(x, x') = \min(1, \frac{\pi(x')q(x|x')}{\pi(x)q(x'|x)})$
  - Move to  $x'$  with probability  $A(x, x')$ , otherwise stay at  $x$
- **Intuition**
  - Propose local moves
  - Accept if they lead to higher probability, or sometimes accept lower-probability states to explore
  - Balances exploration and exploitation to avoid getting stuck in local modes
- **Pros**
  - Very flexible: works with any proposal distribution
  - Can handle high-dimensional spaces