



How to Choose Priors

- **Weakly-informative priors** (aka “flat”, “vague”, “diffuse priors”)
 - Provide minimal information
 - Coefficient of linear regression centered around 0: $\beta \sim \text{Normal}(0, 10)$
- **Regularizing priors**
 - Known information about the parameter
 - Parameter is positive: $\sigma \sim \text{HalfCauchy}(0, 5)$
 - Parameter close to zero, above/below a number, or in a range
 - $\beta \sim \text{Laplace}(0, 1)$ (lasso prior) encourages sparsity
 - $\beta \sim \text{Normal}(0, 1)$ discourages extreme values
- **Informative priors**
 - Strong priors from previous knowledge (expert opinion, studies)
 - From experimental data: $\beta_1 \sim \text{Normal}(2.5, 0.5^2)$
 - From previous data, about 5% of cases positive: $p \sim \text{Beta}(2, 38)$
- **Prior elicitation**
 - Compute least informative distribution given constraints
 - Estimate distribution using maximum entropy to satisfy constraints
 - E.g., beta distribution with 90% of mass between 0.1 and 0.7

- *Communicating a Bayesian Analysis*
- Probabilistic Programming
- Posterior-Based Decisions

Communicating the Model of a Bayesian Analysis

1. Communicate assumptions / hypothesis


- Describe priors and probabilistic models
- E.g., coin-flip distributions:

$$\begin{cases} \theta \sim \text{Beta}(\alpha, \beta) \\ y \sim \text{Binomial}(n = 1, p = \theta) \end{cases}$$

2. Communicate Bayesian analysis result

- Describe posterior distribution
- Summarize location and dispersion
- Mean (or mode, median)
- Std dev
 - Misleading for skewed distributions
- Highest-posterior density (HPD)
 - Shortest interval containing a portion of probability density (e.g., 95% or 50%)
 - Amount is arbitrary (e.g., ArviZ defaults to 94%)

Kruschke diagram



Confidence Intervals vs Credible Intervals

- People confuse:
 - Frequentist **confidence intervals** with
 - Bayesian **credible intervals**
- In the frequentist framework, there is a true (unknown) parameter value
 - A **confidence interval** may or may not contain the true parameter value
 - Interpretation of a 95% confidence interval
 - No: *"There is a 95% probability that the true value is in this interval"*
 - Yes: *"If repeated many times, 95% of intervals would contain the true value"*
- In the Bayesian framework, parameters are random variables
 - Interpretation of a 95% **Bayesian credible interval**
 - *"There is a 95% probability that the true parameter lies within this interval, given the observed data"*
 - Bayesian **credible interval** is intuitive

Confidence Intervals vs Credible Intervals (ELI5)

- **Confidence Interval (Frequentist)**

- Imagine fishing in a lake without seeing the fish
- You throw your net
- 95% confidence interval: *"If I threw this net 100 times, about 95 nets would catch the fish."*
- Important: Once the net is thrown, it either caught the fish or not. The 95% makes sense across many attempts

- **Credible Interval (Bayesian)**

- Imagine a magical map showing where fish *probably* are, based on past observations
- 95% credible interval: *"Given my map, there's a 95% chance the fish is inside this part of the lake."*
- The fish's location is uncertain, and probability describes your belief


- **Key Difference**

- Confidence interval (Frequentist): Probability from repeating experiments
 - It's about the procedure, not the specific interval
- Credible interval (Bayesian): Probability describes your belief about the value, given the data
 - It's about *this interval*

- Communicating a Bayesian Analysis
- ***Probabilistic Programming***
- Posterior-Based Decisions

Bayesian Statistics

- Given:
 - The **"knows"**
 - Model structure (modeled as a graph of probability distributions)
 - Data, observations (modeled as constants)
 - The **"unknowns"**
 - Model parameters (modeled as probability distributions)
- Use Bayes' theorem to condition unknowns to knowns hoping to reduce the uncertainty about the unknowns
- **Problem**
 - Most probabilistic models are analytically intractable
- **Solution**
 - Probabilistic programming
 - Specify a probabilistic model using code
 - Solve models using numerical techniques



Probabilistic Programming Languages

- **Steps:**
 1. Specify models using code
 2. Numerical models solve inference problems without user understanding
 - Universal inference engines
 - PyMC3: flexible Python library for probabilistic programming
 - Theano: library to define, optimize, evaluate mathematical expressions using tensors
 - ArviZ: library to interpret probabilistic model results
- **Pros:**
 - Compute results without analytical closed form
 - Treat model solving as a black box
 - Focus on model design, evaluation, interpretation
- Probabilistic programming languages impact like Fortran on scientific computing
 - Build algorithms, ignore computational details

Coin Example: Numerical Solution (1/3)

- Assume you know the true value of θ (not true in general)
- Observe samples of the variable y
- Model the prior θ and the likelihood $y|\theta$

$$\begin{cases} \theta \sim \text{Beta}(\alpha = 1, \beta = 1) \\ Y \sim \text{Binomial}(n = 1, p = \theta) \end{cases}$$

- Run inference
- Generate samples of the posterior
- Summarize posterior
 - E.g., Highest-Posterior Density (HPD)

Coin Example: Numerical Solution (2/3)

```
[18]: np.random.seed(123)
n = 4
# Unknown value.
theta_real = 0.35

# Generate some observational data.
data = stats.bernoulli.rvs(p=theta_real, size=n)
data
```

```
[18]: array([1, 0, 0, 0])
```


```
[19]: with pm.Model() as our_first_model:
    # Prior.
    theta = pm.Beta('theta', alpha=1., beta=1.)
    # Likelihood.
    y = pm.Bernoulli('y', p=theta, observed=data)
    # (Numerical) Inference to estimate the posterior distribution through samples.
    idata = pm.sample(1000, random_seed=123)
```

```
Auto-assigning NUTS sampler...
Initializing NUTS using jitter+adapt_diag...
Multiprocess sampling (4 chains in 4 jobs)
NUTS: [theta]
```

Sampling 4 chains, 0 divergences ————— 100% 0:00:00 / 0:00:00

Sampling 4 chains for 1_000 tune and 1_000 draw iterations (4_000 + 4_000 draws total) took 1 seconds.

```
[20]: az.plot_trace(idata);
```




```
[21]: az.summary(idata)
```

```
[21]:      mean    sd  hdi_3%  hdi_97%  mcse_mean  mcse_sd  ess_bulk  ess_tail  r_hat
theta  0.324  0.179   0.031   0.653     0.005   0.003   1500.0   1737.0    1.0
```


- Generate data from ground truth model
- Build PyMC model matching mathematical model
- PyMC uses NUTS sampler, computes 4 chains
- No trace diverges
- Kernel density estimation (KDE) for posterior (should be Beta)
- Traces appear “noisy” and non-diverging (good)
- Numerical summary of posterior: mean, std dev, HDI
- $\mathbb{E}[\hat{\theta}] \approx 0.324$
- $\Pr(\hat{\theta} \in [0.031, 0.653]) = 0.94$

Coin Example: Numerical Solution (3/3)


```
[22]: az.plot_trace(idata, kind="rank_bars", combined=True);
```



```
[23]: az.plot_posterior(idata);
```



- Compute single KDE for all chains
- Rank plot to check results
- Histograms should look uniform, exploring different (and all) posterior regions
- Plot single KDE with all statistics

- 
- Communicating a Bayesian Analysis
 - Probabilistic Programming
 - *Posterior-Based Decisions*

Posterior-Based Decisions

- Sometimes describing the posterior is not enough
 - You need to make decisions based on our inference
- E.g., is the coin fair ($\theta = 0.5$) or biased?
 - Since $\mathbb{E}[\hat{\theta}] = 0.324$ it seems that the coin is biased
 - You can't rule out that the coin is unbiased since
 - $HPI = [0.03, 0.65]$
 - $0.5 \in HPI$
- If you want a sharper decision, you need to:
 - Collect more data to reduce the spread of the posterior
 - Define a more informative prior

Savage-Dickey Density Ratio

- The Savage-Dickey ratio tests a point null-hypotheses in Bayesian inference
- Idea:** compare prior and posterior densities at a single point θ_0

$$BF_{01} = \frac{p(\theta_0|H_1)}{p(\theta_0|\mathcal{D}, H_1)}$$

where:


- $p(\theta_0|H_1)$ is the *prior* density θ under the alternative hypothesis H_1 , evaluated at θ_0
- $p(\theta_0|\mathcal{D}, H_1)$ is the *posterior* density θ under H_1 evaluated at θ_0

Bayes Factor (BF)	Interpretation
1 - 3	Not enough evidence
3 - 10	Substantial evidence
10 - 100	Strong evidence
> 100	Decisive evidence

- Intuition:** this ratio shows how much data changes belief about θ_0
- If posterior density at θ_0 is much smaller than prior density, Bayes factor suggests strong evidence against H_0


Savage-Dickey Density Ratio: Example

```
[29]: az.plot_bf(idata1, var_name="theta", prior=np.random.uniform(0, 1, 10000), ref_val=0.5);
```



```
[30]: az.plot_bf(idata2, var_name="theta", prior=np.random.uniform(0, 1, 10000), ref_val=0.5);
```

The reference value is outside of the posterior. This translate into infinite support for H1



- H_0 : "coin is fair"
- The prior for H_0 is 0.87
- The posterior for H_0 is 1.15
- $BF_{10} = 0.87$

ROPE: Region of Practical Equivalence

- **ROPE** = an interval for a parameter where all values inside are considered “equivalent”
 - H_0 : “*coin is fair*” iff $\theta = 0.5$ is impractical
 - ROPE: $\theta \in [0.45, 0.55]$ is equivalent to 0.5
- **Hypothesis testing with ROPE and HPI**
 - Compare ROPE (Region Of Practical Equivalence) with HPI (Highest-Posterior Interval)
 - If HPI is within ROPE, no effect: H_1 is rejected
 - If HPI is outside ROPE, there is an effect: H_0 is rejected
 - If HPI overlaps with ROPE, result is inconclusive
- Decide ROPE before analysis based on domain knowledge
 - Picking it after analysis is like picking the p-value threshold after seeing the p-value

Loss Function: Motivation

- You need to make decisions based on our inference
- For many problems, decision cost is asymmetric
 - E.g., cost of a bad decision $>$ benefit of a good decision
 - E.g., vaccines may cause overreaction, but benefits outweigh risks
- To make the best decision, measure:
 - Benefits of a correct decision
 - Cost of a mistake
 - Decide trade-off between benefits and costs using a loss function
 - Use loss we function for decisions
- Loss quantifies *“how bad is an estimation mistake?”*
 - Larger loss indicates worse estimation

Loss Function

- Aka “cost function”
 - The inverse is known as “objective”, “fitness”, “utility function”
- Use a function to measure the difference between:
 - The true value θ ; and
 - The estimated value $\hat{\theta}$

Loss	Expression	Point estimate
Quadratic loss	$(\theta - \hat{\theta})^2$	Mean of posterior
Absolute loss	$ \theta - \hat{\theta} $	Median of posterior
1-0 loss	$I(\theta \neq \hat{\theta})$	Mode of posterior

- Making decisions in Bayesian statistics using loss function
 - Goal: pick a single value $\hat{\theta}$
 - You don't know the true value θ
 - Estimate θ in terms of the posterior distribution
 - Find the value $\hat{\theta}$ that minimizes the expected loss function