




# Tutorial: Bayesian Coin


---

- Bayesian Coin

- 
- ***Chemical Shift: Example***
  - Posterior Predictive Checks
  - Groups Comparison

# Chemical Shift

- Nuclear magnetic resonance (NMR)
  - Used to study molecules of living things
  - Measures observable quantities related to unobservable molecular properties (e.g., chemical shift)
- Data looks Gaussian with a couple of outliers



# Use of Gaussians in Statistics


---

- Aka “normal” distribution
- **Gaussians are easy to work with and abundant in nature**
- **Pros:**
  - Average of large sample size tends to be Gaussian (by Central Limit Theorem)
  - Many phenomena approximated using Gaussians (since they are average of effects)
  - Conjugate prior of Gaussian is Gaussian
- **Cons:**
  - Not robust to outliers
  - Important to relax assumption of Gaussianity

# Chemical Shift: Example

- Assume Gaussian is a decent approximation of the data
- The **likelihood**  $y|\mu, \sigma$  comes from a normal distribution:

$$Y \sim N(\mu, \sigma)$$



- Priors** for mean and sigma of  $Y$ 
  - Mean** from uniform distribution  $U(l, h)$ :

$$\mu \sim U(l, h)$$

- Set  $\mu$  larger than data range, e.g., [40, 70]
- Std dev from half-normal distribution (i.e., a regular normal, but restricted to non-negative values):

$$\sigma \sim \text{HalfNormal}(0, \sigma_\sigma)$$

- If unknown, set  $\sigma_\sigma = 10$

# Chemical Shift: PyMC

```
[87]: with pm.Model() as model_g:
      # The mean is Uniform in [40, 70] (which is larger than the data).
      mu = pm.Uniform("mu", lower=40, upper=70)
      # The std dev is half normal with a large value (which is a large value based on the data).
      sigma = pm.HalfNormal("sigma", sigma=10)
      # The model is N(mu, sigma).
      y = pm.Normal("y", mu=mu, sigma=sigma, observed=data)
      # Sample.
      idata_g = pm.sample(1000)
```

Auto-assigning NUTS sampler...  
Initializing NUTS using jitter+adapt\_diag...  
Multiprocess sampling (4 chains in 4 jobs)  
NUTS: [mu, sigma]

Sampling 4 chains, 0 divergences ————— 100% 0:00:00 / 0:00:00

Sampling 4 chains for 1\_000 tune and 1\_000 draw iterations (4\_000 + 4\_000 draws total) took 1 seconds.

- The PyMC code is **one-to-one with the model**:


$$\begin{cases} \mu \sim U(l = 40, h = 70) \\ \sigma \sim \text{HalfNormal}(0, \sigma_\sigma = 10) \\ Y \sim N(\mu, \sigma) \end{cases}$$

- Get 1000 samples from the posterior


# Chemical Shift: PyMC

- Compute 4 traces for 2 variables  $\mu$ ,  $\sigma$
- Results are well-formed
- $\mu \in [52.55, 54.45]$
- $\sigma \in [2.86, 4.23]$
- Is the model good?

```
[92]: # There are 4 traces for 2 variables.
      az.plot_trace(idata_g);
```



```
[93]: # The posterior distribution of the params is bi-dimensional, since it has mu and sigma.
      az.plot_pair(idata_g, kind='kde', marginals=True);
```



```
[94]: # Report a summary of the inference.
      az.summary(idata_g, kind="stats").round(2)
```

	mean	sd	hdi_3%	hdi_97%
<b>mu</b>	53.50	0.51	52.55	54.46
<b>sigma</b>	3.55	0.38	2.86	4.23



- Chemical Shift: Example
- *Posterior Predictive Checks*
- Groups Comparison

# Samples from Posterior Distribution

- Given a **posterior distribution**  $\Pr(\theta|y)$ , you can generate predictions  $\tilde{y}$  based on the data  $y$  and the estimated parameters  $\hat{\theta}$ :

$$\Pr(\tilde{y}|y) = \int_{\theta} \Pr(\tilde{y}|\theta) \Pr(\theta|y) d\theta = \int \text{model} \times \text{posterior}$$

- This is called the “**posterior predictive distribution**” as it predicts **future data** using the **posterior distribution**
- Conceptually:
  - Sample a value of  $\theta$  from the posterior  $\Pr(\theta|y)$
  - Feed the value of  $\theta$  to the likelihood  $\Pr(y|\theta)$
  - Obtain  $\tilde{y}$
- This process has two sources of uncertainty:
  - Parameter uncertainty
    - Captured by the posterior  $\Pr(\theta|y)$
  - Sampling uncertainty
    - Captured by the likelihood  $\Pr(y|\theta)$

# Posterior Predictive Check (PPC)

---

- **Intuition:** can the model reproduce observed data?
- **PPC approach:**
  - Generate predictions  $\tilde{y}$  with observed data  $y$
  - Check consistency between predicted values and observed data
- Models should always be checked
- Differences can arise by:
  - Mistakes
  - Limitations of the model
    - E.g., the model works well for average behavior but fails to predict rare values
  - Limitations of the data

# Bayesian Workflow Using PPC


---

1. Given a process
  - True distribution of the process is unknown / unknowable
2. **Sample the process**
  - Get finite sample  $y$  through sampling
  - E.g., experiment, survey, simulation
3. **Inference**
  - Build probabilistic model using prior  $\Pr(\theta)$  and likelihood  $\Pr(y|\theta)$  to get posterior distribution  $\Pr(\theta|y)$
  - Posterior distribution: distribution of model parameters  $\theta$  given data
4. **Predictive distribution**
  - Compute predictions from posterior distribution (i.e., posterior predictive distribution)
  - Posterior predictive distribution is the distribution of predicted samples averaged over posterior distribution
5. **Validation**
  - Validate model by comparing original samples vs predicted samples

# Chemical Shift Example: PPC

```
[95]: # Compute 100 posterior predictive samples.  
y_pred_g = pm.sample_posterior_predictive(idata_g, model=model_g)  
  
Sampling: [y]  
  
Sampling ... 100% 0:00:00 / 0:00:00
```

```
•[96]: # Black: KDE of the data (observed)  
# Blue: KDEs of the posterior predictive samples  
# Orange: KDE of the posterior predictive mean  
az.plot_ppc(y_pred_g, mean=True, num_pp_samples=100);
```




- Sample the posterior
- Apply the model
- Get predictive posterior distribution (dashed orange distribution)
- Compare to data (black distribution)
- **Is the PPC model good?**
- **No**
  - Posterior mean is more to the right than data
  - Posterior std dev is larger

# Chemical Shift: Model Critique

```
[95]: # Compute 100 posterior predictive samples.  
y_pred_g = pm.sample_posterior_predictive(idata_g, model=model_g)  
Sampling: [y]  
Sampling ... 100% 0:00:00 / 0:00:00
```

```
•[96]: # Black: KDE of the data (observed)  
# Blue: KDEs of the posterior predictive samples  
# Orange: KDE of the posterior predictive mean  
az.plot_ggcy_pred_g, mean=True, num_pp_samples=100;
```



## • Problem


- Two data points on the tails of the distribution
- The normal distribution is “surprised” by these points and “reacts” by adjusting the mean towards them and increasing the standard deviation

## • Solution

- Declare points as outliers and discard
  - E.g., equipment malfunction (need evidence)
- Change the model
- Bayesians prefer to encode assumptions into the model (e.g., priors, likelihoods)
  - Rather than ad-hoc heuristics (e.g., outlier removal rules)

# Student's t-distribution: Recap


- Student's t-distribution has 3 params:
  1. Mean  $\mu$ 
    - It doesn't always exist
  2. Scale  $\sigma$ 
    - Similar to std dev, but it doesn't always exist
  3. Degrees of freedom  $\nu \in [0, \infty]$ 
    - Aka “normality parameter”, since it controls how “normal” is the distribution
    - With  $\nu = 1$ : heavy tails and no mean (Cauchy)
    - With  $\nu \rightarrow \infty$  we recover the Gaussian
- Heavy tails (high kurtosis) means *“values are more likely to be far from the mean compared to a Normal”*



# Chemical Shift: Use Student's t-dist (1/3)

- Use Student's t-distribution model instead of Normal

$$\begin{cases} \mu \sim U(l, h) \\ \sigma \sim \text{HalfNormal}(0, \sigma) \\ \nu \sim \text{Exp}(\lambda) \\ y \sim \text{StudentT}(\mu, \sigma, \nu) \end{cases}$$



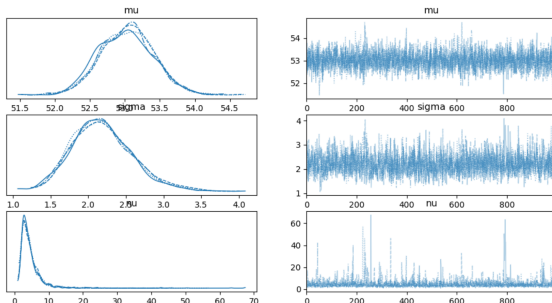


# Chemical Shift: Use Student's t-dist (2/3)

```
[102]: # Use a Student-T model.
with pm.Model() as model_t:
    mu = pm.Uniform("mu", 40, 75)
    sigma = pm.HalfNormal("sigma", sigma=10)
    # A student with nu = 30 is close to a Gaussian.
    nu = pm.Exponential("nu", 1/30)
    #
    y = pm.StudentT("y", mu=mu, sigma=sigma, nu=nu, observed=data)
    idata_t = pm.sample(1_000)
```

Auto-assigning NUTS sampler... \*\*\*

```
az.plot_trace(idata_t);***
```



```
[104]: az.summary(idata_t, kind="stats").round(2)
```

```
[104]:
```


	mean	sd	hdi_3%	hdi_97%
mu	53.03	0.38	52.35	53.76
sigma	2.19	0.40	1.45	2.95
nu	4.65	3.92	1.20	9.20

- Outliers decrease  $\nu$  (less Gaussian) instead of increasing mean and standard deviation
  - $\mu$  similar to Gaussian estimate
  - $\sigma$  smaller
  - $\nu \approx 5$  (not very Gaussian)
- Estimation more robust; outliers have less effect


# Chemical Shift: Use Student's t-dist (3/3)

```
[105]: # Compute 100 posterior predictive samples.  
y_ppc_t = pm.sample_posterior_predictive(idata_t, model_t);  
Sampling: [y]  
Sampling ... 100% 0:00:00 / 0:00:00
```

```
[106]: ax = az.plot_ppc(y_ppc_t, num_pp_samples=100, mean=True)  
ax.set_xlim(40, 70);
```




- PPC (posterior predictive check) fits better than Normal model
- Plot is “hairy” because KDE is estimated only in data interval and 0 outside

- 
- Chemical Shift: Example
  - Posterior Predictive Checks
  - ***Groups Comparison***

# Group Comparison

---

- **Group comparison** tests for statistically significant results between “treatment” and “control group”
- E.g., 
  - How well do patients respond to a new drug vs a placebo?
  - Is there a reduction in car accidents after new traffic regulation?
  - Does college student performance improve without cellphones at school?
- **Effect size** quantifies the difference between two groups
  - Moves from “does it work?” (hypothesis testing) to “how well does it work?” (estimate effect size)

# Bogus Control Groups

---

- When something is claimed to be harder/better/faster/stronger, ask for the baseline used for comparison
  - E.g.,
    - Sell sugary yogurts to boost the immune system by comparing it to using milk
    - A better control group would be a less sugary yogurt
- **Placebo** is a psychological phenomenon where a patient experiences improvements after receiving an inactive treatment
  - Using a placebo is better than “no treatment”
  - Shows difficulty in accounting for all factors in an experiment

# Group Comparison Bayesian-Style

---

- **Frequentist approach**
  - Compare p-value of difference of means in each group
- **Bayesian approach**
  - Compare posterior distribution of means between groups using:
    - Plot of posterior
    - Cohen's d
    - Probability of superiority

# Sample Size Effect

---

- **Sample size effect** is the impact of the number of observations on statistical results (e.g., p-values, confidence intervals)
  - Small sample: Large mean difference might not be statistically significant (low p-value)
  - Large sample: Tiny mean difference can be highly significant (small p-value) but meaningless
  - E.g., Cohen's d

# Cohen's d

---

- **Cohen's d** is the difference of means relative to pooled standard deviation

$$\frac{\mu_2 - \mu_1}{\sqrt{(\sigma_1^2 + \sigma_2^2)/2}}$$

- Normalizes effect by variability for pooled std dev
- Bayesian analysis estimates actual std dev
- Variability of each group normalizes mean difference
  - Similar to a Z-score, number of std dev values differ
- Compute posterior distribution of means and std
- Compute distribution of Cohen's d or use formula